

Transformers

Naeemullah Khan


naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

August 9, 2026



Source: [Roboflow — What is DETR?](#) (Carion et al., 2020)

1. Motivation
2. Learning Outcomes
3. General Attention Layer
4. Self-Attention Layer
5. Permutation Equivariance
6. Transformer Overview
7. RNNs vs Transformers
8. Limitations and Future Directions
9. References

Why this topic matters:


- ▶ RNNs are sequential and struggle with long-range dependencies.
- ▶ Attention lets a model focus on the most relevant parts of the input.
- ▶ Transformers replace recurrence with self-attention, making them fully parallelizable.
- ▶ They are the backbone of modern NLP and, increasingly, computer vision.
- ▶ *“Attention is all you need.”* — Vaswani et al., 2017.

By the end of this session, you will be able to:

- ▶ Explain the general attention layer and the role of queries, keys, and values.
- ▶ Describe self-attention and why it is permutation equivariant.
- ▶ Understand positional encodings, masked attention, and multi-head attention.
- ▶ Describe the Transformer block and how blocks stack into a full Transformer.
- ▶ Compare Transformers with RNNs in terms of parallelism and long-range modeling.

Transformers: **General Attention Layer**

Attention We Just Saw in Image Captioning




Outputs:
context vector: $\mathbf{c}$ (shape: D)

Operations:
Alignment: $e_{i,j} = f_{\text{att}}(\mathbf{h}, \mathbf{z}_{i,j})$
Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$
Output: $\mathbf{c} = \sum_{i,j} a_{i,j} \mathbf{z}_{i,j}$

Inputs:
Features: $\mathbf{z}$ (shape: $H \times W \times D$)
Query: $\mathbf{h}$ (shape: D)

Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers




Outputs:
context vector: $\mathbf{c}$ (shape: D)

Operations:
Alignment: $\mathbf{e}_i = f_{\text{att}}(\mathbf{h}, \mathbf{x}_i)$
Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$
Output: $\mathbf{c} = \sum_i a_i \mathbf{x}_i$

Inputs:
Input vectors: $\mathbf{x}$ (shape: $N \times D$)
Query: $\mathbf{h}$ (shape: D)

- The attention operation is permutation invariant.
- It does not depend on the ordering of the features.
- Stretch $H \times W = N$ into N vectors.


Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers

Change $f_{att}(\cdot)$ to a scaled simple dot product:

- ▶ Larger dimensions mean more terms in the dot product sum.
- ▶ Therefore, the variance of the logits is higher. Large-magnitude vectors will produce much higher logits.
- ▶ As a result, the post-softmax distribution has lower entropy, assuming logits are IID.
- ▶ Ultimately, these large-magnitude vectors will cause softmax to peak and assign very little weight to all others.
- ▶ Divide by $\sqrt{D}$ to reduce the effect of large-magnitude vectors.



Outputs:

context vectors: y (shape: D)

Operations:

Alignment: $e_{i,j} = q_j \cdot x_i / \sqrt{D}$

Attention: $a = \text{softmax}(e)$

Output: $y_j = \sum_i a_{i,j} x_i$


- We can have multiple query vectors.
- Each query creates a new output context vector.

Inputs:

Input vectors: x (shape: $N \times D$)

Queries: q (shape: $M \times D$)

Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers




Outputs:
context vectors: $\mathbf{y}$ (shape: D)

Operations:
Alignment: $e_{i,j} = q_j \cdot x_i / \sqrt{D}$
Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$
Output: $y_i = \sum_j a_{i,j} x_j$

Inputs:
Input vectors: $\mathbf{x}$ (shape: $N \times D$)
Queries: $\mathbf{q}$ (shape: $M \times D$)


- Notice that the input vectors are used for both the alignment and the attention calculations.

Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers

- Notice that the input vectors are used for both the alignment and the attention calculations.
- We can add more expressivity to the layer by adding a different fully connected (FC) layer before each of the two steps.



Outputs:

context vectors: $\mathbf{y}$ (shape: D_v)

Operations:

Key vectors: $\mathbf{k} = \mathbf{x}W_k$

Value vectors: $\mathbf{v} = \mathbf{x}W_v$

Alignment: $e_{i,j} = q_i \cdot k_j / \sqrt{D}$

Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$

Output: $y_j = \sum_i a_{i,j} v_i$

Inputs:

Input vectors: $\mathbf{x}$ (shape: $N \times D$)


Queries: $\mathbf{q}$ (shape: $M \times D_k$)

- Notice that the input vectors are used for both the alignment and the attention calculations.
- We can add more expressivity to the layer by adding a different fully connected (FC) layer before each of the two steps.
- The input and output dimensions can now change depending on the key and value FC layers.

Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers


Transformers: **Self-Attention Layer**

- ▶ **Self-Attention:** Computes attention within a single sequence.
- ▶ **How it works:**
 - Each token attends to every other token.
 - Captures global context, not just local dependencies.



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers

- Recall that the query vector is a function of the input vectors.
- We can calculate the query vectors from the input vectors, thus defining a “self-attention” layer.
- There are no input query vectors anymore.
- Instead, query vectors are calculated using a fully connected (FC) layer.



Outputs:
context vectors: y (shape: D_v)

Operations:

Key vectors: $k = xW_k$

Value vectors: $v = xW_v$

Query vectors: $q = xW_q$

Alignment: $e_{i,j} = q_i \cdot k_j / \sqrt{D}$

Attention: $a = \text{softmax}(e)$



Output: $y_j = \sum_i a_{i,j} v_i$

Inputs:
Input vectors: x (shape: $N \times D$)

- Recall that the query vector is a function of the input vectors.
- We can calculate the query vectors from the input vectors, thus defining a “self-attention” layer.
- There are no input query vectors anymore.
- Instead, query vectors are calculated using a fully connected (FC) layer.

Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers


Self-Attention Layer



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers

Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers


Transformers: **Permutation Equivariance**



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers


- ▶ Self-Attention Layer is Permutation equivariant
- ▶ It doesn't care about the order of the inputs!
- ▶ **Problem:** How can we encode ordered sequences like language or spatially ordered image features?

- ▶ Concatenate/add special positional encoding p_j to each input vector x_j
- ▶ We use a function $\text{pos}: \mathbb{N} \rightarrow \mathbb{R}^D$ to process the position j of the vector into a D -dimensional vector
- ▶ So, $p_j = \text{pos}(j)$



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers


Masked self-attention layer



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers



- Prevent vectors from looking at future vectors.
- Manually set alignment scores to $-\infty$

- Multiple self-attention heads in parallel



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers

General attention versus self-attention



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers

Example: CNN with Self-Attention

Input Image



Cat image is free to use under the [Pixabay License](#)




CNN



Features:
 $C \times H \times W$


Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers; module after Zhang et al., "Self-Attention Generative Adversarial Networks," ICML 2019

Example: CNN with Self-Attention




Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers; module after Zhang et al., "Self-Attention Generative Adversarial Networks," ICML 2019

Example: CNN with Self-Attention




Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers; module after Zhang et al., "Self-Attention Generative Adversarial Networks," ICML 2019

Example: CNN with Self-Attention




Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers; module after Zhang et al., "Self-Attention Generative Adversarial Networks," ICML 2019

Example: CNN with Self-Attention



Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers; module after Zhang et al., "Self-Attention Generative Adversarial Networks," ICML 2019

Example: CNN with Self-Attention




Source: Stanford CS231n (2022), Lecture 11: Attention and Transformers; module after Zhang et al., "Self-Attention Generative Adversarial Networks," ICML 2019

Transformers: **Transformer Overview**

Attention is all you need


Vaswani et al, NeurIPS 2017

- ▶ **Input:** Set of vectors x
- ▶ **Output:** Set of vectors y
- ▶ Self-attention is the only interaction between vectors!
- ▶ Layer normalization and MLP operate independently on each vector.
- ▶ Highly scalable and highly parallelizable.



Source: UMich EECS 498 (2022), Lecture 17: Attention, J. Johnson

- ▶ A Transformer consists of a sequence of transformer blocks.
- ▶ Vaswani et al.'s base model used 6 encoder and 6 decoder blocks, $d_{\text{model}} = 512$, and 8 attention heads.



Source: UMich EECS 498 (2022), Lecture 17: Attention, J. Johnson

Transformers: **RNNs vs Transformers**

Aspect	RNN / LSTM	Transformer
Computation	Sequential	Parallel
Long-range dependencies	Hard (signal decays)	Direct (any-to-any)
Path length between tokens	$O(n)$	$O(1)$
Complexity per layer	$O(n \cdot d^2)$	$O(n^2 \cdot d)$
Order information	Built-in (recurrence)	Added (positional encoding)

- ▶ Transformers trade quadratic attention cost for full parallelism and shorter gradient paths.
- ▶ This makes them far more scalable on modern hardware, which is the key to large pretrained models.

Transformers: **Limitations and Future Directions**

- ▶ **Quadratic complexity:** self-attention scales as $O(n^2)$ in sequence length.
- ▶ **Data hungry:** need large datasets or pretraining to perform well.
- ▶ **Weak inductive biases:** unlike CNNs/RNNs, no built-in locality or order bias; order must be injected via positional encodings.
- ▶ **Memory intensive:** long sequences are expensive to store and compute.

- ▶ **Efficient attention:** Linformer, Performer, Longformer, FlashAttention.
- ▶ **Vision:** extend self-attention to images (Vision Transformers, covered later in the course).
- ▶ **Multimodality:** unify text, image, and audio in a single architecture (e.g., Flamingo).

Transformers: **References**

- [1] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is All You Need. *Advances in Neural Information Processing Systems*, 30. <https://arxiv.org/abs/1706.03762>
- [2] Bahdanau, D., Cho, K., & Bengio, Y. (2015). Neural Machine Translation by Jointly Learning to Align and Translate. *International Conference on Learning Representations*. <https://arxiv.org/abs/1409.0473>
- [3] Zhang, H., Goodfellow, I., Metaxas, D., & Odena, A. (2019). Self-Attention Generative Adversarial Networks. *International Conference on Machine Learning*, 7354–7363. <https://arxiv.org/abs/1805.08318>
- [4] Fei-Fei Li, Johnson, J., & Yeung, S. Stanford CS231n: Deep Learning for Computer Vision. <http://cs231n.stanford.edu/>

- [5] Johnson, J. UMich EECS 498.008/598.008: Deep Learning for Computer Vision.
<https://web.eecs.umich.edu/~justincj/teaching/eecs498/WI2022/>