

Vision Transformers

Naeemullah Khan


naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

June 20, 2026



Source: miro.medium.com

1. Motivation
2. Learning Outcomes
3. Vision Transformers
4. Improving ViT: Distillation
5. CNN vs ViT
6. Hierarchical ViT: Swin Transformer
7. Other Hierarchical Vision Transformers
8. Limitations and Future Directions
9. References

Why this topic matters:

- ▶ Transformers dominate NLP — can the same idea work for images?
- ▶ Vision Transformers (ViTs) treat an image as a sequence of patches.
- ▶ With enough data, ViTs match or beat CNNs on many benchmarks.
- ▶ Hierarchical and hybrid designs make them practical for detection and segmentation.
- ▶ *"An image is worth 16x16 words."* — Dosovitskiy et al., 2020.

By the end of this session, you will be able to:

- ▶ Explain how an image is converted into patch embeddings for a Transformer.
- ▶ Describe the ViT architecture (patches, class token, positional encodings, MLP head).
- ▶ Understand data-efficient training via distillation (DeiT).
- ▶ Compare CNNs and ViTs in terms of inductive bias and data efficiency.
- ▶ Describe hierarchical ViTs (Swin) and Transformer-based detection (DETR).
- ▶ Recognize that modern CNNs (ConvNeXt) remain competitive.

Advanced Computer Vision: **Vision Transformers**


Key idea:

- ▶ Split image into fixed-size patches (e.g., 16×16)
- ▶ Flatten and project patches into embeddings
- ▶ Add positional encodings to embeddings
- ▶ Apply standard Transformer encoder

Architecture:

- ▶ Patch embedding
- ▶ Class token
- ▶ Transformer blocks
- ▶ MLP Head

Source: Dosovitskiy et al., "An Image is Worth 16×16 Words", 2020




N input patches, each
of shape $3 \times 16 \times 16$



Linear projection to
D-dimensional vector


N input patches, each
of shape $3 \times 16 \times 16$





Add positional
embedding: learned D-
dim vector per position


Linear projection to
D-dimensional vector

N input patches, each
of shape 3x16x16










⁰Dosovitskiy et al, "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale", ICLR 2021

- ▶ **Patch Embedding:** Achieved via a linear projection of flattened patches ($P = 16$ typical)
- ▶ **Differences from CNN stems:**
 - Large stride causes optimization instability
 - Mitigated via convolutional preprocessing (e.g., small conv stem)



- ▶ ViT learns the grid-like structure of the image patches via its position embeddings.
- ▶ The lower layers contain both global and local features, while the higher layers contain only global features.

Advanced Computer Vision: **Improving ViT:** **Distillation**

Step 1: Train a teacher CNN on ImageNet



$P(\text{cat}) = 0.9$
 $P(\text{dog}) = 0.1$

→ Cross
Entropy
Loss


← GT label:
Cat

Step 2: Train a student ViT to match ImageNet predictions from the teacher CNN (and match GT labels)



$P(\text{cat}) = 0.1$
 $P(\text{dog}) = 0.9$


→ KL Divergence Loss



$P(\text{cat}) = 0.2$
 $P(\text{dog}) = 0.8$

→ Cross
Entropy
Loss

← GT label:
Dog



⁰Touvron et al, "Training data-efficient image transformers distillation through attention", ICML 2021

Advanced Computer Vision: **CNN** vs **ViT**

CNN



1. Maintain 2D structure logic



2. Shift equivariant



3. Consider only local correlations



4. Hierarchically growing field of view



5. Hierarchically progressing complexity




6. Reasonable amount of params



7. Global representation

ViT






In most CNNs (including ResNets), **decrease** resolution and **increase** channels as you go deeper in the network (Hierarchical architecture)

Useful since objects in images can occur at various scales

In a ViT, all blocks have same resolution and number of channels (Isotropic architecture)


Can we build a **hierarchical** ViT model?




Feature	CNNs	Vision Transformers
Locality	Strong inductive bias	Weak (learns from data)
Parameter sharing	Yes	Yes
Global context	Limited	Native support
Training data	Efficient on small data	Needs large data or pretraining
Flexibility	Less modular	More modular

Advanced Computer Vision: **Hierarchical ViT: Swin Transformer**


Hierarchical ViT: Swin Transformer






⁰Liu et al, "Swin Transformer: Hierarchical Vision Transformer using Shifted Windows", CVPR 2021

Hierarchical ViT: Swin Transformer



- ▶ Swin Transformer introduces windowed self-attention with shifting windows, merging tokens akin to pooling.¹
- ▶ Builds a pyramid representation, enabling linear compute complexity and high performance on detection/segmentation (e.g. 58.7 COCO box AP).
- ▶ Hierarchical design allows for multi-scale feature extraction, similar to CNNs.


Hierarchical ViT: Swin Transformer (cont.)



¹<https://arxiv.org/abs/2103.14030>


²Liu et al, "Swin Transformer: Hierarchical Vision Transformer using Shifted Windows", CVPR 2021

MViT




Fan et al, "Multiscale Vision Transformers", ICCV 2021

Swin-V2



Liu et al, "Swin Transformer V2: Scaling up Capacity and Resolution", CVPR 2022


Improved MViT



Li et al, "Improved Multiscale Vision Transformers for Classification and Detection", arXiv 2021


Advanced Computer Vision: **Other Hierarchical Vision Transformers**

MViT




Fan et al, "Multiscale Vision Transformers", ICCV 2021

Swin-V2



Liu et al, "Swin Transformer V2: Scaling up Capacity and Resolution", CVPR 2022

Improved MViT





Li et al, "Improved Multiscale Vision Transformers for Classification and Detection", arXiv 2021

- ▶ Simple object detection pipeline: directly outputs a set of boxes from a Transformer
- ▶ No anchors, no regression of box transforms
- ▶ Matches predicted boxes to ground truth boxes with bipartite matching; trains to regress box coordinates

²Carion et al., “End-to-End Object Detection with Transformers,” ECCV 2020



Object Detection with Transformers: DETR



²Carion et al., "End-to-End Object Detection with Transformers," ECCV 2020

- ▶ DETR (DEtection TRansformer, Carion et al., 2020) introduces a simple and unified object detection pipeline using Transformers.
- ▶ Transformers replace region proposal networks (RPNs) with attention-based object queries for end-to-end detection.
- ▶ DETR directly predicts a fixed set of bounding boxes and class labels, eliminating the need for anchors, hand-crafted box proposals, or non-maximum suppression (NMS).
- ▶ The approach simplifies the pipeline and enables global reasoning, though it may struggle with small objects and require higher compute.
- ▶ DETR uses bipartite matching (Hungarian algorithm) to uniquely match predicted boxes to ground truth boxes, and is trained end-to-end to regress box coordinates and classify objects.

Object Detection with Transformers: DETR (cont.)



³Carion et al., "End-to-End Object Detection with Transformers", ECCV 2020

<https://arxiv.org/abs/2005.12872>

⁴<https://docs.google.com/presentation/d/1X0BDhpJ0a3IOf29DU8vAB4WJruyLugIX/edit?slide=id.pi#>

- ▶ ConvNeXt is a modern CNN architecture that incorporates insights from transformer models.
- ▶ It uses a hierarchical design similar to Swin Transformer, with a focus on simplicity and efficiency.
- ▶ Key features include:
 - Depthwise separable convolutions
 - Layer normalization
 - Global average pooling
- ▶ Achieves competitive performance on image classification tasks while maintaining the strengths of CNNs.
- ▶ ConvNeXt demonstrates that CNNs can still be effective with modern design principles, even in the era of transformers.




Figure 4. **Block designs** for a ResNet, a Swin Transformer, and a ConvNeXt. Swin Transformer's block is more sophisticated due to the presence of multiple specialized modules and two residual connections. For simplicity, we note the linear layers in Transformer MLP blocks also as " 1×1 convs" since they are equivalent.

5

⁵Liu et al., "A ConvNet for the 2020s", CVPR 2022 <https://arxiv.org/abs/2201.03545>

Vision Transformers: **Limitations and Future Directions**

- ▶ **Data hungry:** Performance drops significantly on small datasets.
- ▶ **Computationally expensive:** Quadratic complexity of self-attention.
- ▶ **Lacks built-in locality bias:** Unlike CNNs; some variants (e.g., Swin Transformer) address this.
- ▶ **Requires large-scale pretraining:** Needs datasets like JFT-300M, ImageNet-22K for strong performance.

- ▶ **Efficient transformers:** [Linformer](#), [Performer](#), [Longformer](#).
- ▶ **Hybrid models:** Combining CNNs and Transformers (e.g., [CoaT](#), [CvT](#)).
- ▶ **Low-resource ViTs:** Data-efficient training (e.g., [DeiT](#)).
- ▶ **Vision-language models:** [CLIP](#), [Flamingo](#), [GPT-4V](#).
- ▶ **Hardware acceleration:** [Edge ViTs](#), quantized ViTs for deployment.

Vision Transformers: **References**

- [1] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is All You Need. *Advances in Neural Information Processing Systems*, 30. <https://arxiv.org/abs/1706.03762>
- [2] Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., ... & Houlsby, N. (2020). An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. *International Conference on Learning Representations*.
<https://arxiv.org/abs/2010.11929>
- [3] Touvron, H., Cord, M., Douze, M., Massa, F., Sablayrolles, A., & Jégou, H. (2021). Training data-efficient image transformers distillation through attention. *International Conference on Machine Learning*. <https://arxiv.org/abs/2012.12877>
- [4] Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., Zhang, Z., ... & Guo, B. (2021). Swin Transformer: Hierarchical Vision Transformer using Shifted Windows. *CVPR 2021*.
<https://arxiv.org/abs/2103.14030>

- [5] Chen, J., Chen, J., Chao, H., & Yang, M. (2021). Pre-Trained Image Processing Transformer. *CVPR 2021*. <https://arxiv.org/abs/2012.00364>
- [6] Lu, Y., Dosovitskiy, A., Houlsby, N., & Beyer, L. (2021). Vision Transformer: A Survey. *arXiv preprint arXiv:2111.06091*. <https://arxiv.org/abs/2111.06091>
- [7] Xie, E., Wang, W., Yu, Z., Anandkumar, A., Alvarez, J. M., & Luo, P. (2021). SegFormer: Simple and Efficient Design for Semantic Segmentation with Transformers. *NeurIPS 2021*. <https://arxiv.org/abs/2105.15203>
- [8] Zheng, S., Lu, J., Zhao, H., Zhu, X., Luo, Z., Wang, Y., ... & Fu, Y. (2021). Rethinking Semantic Segmentation from a Sequence-to-Sequence Perspective with Transformers (SETR). *CVPR 2021*. <https://arxiv.org/abs/2012.15840>
- [9] Strudel, R., Garcia, R., Laptev, I., & Schmid, C. (2021). Segmenter: Transformer for Semantic Segmentation. *ICCV 2021*. <https://arxiv.org/abs/2105.05633>

- [10] Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021). Learning Transferable Visual Models From Natural Language Supervision (CLIP). *ICML 2021*. <https://arxiv.org/abs/2103.00020>
- [11] Alayrac, J.-B., Donahue, J., Luc, P., Miech, A., Barr, I., Hasson, Y., ... & Simonyan, K. (2022). Flamingo: a Visual Language Model for Few-Shot Learning. *arXiv preprint arXiv:2204.14198*. <https://arxiv.org/abs/2204.14198>
- [12] Oquab, M., Darcet, T., Moutakanni, T., Vo, D. H., Szafraniec, M., Khalidov, V., ... & Jegou, H. (2023). DINOv2: Learning Robust Visual Features without Supervision. *arXiv preprint arXiv:2304.07193*. <https://arxiv.org/abs/2304.07193>
- [13] Fei-Fei Li, Y., Li, Y., & Gao, R. Stanford CS231n: Deep Learning for Computer Vision. <http://cs231n.stanford.edu/>

- [14] Shocher, A., Bagon, S., Galun, M., & Dekel, T. WAIC DL4CV: Deep Learning for Computer Vision: Fundamentals and Applications.
<https://www.cs.tau.ac.il/~assafshocher/WAIC-DL4CV/>
- [15] Johnson, J. UMich EECS 498.008/598.008: Deep Learning for Computer Vision.
<https://web.eecs.umich.edu/~justincj/teaching/eecs498/WI2023/>
- [16] Hugging Face. Vision Transformer Playground.
<https://huggingface.co/spaces/akhaliq/ViT>
- [17] Stanford CS231n. Lecture Slides: Convolutional Neural Networks for Visual Recognition.
<http://cs231n.stanford.edu/slides/>
- [18] Justin Johnson. UMich EECS 498.008/598.008: Deep Learning for Computer Vision Slides. <https://web.eecs.umich.edu/~justincj/teaching/eecs498/slides.html>

[19] PyTorch. PyTorch Tutorials: Vision.

https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html