

Advanced Computer Vision

Naeemullah Khan

naeemullah.khan@kaust.edu.sa




جامعة الملك عبد الله
للعلوم والتقنية

King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

June 18, 2026



1. Motivation
2. Learning Outcomes
3. Recurrent Neural Networks (RNN)
4. Image Captioning
5. Sequence-to-Sequence with RNNs
6. Image Captioning with RNNs and Attention
7. General Attention Layer
8. Self Attention Layer
9. Permutation Invariance
10. Transformer Overview
11. Vision Transformers
12. Improving ViT: Distillation
13. CNN vs ViT

14. Hierarchical ViT: Swin Transformer

15. Other Hierarchical Vision Transformers

16. Limitations and Future Directions

17. References

Why this topic matters:

- ▶ Traditional CNNs struggle with sequential and long-range dependencies.
- ▶ RNNs were a start, but they come with their own baggage.
- ▶ Attention mechanisms and Transformers revolutionized NLP – now transforming vision tasks.
- ▶ Vision Transformers (ViTs) are outperforming CNNs on many benchmarks.
- ▶ *"Understanding the Transformer is a must for any modern deep learning practitioner."*

By the end of this session, you will be able to:

- ▶ Explain the working and limitations of RNNs.
- ▶ Understand attention and self-attention mechanisms.
- ▶ Grasp the Transformer architecture.
- ▶ Apply Transformers to computer vision tasks.
- ▶ Compare CNNs and ViTs in terms of performance and design.
- ▶ Explore current applications and future directions.

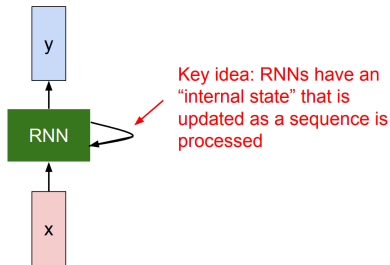
Advanced Computer Vision: **Recurrent Neural Networks (RNN)**

What are RNNs?

- ▶ Neural networks designed for sequence data.
- ▶ Maintain a hidden state h_t that carries information from previous time steps.
- ▶ Recurrent architecture: $h_t = f(h_{t-1}, x_t)$

Use Cases:

- ▶ Language modeling
- ▶ Time series prediction
- ▶ Sequence classification




$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state / old state input vector at some time step

some function with parameters W


1

Recurrent Neural Networks (RNN) (cont.)




Recurrent Neural Networks (RNN) (cont.)


one to one




one to many




many to one



many to many



many to many



RNN Advantages:

- ▶ Can process any length input
- ▶ Computation for step t can (in theory) use information from many steps back
- ▶ Model size doesn't increase for longer input
- ▶ Same weights applied on every timestep, so there is symmetry in how inputs are processed.

RNN Disadvantages:

- ▶ Recurrent computation is slow
- ▶ In practice, difficult to access information from many steps back

Limitations of RNNs:


- ▶ Vanishing/exploding gradients in long sequences.
- ▶ Sequential computation — slow to train.
- ▶ Difficulty in capturing long-term dependencies.
- ▶ Cannot parallelize easily due to recursive nature.

Advanced Computer Vision: **Image Captioning**

- A computer Vision task in which we generate captions for images



A cat sitting on a suitcase on the floor



A cat is sitting on a tree branch



A dog is running in the grass with a frisbee



A white teddy bear sitting in the grass



Two people walking on the beach with surfboards



A tennis player in action on the court




Two giraffes standing in a grassy field



A man riding a dirt bike on a dirt track

- ▶ The task is to generate a natural language description of the content of an image
- ▶ It combines techniques from computer vision and natural language processing
- ▶ The process typically involves:
 - Extracting features from the image using a convolutional neural network (CNN)
 - Using these features to generate a sequence of words that describe the image, often with a recurrent neural network (RNN) or transformer model

Recurrent Neural Network



Convolutional Neural Network

Advanced Computer Vision: **Sequence-to-Sequence with RNNs**

Input: Sequence $x_1, \dots, x_T$

Output: Sequence $y_1, \dots, y_T$


Decoder: $s_t = g_U(y_{t-1}, s_{t-1}, c)$

Encoder: $h_t = f_W(x_t, h_{t-1})$

From final hidden state predict:

Initial decoder state s_0

Context vector c (often $c=h_T$)



Sequence-to-Sequence with RNNs (cont.)


Input: Sequence $x_1, \dots, x_T$

Output: Sequence $y_1, \dots, y_T$

Decoder: $s_t = g_U(y_{t-1}, s_{t-1}, c)$

Encoder: $h_t = f_W(x_t, h_{t-1})$

From final hidden state predict:
Initial decoder state s_0
Context vector c (often $c=h_T$)



Problem: Input sequence bottlenecked through fixed-sized vector. What if $T=1000$?

Sutskever et al, "Sequence to sequence learning with neural networks" NeurIPS 2014

Sequence-to-Sequence with RNNs (cont.)


Input: Sequence $x_1, \dots, x_T$

Output: Sequence $y_1, \dots, y_T$

Decoder: $s_t = g_U(y_{t-1}, s_{t-1}, c)$

Encoder: $h_t = f_W(x_t, h_{t-1})$

From final hidden state predict:
Initial decoder state s_0
Context vector c (often $c=h_T$)



Problem: Input sequence bottlenecked through fixed-sized vector. What if $T=1000$?

Idea: use new context vector at each step of decoder!

Sutskever et al, "Sequence to sequence learning with neural networks" NeurIPS 2014

- ▶ **Key idea:** “Don’t process everything — just focus on the most relevant parts.”
- ▶ Attention helps models decide where to look in a sequence.

Equation (simplified):


$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

- ▶ Q = Queries
- ▶ K = Keys
- ▶ V = Values

Input: Sequence $x_1, \dots, x_T$


Output: Sequence $y_1, \dots, y_T$


Encoder: $h_t = f_W(x_t, h_{t-1})$ From final hidden state:
Initial decoder state s_0



Compute (scalar) **alignment scores**

$$e_{t,i} = f_{\text{att}}(s_{t-1}, h_i) \quad (f_{\text{att}} \text{ is an MLP})$$






Compute (scalar) **alignment scores**

$$e_{t,i} = f_{\text{att}}(s_{t-1}, h_i) \quad (f_{\text{att}} \text{ is an MLP})$$

Normalize alignment scores
to get **attention weights**

$$0 < a_{t,i} < 1 \quad \sum_i a_{t,i} = 1$$

Sequence-to-Sequence with RNNs and Attention (cont.)




Compute (scalar) **alignment scores**
 $e_{t,i} = f_{\text{att}}(s_{t-1}, h_i)$ (f_{att} is an MLP)

Normalize alignment scores
to get **attention weights**
 $0 < a_{t,i} < 1 \quad \sum_i a_{t,i} = 1$

Compute context vector as
linear combination of hidden
states
 $c_t = \sum_i a_{t,i} h_i$

Sequence-to-Sequence with RNNs and Attention (cont.)



From final hidden state:
Initial decoder state s_0

Intuition: Context vector attends to the relevant part of the input sequence
"estamos" = "we are"
so maybe $a_{11}=a_{12}=0.45$,
 $a_{13}=a_{14}=0.05$

Compute (scalar) **alignment scores**

$$e_{t,i} = f_{\text{att}}(s_{t-1}, h_i) \quad (f_{\text{att}} \text{ is an MLP})$$

Normalize alignment scores to get **attention weights**


$$0 < a_{t,i} < 1 \quad \sum_i a_{t,i} = 1$$

Compute context vector as linear combination of hidden states

$$c_t = \sum_i a_{t,i} h_i$$


Use context vector in decoder: $s_t = g_U(y_{t-1}, s_{t-1}, c_t)$

estamos



[START]

Sequence-to-Sequence with RNNs and Attention (cont.)



From final hidden state:
Initial decoder state s_0

Intuition: Context vector attends to the relevant part of the input sequence
"estamos" = "we are"
so maybe $a_{11}=a_{12}=0.45$,
 $a_{13}=a_{14}=0.05$

Compute (scalar) **alignment scores**
 $e_{t,i} = f_{\text{att}}(s_{t-1}, h_i)$ (f_{att} is an MLP)

Normalize alignment scores to get **attention weights**
 $0 < a_{t,i} < 1 \quad \sum_i a_{t,i} = 1$


Compute context vector as linear combination of hidden states

$$c_t = \sum_i a_{t,i} h_i$$

Use context vector in decoder: $s_t = g_U(y_{t-1}, s_{t-1}, c_t)$


This is all differentiable! No supervision on attention weights – backprop through everything

Sequence-to-Sequence with RNNs and Attention (cont.)



Sequence-to-Sequence with RNNs and Attention (cont.)

- Use a different context vector at each timestep of the decoder.
- The input sequence is not bottlenecked through a single vector.
- At each timestep of the decoder, the context vector "looks at" different parts of the input sequence.



Example: English to French translation

Input: “The agreement on the European Economic Area was signed in August 1992.”


Output: “L’accord sur la zone économique européenne a été signé en août 1992.”

Diagonal attention means words correspond in order

Attention figures out different word orders


Diagonal attention means words correspond in order

Visualize attention weights $a_{t,i}$




Bahdanau et al, "Neural machine translation by jointly learning to align and translate", ICLR 2015

Advanced Computer Vision: **Image Captioning with RNNs and Attention**




Use a CNN to compute a
grid of features for an image

$$e_{t,i,j} = f_{\text{att}}(s_{t-1}, h_{i,j})$$



Use a CNN to compute a
grid of features for an image

Image Captioning with RNNs and Attention



Use a CNN to compute a grid of features for an image

Image Captioning with RNNs and Attention






Image Captioning with RNNs and Attention



$$e_{t,i,j} = f_{\text{att}}(s_{t-1}, h_{i,j})$$
$$a_{t,:} = \text{softmax}(e_{t,:})$$
$$c_t = \sum_{i,j} a_{t,i,j} h_{i,j}$$



$h_{1,1}$	$h_{1,2}$	$h_{1,3}$
$h_{2,1}$	$h_{2,2}$	$h_{2,3}$
$h_{3,1}$	$h_{3,2}$	$h_{3,3}$

Use a CNN to compute a grid of features for an image




Image Captioning with RNNs and Attention





Image Captioning with RNNs and Attention



$$e_{t,i,j} = f_{\text{att}}(s_{t-1}, h_{i,j})$$
$$a_{t,:} = \text{softmax}(e_{t,:})$$
$$c_t = \sum_{i,j} a_{t,i,j} h_{i,j}$$

Each timestep of decoder uses a different context vector that looks at different parts of the input image




Image Captioning with RNNs and Attention




A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A group of people sitting on a boat in the water.




A giraffe standing in a forest with trees in the background.



Advanced Computer Vision: **General Attention Layer**

Attention We Just Saw in Image Captioning



Outputs:

context vector: $\mathbf{c}$ (shape: D)

Operations:

Alignment: $e_{i,j} = f_{\text{att}}(\mathbf{h}, \mathbf{z}_{i,j})$


Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$

Output: $\mathbf{c} = \sum_{i,j} a_{i,j} \mathbf{z}_{i,j}$

Inputs:

Features: $\mathbf{z}$ (shape: $H \times W \times D$)

Query: $\mathbf{h}$ (shape: D)



Outputs:

context vector: $\mathbf{c}$ (shape: D)

Operations:

Alignment: $e_i = f_{\text{att}}(h, x_i)$

Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$


Output: $\mathbf{c} = \sum_i a_i x_i$

- ▶ The attention operation is permutation invariant.
- ▶ It does not depend on the ordering of the features.
- ▶ Stretch $H \times W = N$ into N vectors.

Inputs:

Input vectors: $\mathbf{x}$ (shape: $N \times D$)

Query: $\mathbf{h}$ (shape: D)



Outputs:

context vector: c (shape: D)

Operations:

Alignment: $e_i = h \cdot x_i / \sqrt{D}$

Attention: $a = \text{softmax}(e)$

Output: $c = \sum_i a_i x_i$


Inputs:

Input vectors: x (shape: $N \times D$)

Query: h (shape: D)

Change $f_{att}(\cdot)$ to a scaled simple dot product:

- ▶ Larger dimensions mean more terms in the dot product sum.
- ▶ Therefore, the variance of the logits is higher. Large-magnitude vectors will produce much higher logits.
- ▶ As a result, the post-softmax distribution has lower entropy, assuming logits are IID.
- ▶ Ultimately, these large-magnitude vectors will cause softmax to peak and assign very little weight to all others.
- ▶ Divide by $\sqrt{D}$ to reduce the effect of large-magnitude vectors.



Outputs:

context vectors: $\mathbf{y}$ (shape: D)

Operations:

Alignment: $e_{i,j} = q_j \cdot x_i / \sqrt{D}$

Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$


Output: $y_j = \sum_i a_{i,j} x_i$

- We can have multiple query vectors.
- Each query creates a new output context vector.

Inputs:

Input vectors: $\mathbf{x}$ (shape: $N \times D$)

Queries: $\mathbf{q}$ (shape: $M \times D$)




Outputs:
context vectors: $\mathbf{y}$ (shape: D)

Operations:
Alignment: $e_{i,j} = q_j \cdot x_i / \sqrt{D}$
Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$
Output: $y_j = \sum_i a_{i,j} x_i$

- Notice that the input vectors are used for both the alignment and the attention calculations.

Inputs:
Input vectors: $\mathbf{x}$ (shape: N x D)
Queries: $\mathbf{q}$ (shape: M x D)



Operations:

Key vectors: $\mathbf{k} = \mathbf{x}\mathbf{W}_k$


Value vectors: $\mathbf{v} = \mathbf{x}\mathbf{W}_v$

Inputs:

Input vectors: $\mathbf{x}$ (shape: $N \times D$)

Queries: $\mathbf{q}$ (shape: $M \times D_k$)

- ▶ Notice that the input vectors are used for both the alignment and the attention calculations.
- ▶ We can add more expressivity to the layer by adding a different fully connected (FC) layer before each of the two steps.



Outputs:

context vectors: y (shape: D_v)

Operations:

Key vectors: $k = xW_k$

Value vectors: $v = xW_v$

Alignment: $e_{i,j} = q_j \cdot k_i / \sqrt{D}$

Attention: $a = \text{softmax}(e)$

Output: $y_j = \sum_i a_{i,j} v_i$

Inputs:


Input vectors: x (shape: $N \times D$)

Queries: q (shape: $M \times D_k$)

- Notice that the input vectors are used for both the alignment and the attention calculations.
- We can add more expressivity to the layer by adding a different fully connected (FC) layer before each of the two steps.
- The input and output dimensions can now change depending on the key and value FC layers.

Advanced Computer Vision: **Self Attention Layer**

- ▶ **Self-Attention:** Computes attention within a single sequence.
- ▶ **How it works:**
 - Each token attends to every other token.
 - Captures global context, not just local dependencies.



Operations:

Key vectors: $\mathbf{k} = \mathbf{x}\mathbf{W}_k$

Value vectors: $\mathbf{v} = \mathbf{x}\mathbf{W}_v$

Query vectors: $\mathbf{q} = \mathbf{x}\mathbf{W}_q$

Alignment: $e_{i,j} = \mathbf{q}_i \cdot \mathbf{k}_j / \sqrt{D}$

Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$


Output: $\mathbf{y}_j = \sum_i a_{i,j} \mathbf{v}_i$

Inputs:

Input vectors: $\mathbf{x}$ (shape: $N \times D$)

Queries: $\mathbf{q}$ (shape: $M \times D_q$)

- Recall that the query vector is a function of the input vectors.
- We can calculate the query vectors from the input vectors, thus defining a "self-attention" layer.
- There are no input query vectors anymore.
- Instead, query vectors are calculated using a fully connected (FC) layer.




Outputs:
context vectors: $\mathbf{y}$ (shape: D_v)

Operations:
Key vectors: $\mathbf{k} = \mathbf{x}\mathbf{W}_k$
Value vectors: $\mathbf{v} = \mathbf{x}\mathbf{W}_v$
Query vectors: $\mathbf{q} = \mathbf{x}\mathbf{W}_q$
Alignment: $e_{i,j} = q_{i,j} \cdot k_{i,j} / \sqrt{D}$
Attention: $\mathbf{a} = \text{softmax}(\mathbf{e})$
Output: $y_j = \sum_i a_{i,j} v_i$

Inputs:
Input vectors: $\mathbf{x}$ (shape: $N \times D$)

- Recall that the query vector is a function of the input vectors.
- We can calculate the query vectors from the input vectors, thus defining a "self-attention" layer.
- There are no input query vectors anymore.
- Instead, query vectors are calculated using a fully connected (FC) layer.


Self Attention Layer






Outputs:
context vectors: y (shape: D_v)

Operations:
Key vectors: $k = xW_k$
Value vectors: $v = xW_v$
Query vectors: $q = xW_q$
Alignment: $e_{i,j} = q_i \cdot k_j / \sqrt{D}$
Attention: $a = \text{softmax}(e)$
Output: $y_j = \sum_i a_{i,j} v_i$

Inputs:
Input vectors: x (shape: $N \times D$)




Advanced Computer Vision: **Permutation Invariance**




- ▶ Self-Attention Layer is Permutation equivariant
- ▶ It doesn't care about the orders of the inputs!
- ▶ **Problem:** How can we encode ordered sequences like language or spatially ordered image features?

- ▶ Concatenate/add special positional encoding p_j to each input vector x_j
- ▶ We use a function $\text{pos}: N \rightarrow \mathbb{R}^D$ to process the position j of the vector into a d -dimensional vector
- ▶ So, $p_j = \text{pos}(j)$



Masked self-attention layer



Outputs:

context vectors: y (shape: D_v)

Operations:

Key vectors: $k = xW_k$

Value vectors: $v = xW_v$

Query vectors: $q = xW_q$

Alignment: $e_{i,j} = q_i \cdot k_j / \sqrt{D}$

Attention: $a = \text{softmax}(e)$


Output: $y_j = \sum_i a_{i,j} v_i$

Inputs:



Input vectors: x (shape: $N \times D$)

- ▶ Prevent vectors from looking at future vectors.
- ▶ Manually set alignment scores to $-\infty$

- Multiple self-attention heads in parallel



General attention versus self-attention



Example: CNN with Self-Attention

Input Image




[Cat image](#) is free to use under the [Pixabay License](#)




Features:
 $C \times H \times W$


Example: CNN with Self-Attention




Example: CNN with Self-Attention




Example: CNN with Self-Attention



Example: CNN with Self-Attention



Example: CNN with Self-Attention




Advanced Computer Vision: **Transformer Overview**


Attention is all you need

Vaswani et al, NeurIPS 2017

- ▶ **Input:** Set of vectors $\mathbf{x}$
- ▶ **Output:** Set of vectors $\mathbf{y}$
- ▶ Self-attention is the only interaction between vectors!
- ▶ Layer normalization and MLP operate independently on each vector.
- ▶ Highly scalable and highly parallelizable.



- ▶ A Transformer consists of a sequence of transformer blocks.
- ▶ Vaswani et al. used 12 blocks, $D_Q = 512$, and 6 attention heads.



Advanced Computer Vision: **Vision Transformers**


Key idea:

- ▶ Split image into fixed-size patches (e.g., 16×16)
- ▶ Flatten and project patches into embeddings
- ▶ Add positional encodings to embeddings
- ▶ Apply standard Transformer encoder

Architecture:

- ▶ Patch embedding
- ▶ Class token
- ▶ Transformer blocks
- ▶ MLP Head

Source: Dosovitskiy et al., "An Image is Worth 16×16 Words", 2020




N input patches, each
of shape 3x16x16



Linear projection to
D-dimensional vector


N input patches, each
of shape $3 \times 16 \times 16$





Add positional
embedding: learned D-
dim vector per position


Linear projection to
D-dimensional vector

N input patches, each
of shape 3x16x16










⁰Dosovitskiy et al, "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale", ICLR 2021

- ▶ **Patch Embedding:** Achieved via a linear projection of flattened patches ($P = 16$ typical)
- ▶ **Differences from CNN stems:**
 - Large stride causes optimization instability
 - Mitigated via convolutional preprocessing (e.g., small conv stem)




- ▶ ViT learns the grid-like structure of the image patches via its position embeddings.
- ▶ The lower layers contain both global and local features, while the higher layers contain only global features.

Advanced Computer Vision: **Improving ViT:** **Distillation**

Step 1: Train a teacher CNN on ImageNet



$$P(\text{cat}) = 0.9$$
$$P(\text{dog}) = 0.1$$




Step 2: Train a student ViT to match ImageNet predictions from the teacher CNN (and match GT labels)





$$P(\text{cat}) = 0.1$$
$$P(\text{dog}) = 0.9$$

KL Divergence Loss



$$P(\text{cat}) = 0.2$$
$$P(\text{dog}) = 0.8$$





⁰Touvron et al, "Training data-efficient image transformers distillation through attention", ICML 2021

Advanced Computer Vision: **CNN** vs **ViT**

CNN



1. Maintain 2D structure logic



2. Shift equivariant



3. Consider only local correlations



4. Hierarchically growing field of view



5. Hierarchically progressing complexity




6. Reasonable amount of params



7. Global representation

ViT






In most CNNs (including ResNets), **decrease** resolution and **increase** channels as you go deeper in the network (Hierarchical architecture)

Useful since objects in images can occur at various scales

In a ViT, all blocks have same resolution and number of channels (Isotropic architecture)


Can we build a **hierarchical** ViT model?




Feature	CNNs	Vision Transformers
Locality	Strong inductive bias	Weak (learns from data)
Parameter sharing	Yes	Yes
Global context	Limited	Native support
Training data	Efficient on small data	Needs large data or pretraining
Flexibility	Less modular	More modular


Advanced Computer Vision: **Hierarchical ViT: Swin Transformer**

Hierarchical ViT: Swin Transformer






⁰Liu et al, "Swin Transformer: Hierarchical Vision Transformer using Shifted Windows", CVPR 2021



- ▶ Swin Transformer introduces windowed self-attention with shifting windows, merging tokens akin to pooling.¹
- ▶ Builds a pyramid representation, enabling linear compute complexity and high performance on detection/segmentation (e.g. 58.7 COCO box AP).
- ▶ Hierarchical design allows for multi-scale feature extraction, similar to CNNs.


Hierarchical ViT: Swin Transformer (cont.)



¹<https://arxiv.org/abs/2103.14030>


²Liu et al, "Swin Transformer: Hierarchical Vision Transformer using Shifted Windows", CVPR 2021

MViT




Fan et al, "Multiscale Vision Transformers", ICCV 2021

Swin-V2



Liu et al, "Swin Transformer V2: Scaling up Capacity and Resolution", CVPR 2022


Improved MViT



Li et al, "Improved Multiscale Vision Transformers for Classification and Detection", arXiv 2021


Advanced Computer Vision: **Other Hierarchical Vision Transformers**

MViT




Fan et al, "Multiscale Vision Transformers", ICCV 2021

Swin-V2



Liu et al, "Swin Transformer V2: Scaling up Capacity and Resolution", CVPR 2022

Improved MViT





Li et al, "Improved Multiscale Vision Transformers for Classification and Detection", arXiv 2021

- ▶ Simple object detection pipeline: directly outputs a set of boxes from a Transformer
- ▶ No anchors, no regression of box transforms
- ▶ Matches predicted boxes to ground truth boxes with bipartite matching; trains to regress box coordinates

²Carion et al., “End-to-End Object Detection with Transformers,” ECCV 2020



Object Detection with Transformers: DETR



²Carion et al., "End-to-End Object Detection with Transformers," ECCV 2020

- ▶ DETR (DEtection TRansformer, Carion et al., 2020) introduces a simple and unified object detection pipeline using Transformers.
- ▶ Transformers replace region proposal networks (RPNs) with attention-based object queries for end-to-end detection.
- ▶ DETR directly predicts a fixed set of bounding boxes and class labels, eliminating the need for anchors, hand-crafted box proposals, or non-maximum suppression (NMS).
- ▶ The approach simplifies the pipeline and enables global reasoning, though it may struggle with small objects and require higher compute.
- ▶ DETR uses bipartite matching (Hungarian algorithm) to uniquely match predicted boxes to ground truth boxes, and is trained end-to-end to regress box coordinates and classify objects.

Object Detection with Transformers: DETR (cont.)



³Carion et al., "End-to-End Object Detection with Transformers", ECCV 2020

<https://arxiv.org/abs/2005.12872>

⁴<https://docs.google.com/presentation/d/1X0BDhpJ0a3IOf29DU8vAB4WJruyLugIX/edit?slide=id.pi#>

- ▶ ConvNeXt is a modern CNN architecture that incorporates insights from transformer models.
- ▶ It uses a hierarchical design similar to Swin Transformer, with a focus on simplicity and efficiency.
- ▶ Key features include:
 - Depthwise separable convolutions
 - Layer normalization
 - Global average pooling
- ▶ Achieves competitive performance on image classification tasks while maintaining the strengths of CNNs.
- ▶ ConvNeXt demonstrates that CNNs can still be effective with modern design principles, even in the era of transformers.




Figure 4. **Block designs** for a ResNet, a Swin Transformer, and a ConvNeXt. Swin Transformer's block is more sophisticated due to the presence of multiple specialized modules and two residual connections. For simplicity, we note the linear layers in Transformer MLP blocks also as " 1×1 convs" since they are equivalent.

5

⁵Liu et al., "A ConvNet for the 2020s", CVPR 2022 <https://arxiv.org/abs/2201.03545>

Advanced Computer Vision: **Limitations and Future Directions**

- ▶ **Data hungry:** Performance drops significantly on small datasets.
- ▶ **Computationally expensive:** Quadratic complexity of self-attention.
- ▶ **Lacks built-in locality bias:** Unlike CNNs; some variants (e.g., Swin Transformer) address this.
- ▶ **Requires large-scale pretraining:** Needs datasets like JFT-300M, ImageNet-22K for strong performance.

- ▶ **Efficient transformers:** Linformer, Performer, Longformer.
- ▶ **Hybrid models:** Combining CNNs and Transformers (e.g., CoaT, CvT).
- ▶ **Low-resource ViTs:** Data-efficient training (e.g., DeiT).
- ▶ **Vision-language models:** CLIP, Flamingo, GPT-4V.
- ▶ **Hardware acceleration:** Edge ViTs, quantized ViTs for deployment.

Advanced Computer Vision: **References**

- [1] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is All You Need. *Advances in Neural Information Processing Systems*, 30. <https://arxiv.org/abs/1706.03762>
- [2] Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., ... & Houlsby, N. (2020). An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. *International Conference on Learning Representations*.
<https://arxiv.org/abs/2010.11929>
- [3] Touvron, H., Cord, M., Douze, M., Massa, F., Sablayrolles, A., & Jégou, H. (2021). Training data-efficient image transformers distillation through attention. *International Conference on Machine Learning*. <https://arxiv.org/abs/2012.12877>
- [4] Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., Zhang, Z., ... & Guo, B. (2021). Swin Transformer: Hierarchical Vision Transformer using Shifted Windows. *CVPR 2021*.
<https://arxiv.org/abs/2103.14030>

- [5] Chen, J., Chen, J., Chao, H., & Yang, M. (2021). Pre-Trained Image Processing Transformer. *CVPR 2021*. <https://arxiv.org/abs/2012.00364>
- [6] Lu, Y., Dosovitskiy, A., Houlsby, N., & Beyer, L. (2021). Vision Transformer: A Survey. *arXiv preprint arXiv:2111.06091*. <https://arxiv.org/abs/2111.06091>
- [7] Xie, E., Wang, W., Yu, Z., Anandkumar, A., Alvarez, J. M., & Luo, P. (2021). SegFormer: Simple and Efficient Design for Semantic Segmentation with Transformers. *NeurIPS 2021*. <https://arxiv.org/abs/2105.15203>
- [8] Zheng, S., Lu, J., Zhao, H., Zhu, X., Luo, Z., Wang, Y., ... & Fu, Y. (2021). Rethinking Semantic Segmentation from a Sequence-to-Sequence Perspective with Transformers (SETR). *CVPR 2021*. <https://arxiv.org/abs/2012.15840>
- [9] Strudel, R., Garcia, R., Laptev, I., & Schmid, C. (2021). Segmenter: Transformer for Semantic Segmentation. *ICCV 2021*. <https://arxiv.org/abs/2105.05633>

- [10] Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021). Learning Transferable Visual Models From Natural Language Supervision (CLIP). *ICML 2021*. <https://arxiv.org/abs/2103.00020>
- [11] Alayrac, J.-B., Donahue, J., Luc, P., Miech, A., Barr, I., Hasson, Y., ... & Simonyan, K. (2022). Flamingo: a Visual Language Model for Few-Shot Learning. *arXiv preprint arXiv:2204.14198*. <https://arxiv.org/abs/2204.14198>
- [12] Oquab, M., Darcet, T., Moutakanni, T., Vo, D. H., Szafraniec, M., Khalidov, V., ... & Jegou, H. (2023). DINOv2: Learning Robust Visual Features without Supervision. *arXiv preprint arXiv:2304.07193*. <https://arxiv.org/abs/2304.07193>
- [13] Fei-Fei Li, Y., Li, Y., & Gao, R. Stanford CS231n: Deep Learning for Computer Vision. <http://cs231n.stanford.edu/>

- [14] Shocher, A., Bagon, S., Galun, M., & Dekel, T. WAIC DL4CV: Deep Learning for Computer Vision: Fundamentals and Applications.
<https://www.cs.tau.ac.il/~assafshocher/WAIC-DL4CV/>
- [15] Johnson, J. UMich EECS 498.008/598.008: Deep Learning for Computer Vision.
<https://web.eecs.umich.edu/~justincj/teaching/eecs498/WI2023/>
- [16] Hugging Face. Vision Transformer Playground.
<https://huggingface.co/spaces/akhaliq/ViT>
- [17] Stanford CS231n. Lecture Slides: Convolutional Neural Networks for Visual Recognition.
<http://cs231n.stanford.edu/slides/>
- [18] Justin Johnson. UMich EECS 498.008/598.008: Deep Learning for Computer Vision Slides. <https://web.eecs.umich.edu/~justincj/teaching/eecs498/slides.html>

[19] PyTorch. PyTorch Tutorials: Vision.

https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html