

Unsupervised Learning

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

August 9, 2026

By the end of this session, you will be able to:

- ▶ Differentiate between Supervised and Unsupervised Learning paradigms.
- ▶ Understand K-Means clustering, and how it could be applied to discover hidden groups in data.
- ▶ Analyze high-dimensional data using PCA and t-SNE for visualization and compression.
- ▶ Understand how Autoencoders utilize neural networks for tasks like denoising, generation, and anomaly detection.

1. Learning Outcomes
2. Unsupervised Learning
3. Introduction
4. Clustering: K-Means
5. Clustering: K-Means
6. Dimensionality Reduction: PCA
7. Dimensionality Reduction: PCA
8. Dimensionality Reduction: t-SNE
9. t-SNE
10. Autoencoders
11. Autoencoders
12. References
13. References

What is Unsupervised Learning?

So far: We've been doing **Supervised Learning**, learning from labeled data (x, y) .

Supervised Learning

Data: (x, y) pairs

- ▶ x : input (features)
- ▶ y : label (target)

Goal: Learn $f(x) \approx y$

Examples:

- ▶ Image classification
- ▶ House price prediction

Unsupervised Learning

Data: Only x (no labels!)

- ▶ Just raw data
- ▶ No target to predict

Goal: Find hidden patterns or structure

Examples:

- ▶ Customer segmentation
- ▶ Data compression

Key difference: No labels $\rightarrow$ algorithm must discover structure on its own!

Labels are expensive!

- ▶ Labeling data requires human effort, time, and expertise
- ▶ Most real-world data is unlabeled
- ▶ Sometimes we don't even know what labels to look for

The Big Idea

Can we learn useful things from data **without** labels?

Yes! Unsupervised learning finds hidden patterns, structure, and representations.

What can we do without labels?

1. Clustering

Group similar data points together

Example: Customer segmentation, document group

3. Anomaly Detection

Identify unusual or outlier data points

Example: Fraud detection, quality control

2. Dimensionality Reduction

Compress high-dimensional data to lower dimensions

Example: Visualization, feature extraction


4. Generation

Learn to create new, realistic data samples

Example: Image synthesis, data augmentation

We'll explore four powerful algorithms that tackle these tasks.

Main Idea: Represent each cluster by its center (mean)



- ▶ Each cluster has a **center** (centroid) μ_k (marked with $\times$)
- ▶ Points belong to their **nearest** center
- ▶ Goal: Find centers that minimize total distance to assigned points

K-Means Algorithm

K-Means Algorithm

Input: Data $\{x_1, x_2, \dots, x_n\}$ and number of clusters K

Initialize: Randomly place K cluster centers $\mu_1, \mu_2, \dots, \mu_K$

Repeat until convergence:

1. **Assignment Step:** For each point x_i , assign to nearest center

$$c_i = \arg \min_{k=1, \dots, K} \|x_i - \mu_k\|^2$$

(find which center is closest)

2. **Update Step:** Move each center to mean of assigned points

$$\mu_k = \frac{1}{|C_k|} \sum_{i \in C_k} x_i$$

(average all points in cluster k)

Stop when: Assignments don't change OR max iterations reached

K-Means: Step-by-Step Example



Repeat steps 2-3 until centers stop moving (typically 10-100 iterations)

Problem: How many clusters should we use?

Elbow Method

Plot "within-cluster sum of squares" (WCSS) vs K

$$WCSS = \sum_{k=1}^K \sum_{i \in C_k} \|x_i - \mu_k\|^2$$




Look for the "elbow" where adding more clusters doesn't help much (in the above case it is 4)

Use case: Group customers for targeted marketing

Example: E-commerce Data

- Features: frequency, spend, recency
- Run K-Means with $K = 4$
- Interpret clusters



Discovered Segments:

Inactive

Low freq,
low spend

Action:
Re-engage

Bargain

High freq,
low spend

Action:
Discounts

Big Spenders

Low freq,
high spend

Action:
Premium Service


VIP

High freq,
high spend

Action:
Loyalty Program

Application 2: Document Clustering

Use case: Automatically organize news articles by topic.




Process

- ▶ Convert text into numbers (Text Features).
- ▶ Similar stories appear close together in space.
- ▶ K-Means automatically groups them into topics.

Example: Google News grouping stories from different sources.


K-Means assumes spherical, equal-sized clusters

Problem 1: Non-spherical



Fails on concentric circles!

Problem 2: Different sizes



Struggles with size differences!

Other Limitations

- ▶ Sensitive to outliers
- ▶ Must specify K in advance
- ▶ Random initialization matters
- ▶ Bad results for high dimensional data

Alternatives: DBSCAN (density-based), Gaussian Mixture Models, Hierarchical clustering


Real-world data is high-dimensional:

Examples:

- ▶ Images: $1000 \times 1000 = 1\text{M}$ pixels
- ▶ Text: 10,000+ words
- ▶ Genomics: 20,000+ genes
- ▶ Sensors: 100+ measurements

Problems:

- ▶ Hard to visualize
- ▶ Slow to compute
- ▶ Lots of redundancy
- ▶ Noisy measurements



Solution: Find a lower-dimensional representation while keeping important information (Dimensionality Reduction)


Principal Component Analysis (PCA) is a dimensionality reduction algorithm that finds the most important directions in your data.

Core Idea

Instead of using original features, create **new features** (principal components) that:


- ▶ Capture maximum variance
- ▶ Are uncorrelated with each other
- ▶ Are ranked by importance

Then keep only the top few components!



PCA Intuition: Finding the Best Directions

Goal: Find directions where data varies the most.



- ▶ PC1: Direction with most spread
- ▶ PC2: Perpendicular to PC1, next most spread
- ▶ PC3, PC4, ... : Each perpendicular to all previous ones

Step 1: Prepare the Data

1. Center the Data

Shift the data so the center is at the origin (0,0).

$$X_{\text{centered}} = X - \bar{X}$$

(Subtract the mean from every feature)

2. Compute Covariance Matrix

Calculate how features vary with respect to each other.

$$C = \frac{1}{n} X_{\text{centered}}^{\top} X_{\text{centered}}$$

(Finds correlations between features)

Step 2: Find Directions & Project

3. Find Principal Components

Calculate **Eigenvectors** and **Eigenvalues** of the covariance matrix C .

- ▶ **Eigenvectors:** The directions of the new axes.
- ▶ **Eigenvalues:** The amount of variance (information) along those axes.

4. Select & Project

Pick the top k eigenvectors with the largest eigenvalues and transform the data.


$$Z = X_{\text{centered}} \cdot V_k$$

(Reduces data from d dimensions $\rightarrow k$ dimensions)

Key question: How many components are enough?

Cumulative Variance Explained

(The most common method for selecting K)




Rule of Thumb: Choose the smallest K that preserves **80–95%** of the total variance.

The Data: MNIST Digits, 70,000 handwritten digits (0–9).

The Challenge:


- ▶ Each image is 28×28 pixels.
- ▶ Total features: $28 \times 28 = \mathbf{784}$ dimensions!
- ▶ We cannot "see" in 784 dimensions.

Question: Can we compress this to just **2D** to see how the digits relate?




Applying PCA: We project the 784 dimensions down to the top 2 Principal Components.

Application 1: Data Visualization (cont.)




Is this the best visualization?

- ▶ It captures some structure, but points overlap heavily.
- ▶ **Reason:** PCA is **linear**, it cannot unfold complex, curved structures.
- ▶ PCA is mostly used for compression. We will cover a better visualization algorithm shortly.



Use case: Store data more efficiently



More components = better reconstruction, less compression

10 PCs
99% compression
(blurry)

50 PCs
93% compression
(good)

200 PCs
75% compression
(near perfect)

PCA Limitations:

- ▶ **Linear only:** Finds linear combinations of features
- ▶ **Sensitive to scale:** Must standardize data first
- ▶ **Not Interpretable:** PCs are combinations of all features
- ▶ **Outliers:** Can distort principal components

Let's have a look at a better algorithm for high dimensional data visualization...

PCA is great... but it's linear. Many datasets have **non-linear** structure.


Problem

- ▶ High-dimensional data can be hard to visualize
- ▶ PCA may **mix** classes if the separation is non-linear
- ▶ We mainly want a **good 2D plot** for exploration

Main Idea


t-SNE tries to build a 2D map where: neighbors in high-D remain neighbors in 2D.

High-D




t-SNE
→

2D map



Neighborhood stays together

Example: Visualizing MNIST via t-SNE




This is much better than PCA! We can see how similar digits are grouped together, while different ones are far from each other.

We have covered classical algorithms like K-Means and PCA.

Now, let's explore a powerful Neural Network architecture designed for unsupervised learning...

Autoencoder: a type of neural network that learn to compress data into a compact form and then reconstruct it to match the original input.




The trick: Force information through a narrow **bottleneck** z .

This forces the network to learn compressed, meaningful representations!

Why Does the Bottleneck Work?

What if there's no bottleneck?

Without Bottleneck




Network can just copy:

$$h_i = x_i, \hat{x}_i = h_i$$

Learns nothing useful!


With Bottleneck



Must compress 4D $\rightarrow$ 2D
Forces learning of structure!

Learns meaningful features!


How to use an Autoencoder for compression:



The Process

1. **Train** the full network to reconstruct the input ($x \approx \hat{x}$).
2. **Discard** the Decoder (it's only needed for training).
3. **Use the Encoder** to map high-dimensional data x to low-dimensional code z .


How to use an Autoencoder for generation:



The Process

1. **Train** the full network on images to learn features.
2. **Discard** the Encoder.
3. **Sample** a random vector z (pick random numbers).
4. **Generate** a brand new image by running z through the Decoder.

Idea: Train the network to recover the original signal from corrupted input.




How it works


- ▶ Take clean image x .
- ▶ Corrupt it to get $\tilde{x}$.
- ▶ Force AE to map $\tilde{x} \rightarrow x$.

Concept: The network learns to reconstruct "Normal" data perfectly. When it sees an anomaly, it fails to reconstruct it.

1. Training / Normal Operation



2. Detection Phase



Real-World Use Cases

- ▶ **Credit Card Fraud:** Normal transactions reconstruct well, theft looks weird.
- ▶ **Manufacturing:** Detect defective parts on an assembly line.

- ▶ Aurélien Géron, *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow*
- ▶ Andrew Ng, *Machine Learning Specialization* (Coursera/DeepLearning.AI)

Slides contributed by Mohamed Eltayeb