

Linear Regression

Naeemullah Khan

naeemullah.khan@kaust.edu.sa




جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology


June 18, 2026

- ▶ Introduction to ML
- ▶ Linear Regression
- ▶ Optimization
- ▶ Linear Regression: Probabilistic Interpretation
- ▶ Bias-Variance Tradeoff
- ▶ Regularization

- **Machine Learning** is the science (and art) of programming computers so they can learn from data.



The traditional approach




The Machine Learning approach

- ▶ **Tabular Data** (e.g., spreadsheets, databases)
 - Note: Columns are called **Features**. Rows are called **Samples**.
- ▶ **Time-Series Data** (e.g., stock prices, weather forecasts, IoT sensor data)
- ▶ **Text Data** (Natural Language Processing, e.g., emails, social media posts, documents)
- ▶ **Images and Videos** (Computer Vision, e.g., medical imaging, surveillance, facial recognition)
- ▶ **Audio Data** (Speech Recognition, Music Processing, e.g., voice commands, podcasts, sound classification)

- ▶ **Supervised:** The algorithm learns from labeled data.
 - **Regression:** Predict continuous value (e.g. house prices).
 - **Classification:** Predict discrete value (e.g. spam/not-spam).
- ▶ **Unsupervised:** The algorithm works on unlabeled data. We are interested in things like:
 - **Clustering:** Grouping
 - **Dimensionality Reduction:** Reducing the Dimensions
 - **Anomaly Detection:** Detecting outliers
- ▶ **Reinforcement Learning:** involves learning to make decisions by interacting with an environment.
 - **Reward Signal:** The agent receives feedback in the form of rewards or penalties, guiding its learning.
 - **Policy:** A strategy the agent learns to decide actions based on the current state.


- **Value Function:** An estimate of the expected cumulative reward from a state or state-action pair.
- **Exploration vs Exploitation:** The agent balances exploring new actions to discover rewards and exploiting known actions to maximize them.
- Really popular in video games and robotics! (Also recently in LLMs, see **RLHF**)

Example: Spam Classification



An example of Supervised Learning: Spam Classification

Example: Clustering




An example of Unsupervised Learning: Clustering

- ▶ Most of the ML systems consist of three main components:
- ▶ **Hypothesis (Model):** The function that approximates the target.
 - E.g. Linear Regression, Logistic Regression, SVM, Decision Trees, NN, ...
- ▶ **Optimizer:** The mechanism for improving predictions of our model.
- ▶ **Loss Function:** The measure of how wrong the predictions are.

- ▶ How are they related to each other? 🤔

- ▶ We firstly define our task (classification/regression) then choose an appropriate model.
- ▶ We will use an optimization method to minimize the loss function.
- ▶ Reached a minima?
 - Model is making the least possible number of mistakes.
 - Model trained 🎉.

- ▶ Linear Regression is “still” one of the more widely used ML/DL Algorithms
- ▶ Easy to understand and implement
- ▶ Efficient to Solve
- ▶ We will use Linear Regression to understand the concepts of:
 - Data
 - Models
 - Loss
 - Optimization



Model (*Linear*): $Y = mX + b$

$$\theta = \{m, b\}$$

- ▶ Y: Response Variable
- ▶ X: Covariate / Independent variable / Regressors
- ▶ m: slope

- ▶ **Hypothesis:** $\hat{y}_i = mx_i + b$
- ▶ **Input:** data (x_i, y_i) , $i \in \{1, 2, \dots, N\}$
 - (e.g., house size x and price y)
- ▶ **Goal:** learn values of variable (m, b)


- Some clarification about the notation we will use for this course

$$X_i^{j,[k]}$$

- i is the index of the data, j is the feature number, and k is the power.


Solution Strategy for Solving the Problem

- ▶ There are countless possible lines.
- ▶ We want a line which is in some sense the “average line” that represents the data.
- ▶ Any ideas as to how we can do it?



Click the image to view the animated version.

- ▶ To find the "best line," we should minimize the distances between our line's predictions and all the data points.
- ▶ How to define that mathematically?



- For one sample, this can be represented mathematically by:

$$(y_i - \hat{y}_i) \quad (\text{Error})$$


- But this could result in negative value if $\hat{y} > y$. Let's square it to remove the negative sign:

$$(y_i - \hat{y}_i)^2 \quad (\text{Squared Error})$$


- But we have N samples, not only one. So, let's sum the errors and take the average:

$$\text{Loss (MSE)} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \quad (\text{Mean Squared Error})$$

Cost Function Example



Hypothesis



Loss Function

Notice: Lower is better.

$$J(m = 0) = 14$$

$$J(m = 1) = 0$$


$$J(m = 2) = 14$$

$$h(x) = mx \quad J(m) = \sum_{i=1}^3 (y_i - mx_i)^2$$


- ▶ There are two approaches to find the minima:
 - **Exact (Closed-form):** Directly calculates the solution mathematically by solving for $f'(x) = 0$.
Important Note: can be used only with a very limited number of algorithms.
 - **Approximation (Iterative approach):** Gradually improves the solution step by step. Done by optimizers (e.g., Gradient Descent, ADAM, ... etc).

How to find minima of a function (Review):

Closed-form:




Iterative:



► **Example:** $y = x^2$ (Solution: $x = 0$)

- Closed-form Final Result: $x = 0$
- Iterative Final Result: $x = 0.00001$ (close enough)



(Notice that $y = x$ for our 3 points.)

- ▶ Let's try to solve this using the closed-form here (Assume $\hat{y} = mx$):

$$J(m) = \sum_{i=1}^3 (y_i - mx_i)^2 \quad \Rightarrow \quad J(m) = \sum_{i=1}^3 (i - mi)^2$$

$$\frac{dJ(m)}{dm} = \frac{d}{dm} \sum_{i=1}^3 (i - mi)^2 \Rightarrow \frac{dJ(m)}{dm} = \sum_{i=1}^3 \frac{d}{dm} (i - mi)^2$$

$$\frac{dJ(m)}{dm} = \sum_{i=1}^3 -2i(i - mi) = -2 \sum_{i=1}^3 i^2 + 2m \sum_{i=1}^3 i^2 = 0 \Rightarrow m = 1$$

- ▶ Let's set up regression for a linear function in two variables:
- ▶ The hypothesis function is:

$$\hat{y}_i = mx_i + b$$

- ▶ Similar to the previous problem, our loss function is:

$$J = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- ▶ Let's calculate the partial derivatives of the loss function *w.r.t.* m , b

- We get the following expressions for the gradient of the cost function:

$$\frac{\partial J}{\partial m} = \frac{1}{N} \sum_{i=1}^N -2(y_i - \hat{y}_i)x_i$$

$$\frac{\partial J}{\partial b} = \frac{1}{N} \sum_{i=1}^N -2(y_i - \hat{y}_i)$$

- Simplifying the above expressions, we get:


$$\frac{\partial J}{\partial m} = \frac{-2}{N} \sum_{i=1}^N y_i x_i + \frac{2m}{N} \sum_{i=1}^N x_i^2 + \frac{2b}{N} \sum_{i=1}^N x_i$$

$$\frac{\partial J}{\partial b} = \frac{-2}{N} \sum_{i=1}^N y_i + \frac{2m}{N} \sum_{i=1}^N x_i + \frac{2b}{N} \sum_{i=1}^N 1$$

- Setting the gradient equal to 0 and solving for m and b , we get:

$$\begin{bmatrix} \frac{\sum_i x_i^2}{N} & \frac{\sum_i x_i}{N} \\ \frac{\sum_i x_i}{N} & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} = \begin{bmatrix} \frac{\sum_i x_i y_i}{N} \\ \frac{\sum_i y_i}{N} \end{bmatrix}$$

- ▶ What if y is a non-linear function of x ?
- ▶ Will this approach still work?



- We can transform features x_i :

$$x_i = (x_i^1, x_i^2, x_i^3, \dots, x_i^m)$$



- We will apply some non-linear transformation ϕ :

$$\phi : \mathbb{R}^m \rightarrow \mathbb{R}^M$$

- For example, Polynomial transformation:

$$\phi(x_i) = \left\{ 1, x_i^1, x_i^{1,[2]}, \dots, x_i^{1,[k]}, x_i^2, x_i^{2,[2]}, \dots, x_i^{2,[k]}, \dots, x_i^m, x_i^{m,[2]}, \dots, x_i^{m,[k]} \right\}$$

Transforming the Feature Space (Feature Engineering)



- **Example:** Assume you have:

$$x_i^1 : \text{Length} \quad x_i^2 : \text{Width}$$

- You can add $x_i^3 : \text{Area} = x_i^1 \cdot x_i^2$ to the dataset.

- **Other types:**

- Cosine, splines, radial basis functions, etc.
- Encoding (Label encoding, One-hot, ...)
- Domain-related features (e.g., financial measures)
- Time-related features (Day, month, year, ...)
- Group-level features (e.g., average income per household, total sales per region, median age per team). Often called “Aggregation features”.

- ▶ Let's get back to the gradients. . .

- ▶ Assume we have 100 variables instead of 2.
- ▶ Calculating gradients like this can quickly become tedious.
- ▶ **Notice:** Each term on either side of the expression can be written as a dot product of two vectors (maybe we can calculate it more efficiently)?
- ▶ Let's explore if we can do something better through **vectorization** (Writing equations as matrices).

- ▶ To truly appreciate the power of vectorization, let's make the problem a little more complex. The hypothesis function is now

$$\hat{y}_i = w_0 + w_1 x_i^1 + w_2 x_i^2 + \cdots + w_M x_i^M$$

- ▶ Where w_j ($j = 0, 1, \dots, M$) are the unknown weights of the data, and x_i^j is the j th feature of the i th input.
- ▶ Next, we denote the discrepancy between y_i and $\hat{y}_i$ as ϵ_i

$$y_i = \hat{y}_i + \epsilon_i$$

- Now let's collect the above equation for all N datapoints:

$$y_1 = \hat{y}_1 + \epsilon_1$$

$$y_2 = \hat{y}_2 + \epsilon_2$$

$$\vdots$$

$$y_N = \hat{y}_N + \epsilon_N$$

- Replacing the values of $\hat{y}$, we get:

$$y_1 = w_0 + w_1x_1^1 + w_2x_1^2 + \cdots + w_Mx_1^M + \epsilon_1$$

$$y_2 = w_0 + w_1x_2^1 + w_2x_2^2 + \cdots + w_Mx_2^M + \epsilon_2$$

$\vdots$

$$y_N = w_0 + w_1x_N^1 + w_2x_N^2 + \cdots + w_Mx_N^M + \epsilon_N$$

- Collecting the equations in matrix form:

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_N \end{bmatrix} = \begin{bmatrix} 1 & x_1^1 & x_1^2 & \cdots & x_1^M \\ 1 & x_2^1 & x_2^2 & \cdots & x_2^M \\ 1 & x_3^1 & x_3^2 & \cdots & x_3^M \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_N^1 & x_N^2 & \cdots & x_N^M \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_M \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \vdots \\ \epsilon_N \end{bmatrix}$$

- Notice the rows of the matrix on the right are data samples:

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_N \end{bmatrix} = \begin{bmatrix} \cdots & \mathbf{x}_1 & \cdots \\ \cdots & \mathbf{x}_2 & \cdots \\ \cdots & \mathbf{x}_3 & \cdots \\ \vdots & \vdots & \vdots \\ \cdots & \mathbf{x}_N & \cdots \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_M \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \vdots \\ \epsilon_N \end{bmatrix}$$

$$\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$$

► Let's formalize some notations:

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_N \end{bmatrix} \quad \mathbf{X} = \begin{bmatrix} \cdots & \mathbf{x}_1 & \cdots \\ \cdots & \mathbf{x}_2 & \cdots \\ \cdots & \mathbf{x}_3 & \cdots \\ \vdots & \vdots & \vdots \\ \cdots & \mathbf{x}_N & \cdots \end{bmatrix} \quad \boldsymbol{\theta} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_M \end{bmatrix} \quad \boldsymbol{\epsilon} = \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \vdots \\ \epsilon_N \end{bmatrix}$$

$$\mathbf{y} = \mathbf{X}\boldsymbol{\theta} + \boldsymbol{\epsilon}$$

- Notice that we are using the MSE cost function:

$$J = \frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$$

- Using the definition of epsilon we can write the above as:

$$J = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 = \frac{1}{N} \sum_{i=1}^N (\epsilon_i)^2$$

- Using the definition of dot product the above can be written as:

$$J = \frac{1}{N} \sum_{i=1}^N (\epsilon_i)^2 = \frac{1}{N} \epsilon^T \epsilon$$

- We get:

$$\frac{\partial J}{\partial \theta} = \mathbf{X}^T 2(\mathbf{y} - \mathbf{X}\theta)$$

- Setting it equal to zero we can solve for θ :

$$\theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Closed-form solution for Linear Regression

- ▶ We can look at ML algorithms from two perspectives:
 - **Loss Minimization** Problem (like what we did).
 - **Probability Maximization** Problem (using Maximum Likelihood Estimation).
- ▶ For linear regression, under particular assumptions, these two approaches yield equivalent solutions.

- ▶ We can also look at the probabilistic interpretation of Linear Regression.
- ▶ Keeping everything else same as the previous formulation

$$y_i = \mathbf{x}_i^T \boldsymbol{\theta} + \epsilon_i$$

- ▶ Now assume that $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$, then

$$y_i \mid \mathbf{x}_i \sim \mathcal{N}(\mathbf{x}_i^T \boldsymbol{\theta}, \sigma^2)$$

- ▶ We can write the conditional distribution as:

$$\mathbb{P}(y_i \mid \mathbf{x}_i) \sim \mathcal{N}(\mathbf{x}_i^T \boldsymbol{\theta}, \sigma^2)$$

- ▶ Let's assume that all data point in the dataset are i.i.d. (independent identically distributed). Then we have:

$$\mathbb{P}(\mathcal{D}) = \prod_{i=1}^N \mathbb{P}(\mathbf{x}_i, y_i)$$

- ▶ Using Bayes Theorem we can write:

$$\prod_{i=1}^N \mathbb{P}(\mathbf{x}_i, y_i) = \prod_{i=1}^N \mathbb{P}(\mathbf{x}_i) \mathbb{P}(y_i | \mathbf{x}_i)$$

- ▶ In simple words, given the dataset, we want to find the values of the unknown parameters which maximize the probability of the dataset.
- ▶ Using the definition of the conditional distribution we have:

$$\mathbb{P}(y_i | \mathbf{x}_i) = \frac{1}{\sigma\sqrt{2\pi}} \exp(-(y_i - \mathbf{x}_i^T \boldsymbol{\theta})^2)$$

- ▶ Using the definition we get:

$$\prod_{i=1}^N \mathbb{P}(\mathbf{x}_i, y_i) = \prod_{i=1}^N \mathbb{P}(\mathbf{x}_i) \prod_{i=1}^N \frac{1}{\sigma\sqrt{2\pi}} \exp(-(y_i - \mathbf{x}_i^T \boldsymbol{\theta})^2)$$

- ▶ Let's try to maximize:

$$\prod_{i=1}^N \mathbb{P}(\mathbf{x}_i, y_i) = \prod_{i=1}^N \mathbb{P}(\mathbf{x}_i) \prod_{i=1}^N \frac{1}{\sigma\sqrt{2\pi}} \exp(-(y_i - \mathbf{x}_i^T \boldsymbol{\theta})^2)$$

- ▶ Note that:


$$\arg \max_{\boldsymbol{\theta}} \prod_{i=1}^N \mathbb{P}(\mathbf{x}_i, y_i) = \arg \max_{\boldsymbol{\theta}} \prod_{i=1}^N \exp(-(y_i - \mathbf{x}_i^T \boldsymbol{\theta})^2)$$

- Furthermore, since the right-hand side of the above equation is monotonic in θ , the $\arg \max$ will not change if we take the logarithm of the expression.

$$\arg \max_{\theta} \prod_{i=1}^N \exp \left(-(y_i - \mathbf{x}_i^T \theta)^2 \right) = \arg \max_{\theta} \sum_{i=1}^N \left(-(y_i - \mathbf{x}_i^T \theta)^2 \right)$$


- Notice that the right-hand side is minimizing the MSE.
- Hence, the solution of minimizing the MSE is equivalent to Maximum Likelihood Estimation for linear regression.

- What if Y has a non-linear response?




- Can we still use a linear model?



What is Bias and Variance?




$$\{1, x\}$$



$$\{1, x, x^{[2]}, x^{[3]}, x^{[4]}\}$$



Real Bad Overfit?



- ▶ So far, we have minimized the error (loss) with respect to **training data**
 - Low training error does not imply good expected performance: **over-fitting**
- ▶ We would like to reason about the **expected loss (Prediction Risk)** over:
 - Training Data: $\{(y_1, x_1), \dots, (y_n, x_n)\}$
 - Test point: (y_*, x_*)
- ▶ We will decompose the expected loss into:

$$\mathbb{E}_{D, (y_*, x_*)} \left[(y_* - f(x_* | D))^2 \right] = \text{Noise} + \text{Bias}^2 + \text{Variance}$$




Image from <http://scott.fortmann-roe.com/docs/BiasVariance.html>


R-squared (R^2 or r^2) — Coefficient of Determination

- ▶ **A statistical metric** used to measure how well the independent variables explain the variability in the dependent variable.
- ▶ It provides a measure of the goodness-of-fit (**performance**) for a regression model:

$$R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$$

Where:

- ▶ $SS_{\text{residual}} = \sum (y_i - \hat{y}_i)^2$: sum of squared residuals (difference between observed and predicted values).
- ▶ $SS_{\text{total}} = \sum (y_i - \bar{y})^2$: total variation in the data around the mean.



Interpretation:

- ▶ $R^2 = 1$: The model perfectly explains the data.
- ▶ $R^2 = 0$: The model explains none of the variability (as good as guessing the mean).
- ▶ $R^2 < 0$: The model performs worse than a simple horizontal line at the mean of the target variable.

Range:

- ▶ $0 \leq R^2 \leq 1$
 - $R^2 = 1$: Perfect model; predictions perfectly match the observations.
 - $R^2 = 0$: Model does no better than the mean of the dependent variable.

Interpretability:

- ▶ R^2 indicates the proportion of the variance in the dependent variable that is predictable from the independent variables.

Limitations:

- ▶ R^2 increases with the addition of independent variables, even if they don't improve model prediction significantly.
- ▶ Does not account for overfitting or the complexity of the model.

Negative Values:

- ▶ In rare cases, R^2 can be negative when the model fits the data worse than a horizontal line representing the mean of the dependent variable.

Adjusted R^2 :

- ▶ Adjusted R^2 penalizes for the addition of non-significant predictors and is often a better metric for comparing models.
- ▶ Here, n is the number of data points and p is the number of predictors:

$$R_{\text{adj}}^2 = 1 - \frac{(1 - R^2)(n - 1)}{n - p - 1}$$

R-squared vs. Cost Function:

- ▶ **Cost Function:** A mathematical function used to optimize a model during training by minimizing the prediction error (e.g., Mean Squared Error or MSE).
 - Example: Gradient Descent minimizes MSE for linear regression.
- ▶ **R-squared:** A metric to evaluate how well the model fits the data **after training**. It is not used during model optimization.

Limitations of R^2 :

1. **Doesn't Indicate Causation:** A high R^2 doesn't mean the predictors cause the dependent variable.
2. **Sensitive to Overfitting:** A complex model might have a high R^2 but poor generalization.
3. **Adjusted R^2 :** For models with many predictors, use adjusted R^2 , which accounts for the number of predictors to avoid overestimating performance.

- ▶ To ensure your model doesn't overfit to the training data, you should have another subset called **testing data**.
- ▶ You will evaluate your model against this **subset**, and based on its **metric score** (e.g., **accuracy**) you will decide if it's overfitting or not.
- ▶ But how should I split my data?

► Hold-out set:


- A portion of the dataset set aside and not used during training.
- E.g. 80% for training and 20% for testing.

► Issues:

- Imagine you have these labels: [1, 1, 2, 2, 2, 3, 3, 3, 3, 3] and you took last 30% as test: [3,3,3]. **You didn't include 1 and 2 in test!**
- **Solution:** Always shuffle before split: [3, 2, 3, 1, 3, 2, 3, 1, 3, 2] $\Rightarrow$ test: [1,3,2]
- My dataset is small. Taking 20% as test would not be representative!
- **Solution:** Use KFold.

► K-Fold Cross Validation (CV):

- Split data into k parts (folds)
- Train on $k - 1$ folds, test on the remaining fold
- Repeat k times and average the scores



Regularization

Slides from: <https://harvard-iacs.github.io/2019-CS109A/lectures/lecture8/presentation/Lecture8a-Regularization.pptx>

- ▶ Regularization involves modifying the loss function L to improve model generalization.
- ▶ A regularization term is added to penalize certain properties of the model parameters:

$$L_{\text{reg}}(\beta) = L(\beta) + \lambda R(\beta)$$

- ▶ λ is a scalar that controls the strength of the regularization.
- ▶ The function $R(\beta)$ encodes the properties to penalize (e.g., large coefficients).
- ▶ Minimizing L_{reg} leads to model parameters with more desirable characteristics.

- ▶ LASSO discourages extreme values in model parameters by penalizing their magnitudes.
- ▶ We use MSE as the base loss function.
- ▶ The regularized loss function is:

$$L_{\text{LASSO}}(\beta) = \frac{1}{n} \sum_{i=1}^n |y_i - \beta^\top x_i|^2 + \lambda \sum_{j=1}^J |\beta_j|$$

- ▶ The term $\sum_{j=1}^J |\beta_j|$ is the ℓ_1 norm of the vector β :

$$\sum_{j=1}^J |\beta_j| = \|\beta\|_1$$

- ▶ Regularization term penalizes the **squares** of the parameter magnitudes.
- ▶ The regularized loss function becomes:

$$L_{\text{Ridge}}(\beta) = \frac{1}{n} \sum_{i=1}^n |y_i - \beta^\top x_i|^2 + \lambda \sum_{j=1}^J \beta_j^2$$

- ▶ Note that:

$$\sum_{j=1}^J \beta_j^2 = \|\beta\|_2^2$$

which is the square of the ℓ_2 -norm of β .

- ▶ In both ridge and LASSO regression, the larger the regularization parameter λ , the more we penalize large values in β .
- ▶ If λ is close to zero:
 - The penalty is minimal.
 - We recover the MSE, i.e., ridge and LASSO become ordinary regression.
- ▶ If λ is sufficiently large:
 - The MSE term becomes negligible.
 - The regularization forces β_{ridge} and β_{LASSO} close to zero.
- ▶ To avoid ad-hoc selection, λ should be chosen via cross-validation.

► Ridge Regression Solution:

$$\beta = (X^T X + \lambda I)^{-1} X^T Y$$

► LASSO Regression:

- No closed-form analytical solution.
- The L_1 norm is not differentiable at 0.
- We use the concept of *subdifferential* or *subgradient* to derive manageable expressions.


The solution of the Ridge/Lasso regression involves three steps:

- ▶ Select λ
- ▶ Find the minimum of the ridge/Lasso regression loss function (using the ridge formula), and record the **MSE on the validation/test set**.
- ▶ Find the λ that gives the smallest MSE

The Geometry of Regularization (LASSO)

$$L_{\text{LASSO}}(\beta) = \frac{1}{n} \sum_{i=1}^n |y_i - \beta^T \mathbf{x}|^2 + \lambda \sum_{j=1}^J |\beta_j|$$

$$\hat{\beta}^{\text{LASSO}} = \arg \min L_{\text{LASSO}}(\beta)$$



- ▶ Regularized loss function:



$$L_{\text{LASSO}}(\beta) = \frac{1}{n} \sum_{i=1}^n |y_i - \beta^T \mathbf{x}|^2 + \lambda \sum_{j=1}^J |\beta_j|$$


- ▶ LASSO solution:

$$\hat{\beta}^{\text{LASSO}} = \arg \min L_{\text{LASSO}}(\beta)$$


$$\text{▶ } L_1 = \lambda \sum_{j=1}^J |\hat{\beta}_j^{\text{LASSO}}|$$

$$\text{▶ } L_{\text{MSE}}(\beta) = \frac{1}{n} \sum_{i=1}^n |y_i - \beta^T \mathbf{x}|^2$$






- ▶ The LASSO estimator tends to zero out parameters.
- ▶ The OLS loss can easily intersect with the constraint on one of the axes.




- ▶ The values of the coefficients decrease as λ increases.
- ▶ Coefficients are nullified quickly.


The Geometry of Regularization (Ridge)




- ▶ $L_{\text{Ridge}}(\beta) = \frac{1}{n} \sum_{i=1}^n |y_i - \beta^T \mathbf{x}|^2 + \lambda \sum_{j=1}^J (\beta_j)^2$
- ▶ $\hat{\beta}^{\text{Ridge}} = \arg \min L_{\text{Ridge}}(\beta)$
- ▶ $\lambda \sum_{j=1}^J |\hat{\beta}_j^{\text{Ridge}}|^2 = C$



- ▶ $\frac{1}{n} \sum_{i=1}^n |y_i - \hat{\beta}^{\text{Ridge}^T} \mathbf{x}|^2 = D$
- ▶ The ridge constraint forms a circle in 2D.
- ▶ The solution lies at the point where the MSE contour touches the constraint boundary.




- The ridge estimator is where the constraint and the loss intersect.



- The values of the coefficients decrease as lambda increases, but they are not nullified.

The Geometry of Regularization



Ridge regularization with only **validation**: step by step

1. Split data into $\{\{X, Y\}_{\text{train}}, \{X, Y\}_{\text{validation}}, \{X, Y\}_{\text{test}}\}$
2. For λ in $\{\lambda_{\min}, \dots, \lambda_{\max}\}$:

2.1 Determine the β that minimizes the L_{ridge} :

$$\hat{\beta}_{\text{Ridge}}(\lambda) = (X^T X + \lambda I)^{-1} X^T Y, \quad \text{using the train data}$$

2.2 Record $L_{\text{MSE}}(\lambda)$ using validation data.

3. Select the λ that minimizes the loss on the validation data:

$$\lambda_{\text{ridge}} = \arg \min_{\lambda} L_{\text{MSE}}(\lambda)$$

4. Refit the model using both **train and validation** data:

$$\{X, Y\}_{\text{train}}, \{X, Y\}_{\text{validation}} \Rightarrow \hat{\beta}_{\text{ridge}}(\lambda_{\text{ridge}})$$

5. Report MSE or R^2 on $\{X, Y\}_{\text{test}}$ given the $\hat{\beta}_{\text{ridge}}(\lambda_{\text{ridge}})$

1. Split data into $\{\{X, Y\}_{\text{train}}, \{X, Y\}_{\text{validation}}, \{X, Y\}_{\text{test}}\}$
2. For λ in $\{\lambda_{\min}, \dots, \lambda_{\max}\}$:

2.1 Determine the β that minimizes the L_{lasso} ,

$$\hat{\beta}_{\text{lasso}}(\lambda),$$

using the train data. **This is done using a solver.**

2.2 Record $L_{\text{MSE}}(\lambda)$ using validation data.

3. Select the λ that minimizes the loss on the validation data:


$$\lambda_{\text{lasso}} = \arg \min_{\lambda} L_{\text{MSE}}(\lambda)$$

4. Refit the model using both **train** and **validation** data:

$$\{X, Y\}_{\text{train}}, \{X, Y\}_{\text{validation}} \Rightarrow \hat{\beta}_{\text{lasso}}(\lambda_{\text{lasso}})$$

5. Report MSE or R^2 on $\{X, Y\}_{\text{test}}$ given the $\hat{\beta}_{\text{lasso}}(\lambda_{\text{lasso}})$

Ridge regularization with **validation only**: step by step



1. Remove $\{X, Y\}_{\text{test}}$ from data.
2. Split the rest of the data into K folds: $\{\{X, Y\}_{\text{train}}^{-k}, \{X, Y\}_{\text{val}}^k\}$
3. For $k \in \{1, \dots, K\}$:
 - 3.1 For $\lambda \in \{\lambda_0, \dots, \lambda_n\}$:

- Determine the β that minimizes the L_{ridge} :

$$\hat{\beta}_{\text{ridge}}(\lambda, k) = (X^{\top}X + \lambda I)^{-1}X^{\top}Y$$

using the training data $\{X, Y\}_{\text{train}}^k$.

- Record $L_{\text{MSE}}(\lambda, k)$ using the validation data $\{X, Y\}_{\text{val}}^k$.


At this point, we have a 2D matrix: rows correspond to different k , and columns to different λ values.

4. Average the $L_{\text{MSE}}(\lambda, k)$ for each λ :

$$\bar{L}_{\text{MSE}}(\lambda) = \frac{1}{K} \sum_{k=1}^K L_{\text{MSE}}(\lambda, k)$$

5. Find the λ that minimizes $\bar{L}_{\text{MSE}}(\lambda)$: $\lambda_{\text{ridge}} = \arg \min_{\lambda} \bar{L}_{\text{MSE}}(\lambda)$
6. Refit the model using the full training data $\{\{X, Y\}_{\text{train}}, \{X, Y\}_{\text{val}}\} \Rightarrow \hat{\beta}_{\text{ridge}}(\lambda_{\text{ridge}})$
7. Report MSE or R^2 on $\{X, Y\}_{\text{test}}$ using $\hat{\beta}_{\text{ridge}}(\lambda_{\text{ridge}})$

Ridge regularization with **CV** only: step by step



Validation error MSE_{val} across λ values (log scale) using 5-fold cross-validation.

Each colored line corresponds to one fold. The thick black curve shows the mean validation error

- ▶ Since **LASSO** regression tends to produce zero estimates for a number of model parameters – we say that LASSO solutions are **sparse** – we consider LASSO to be a method for variable selection.
- ▶ Many prefer using LASSO for variable selection (as well as for suppressing extreme parameter values) rather than stepwise selection, as LASSO avoids the statistical problems that arise in stepwise selection.
- ▶ **Question:** What are the pros and cons of the two approaches?