

Reinforcement Learning

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

May 16, 2026

- ▶ Sometimes, it is better to learn a stochastic policy than a deterministic policy
- ▶ Policy gradient methods suffer from poor sample efficiency
- ▶ Importance sampling can help alleviate this problem
- ▶ However, we need to make sure that current policy is not far from policy with which we have collected trajectories
- ▶ Small changes in the policy parameters can unexpectedly lead to big changes in the policy.
- ▶ TRPO uses importance sampling to take multiple gradient steps and constrains the optimization objective in the policy space
- ▶ PPO does the same but approximately enforces KL constraint without computing natural gradients.

► Model-Free RL

- No model
- Learn value function (and/or policy) from experience

► Model-Free RL

- No model
- Learn value function (and/or policy) from experience

► Model-Based RL

- Learn a model from experience
- Plan value function (and/or policy) from model

- ▶ Model-Free RL
 - No model
 - Learn value function (and/or policy) from experience
 - ▶ Model-Based RL
 - Learn a model from experience
 - Plan value function (and/or policy) from model
- ▶ So far, we have only looked at model-free reinforcement learning

- ▶ A model $\mathcal{M}$ is a representation of an MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R} \rangle$
- ▶ We will assume state space $\mathcal{S}$ and action space $\mathcal{A}$ are known
- ▶ So a model $\mathcal{M} = \langle \mathcal{P}_\phi, \mathcal{R}_\phi \rangle$ represents state transitions $\mathcal{P}_\phi \approx \mathcal{P}$ and rewards $\mathcal{R}_\phi \approx \mathcal{R}$

$$s_{t+1} \sim \mathcal{P}_\phi(s_{t+1} \mid s_t, a_t)$$

$$r_{t+1} = \mathcal{R}_\phi(r_{t+1} \mid s_t, a_t)$$

- ▶ A model $\mathcal{M}$ is a representation of an MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R} \rangle$
- ▶ We will assume state space $\mathcal{S}$ and action space $\mathcal{A}$ are known
- ▶ So a model $\mathcal{M} = \langle \mathcal{P}_\phi, \mathcal{R}_\phi \rangle$ represents state transitions $\mathcal{P}_\phi \approx \mathcal{P}$ and rewards $\mathcal{R}_\phi \approx \mathcal{R}$

$$s_{t+1} \sim \mathcal{P}_\phi(s_{t+1} \mid s_t, a_t)$$

$$r_{t+1} = \mathcal{R}_\phi(r_{t+1} \mid s_t, a_t)$$

- ▶ We can also only learn one of state transitions and reward model

- Goal: Estimate model $\mathcal{M}_\phi$ from experience $s_1, a_1, r_1, \dots, s_T$

- ▶ Goal: Estimate model $\mathcal{M}_\phi$ from experience $s_1, a_1, r_1, \dots, s_T$
- ▶ This is a supervised learning problem

- ▶ Goal: Estimate model $\mathcal{M}_\phi$ from experience $s_1, a_1, r_1, \dots, s_T$
- ▶ This is a supervised learning problem
- ▶ Learning $s, a \rightarrow r$ is a regression problem
- ▶ Learning $s, a \rightarrow s'$ is a density estimation problem

- ▶ Goal: Estimate model $\mathcal{M}_\phi$ from experience $s_1, a_1, r_1, \dots, s_T$
- ▶ This is a supervised learning problem
- ▶ Learning $s, a \rightarrow r$ is a regression problem
- ▶ Learning $s, a \rightarrow s'$ is a density estimation problem
- ▶ We pick a loss function like MSE, KL Divergence etc., and then optimize

► Key Idea:

- It would be useful if we could approximately simulate the world! i.e. if we could predict the consequences of our actions

► Key Idea:

- It would be useful if we could approximately simulate the world! i.e. if we could predict the consequences of our actions

► Advantages:

- Can efficiently learn model by supervised learning method
- Can reason about model uncertainty

► Key Idea:


- It would be useful if we could approximately simulate the world! i.e. if we could predict the consequences of our actions

► Advantages:

- Can efficiently learn model by supervised learning method
- Can reason about model uncertainty

► Disadvantages:

- First learn a model, then construct a value function. This means we will have two sources of approximation error



estimate $p(\mathbf{s}' \mid \mathbf{s}, \mathbf{a})$ (model-based)

supervised learning

$$\min_{\phi} \sum_i \|f_{\phi}(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{s}'_i\|^2$$

optimize $\pi_{\theta}(\mathbf{a} \mid \mathbf{s})$ (model-based)

► So, now we have the following algorithm

1. Run some policy (e.g. random policy) to collect data $D = \{(s, a, s')_i\}$
2. Learn model $f_\phi(s, a)$ to minimize $\sum_i \|f_\phi(s_i, a_i) - s'_i\|$
3. Iteratively sample action sequences, run through model $f_\phi(s, a)$ to choose actions

- ▶ So, now we have the following algorithm
 1. Run some policy (e.g. random policy) to collect data $D = \{(s, a, s')_i\}$
 2. Learn model $f_\phi(s, a)$ to minimize $\sum_i \|f_\phi(s_i, a_i) - s'_i\|$
 3. Iteratively sample action sequences, run through model $f_\phi(s, a)$ to choose actions
- ▶ What can go wrong here?

How can this approach fail?



How can this approach fail?



- In the first step, when we sample using random policy we may get trajectories like the red line.

How can this approach fail?



- ▶ In the first step, when we sample using random policy we may get trajectories like the red line.
- ▶ The agent may learn that going right means that we can go higher!

How can this approach fail?



- ▶ In the first step, when we sample using random policy we may get trajectories like the red line.
- ▶ The agent may learn that going right means that we can go higher!
- ▶ Thus, it will learn an incomplete policy and fall at the top.

How can this approach fail?



- ▶ In the first step, when we sample using random policy we may get trajectories like the red line.
- ▶ The agent may learn that going right means that we can go higher!
- ▶ Thus, it will learn an incomplete policy and fall at the top.
- ▶ Data distribution mismatch $p_{\pi_0}(s) \neq p_{\pi_f}(s)$


- ▶ How might we alleviate this issue?


- ▶ How might we alleviate this issue?
- ▶ Let's modify the algorithm a bit
 1. Run some policy (e.g. random policy) to collect data $D = \{(s, a, s')_i\}$
 2. Learn model $f_\phi(s, a)$ to minimize $\sum_i \|f_\phi(s_i, a_i) - s'_i\|$
 3. Iteratively sample action sequences, run through model $f_\phi(s, a)$ to choose actions

- ▶ How might we alleviate this issue?
- ▶ Let's modify the algorithm a bit
 1. Run some policy (e.g. random policy) to collect data $D = \{(s, a, s')_i\}$
 2. Learn model $f_\phi(s, a)$ to minimize $\sum_i \|f_\phi(s_i, a_i) - s'_i\|$
 3. Iteratively sample action sequences, run through model $f_\phi(s, a)$ to choose actions
 4. Execute planned actions, appending visiting tuples (s, a, s') to D
 5. Go to step 2


- ▶ How might we alleviate this issue?
- ▶ Let's modify the algorithm a bit
 1. Run some policy (e.g. random policy) to collect data $D = \{(s, a, s')_i\}$
 2. Learn model $f_\phi(s, a)$ to minimize $\sum_i \|f_\phi(s_i, a_i) - s'_i\|$
 3. Iteratively sample action sequences, run through model $f_\phi(s, a)$ to choose actions
 4. Execute planned actions, appending visiting tuples (s, a, s') to D
 5. Go to step 2
- ▶ Now, we update the model using data collected with planning
- ▶ Then, replan periodically to help account for mistakes.

Revisiting the Cliff






- Initially, it's still the same. The agent learns that going right means that we can go higher!



- ▶ Initially, it's still the same. The agent learns that going right means that we can go higher!
- ▶ Then we re-train the model using the new observations



- ▶ Initially, it's still the same. The agent learns that going right means that we can go higher!
- ▶ Then we re-train the model using the new observations
- ▶ The final policy is different now. It learns to go to the top and stop.

- We can also use model-based RL for policy optimization

- ▶ We can also use model-based RL for policy optimization
- ▶ Key Idea:
 - Augment data with model-simulated roll-outs

- ▶ We can also use model-based RL for policy optimization
- ▶ Key Idea:
 - Augment data with model-simulated roll-outs
- ▶ Algorithm:
 1. Collect data using current policy π_θ , add to D_{env}
 2. Update model $p_\phi(s' | s, a)$ using D_{env}
 3. Collect synthetic roll-outs using π_θ in model p_ϕ from states in D_{env} ; add to D_{model}
 4. Update policy π_θ (and critic Q) using D_{env}
 5. Go to step 1

- ▶ We can also use model-based RL for policy optimization
- ▶ Key Idea:
 - Augment data with model-simulated roll-outs
- ▶ Algorithm:
 1. Collect data using current policy π_θ , add to D_{env}
 2. Update model $p_\phi(s' | s, a)$ using D_{env}
 3. Collect synthetic roll-outs using π_θ in model p_ϕ from states in D_{env} ; add to D_{model}
 4. Update policy π_θ (and critic Q) using D_{env}
 5. Go to step 1
- ▶ Compatible with variety of model-free RL methods
- ▶ Could additionally use D_{env} in policy update (step 4)

How to augment?

- ▶ Generate full trajectories from initial states?

How to augment?

- ▶ Generate full trajectories from initial states?
 - Model may not be accurate for long horizons

How to augment?

- ▶ Generate full trajectories from initial states?
 - Model may not be accurate for long horizons
- ▶ Generate partial trajectories from initial states?

How to augment?

- ▶ Generate full trajectories from initial states?
 - Model may not be accurate for long horizons
- ▶ Generate partial trajectories from initial states?
 - May not get good coverage of later states


How to augment?

- ▶ Generate full trajectories from initial states?
 - Model may not be accurate for long horizons
- ▶ Generate partial trajectories from initial states?
 - May not get good coverage of later states
- ▶ Generate partial trajectories from all states in the data

How to augment?

- ▶ Generate full trajectories from initial states?
 - Model may not be accurate for long horizons
- ▶ Generate partial trajectories from initial states?
 - May not get good coverage of later states
- ▶ Generate partial trajectories from all states in the data

Example real trajectory ●



► Advantages:

- Models are immensely useful if easy to learn
- Model can be trained without reward labels (self-supervised)
- Model is somewhat task-agnostic (can sometimes be transferred across rewards)

► Advantages:

- Models are immensely useful if easy to learn
- Model can be trained without reward labels (self-supervised)
- Model is somewhat task-agnostic (can sometimes be transferred across rewards)

► Disadvantages:

- Models don't optimize for task performance
- Sometimes harder to learn than a policy

► Advantages:

- Models are immensely useful if easy to learn
- Model can be trained without reward labels (self-supervised)
- Model is somewhat task-agnostic (can sometimes be transferred across rewards)

► Disadvantages:

- Models don't optimize for task performance
- Sometimes harder to learn than a policy

Whether to use a model depends on how hard it is to learn!

Case study: Model-based RL for dexterous manipulation

Model-free methods:

SAC: actor-critic method

NPG: policy gradient method





Model-based methods:

PDDM: proposed method

MBPO: RL with model-generated data

PETS: CEM-based planner

Nagabandi et al.: random shooting, no ensembles



Nagabandi, et al. Deep Dynamics Models for Learning Dexterous Manipulation

- ▶ Reinforcement Learning can be categorised into Model-Based and Model-Free.
- ▶ In Model-Based RL, we also learn the model of the world i.e, the state transition model and/or the reward model
- ▶ Model-Based RL can be employed with both value and/or policy based methods
- ▶ We can use the model to augment the real environment data
- ▶ Models can be immensely useful if easy to learn but sometimes they can be harder to learn than a policy
- ▶ Whether to use a model depends on how hard it is to learn

These slides have been adapted from

- ▶ Chelsea Finn & Karol Hausman, [Stanford CS224R: Deep Reinforcement Learning](#)
- ▶ Sergey Levine, [Berkeley CS285 Deep Reinforcement Learning](#)
- ▶ David Silver, Deepmind: [Introduction to Reinforcement Learning](#)