

Introduction to Transformers

Naeemullah Khan

naeemullah.khan@kaust.edu.sa




جامعة الملك عبد الله
للعلوم والتقنية

King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

June 18, 2026



1. Motivation
2. Learning Outcomes
3. Introduction to Transformers
4. Self-Attention Mechanism
5. Multi-Head Attention
6. Positional Encoding
7. Transformer Block
8. Encoder-Decoder Architecture
9. Casual Masking
10. Pre-LayerNorm vs. Post-LayerNorm Architectures
11. Layer Normalization vs RMSNorm
12. Transformer Variants

► **Traditional RNNs/LSTMs suffer from:**

- Sequential computation → slow training
- Vanishing gradients → poor long-range dependencies

► **Need for a model that:**

- Handles global context
- Is parallelizable
- Supports scalability

After this session, you should be able to:

- ▶ Understand the self-attention mechanism
- ▶ Learn how multi-head attention works and why it helps
- ▶ Grasp the role of positional encoding
- ▶ Deconstruct transformer blocks and variants
- ▶ Compare Pre-LayerNorm vs. Post-LayerNorm
- ▶ Evaluate LayerNorm vs. RMSNorm
- ▶ Explore causal masking in decoder-only transformers
- ▶ Understand tradeoffs and future directions


► What is a Transformer?



- Introduced by Vaswani et al. (2017) – “Attention is All You Need”

► Key features:



- No recurrence, all attention
- Encoder-decoder structure
- Enables parallel processing
- Scales well with data and compute

- ▶ Tokenization
- ▶ Input Embeddings
- ▶ Position Encodings
- ▶ Query, Key, & Value
- ▶ Attention
- ▶ Self Attention
- ▶ Multi-Head Attention
- ▶ Feed Forward
- ▶ Add & Norm
- ▶ Encoders
- ▶ Masked Attention
- ▶ Encoder Decoder Attention
- ▶ Linear
- ▶ Softmax
- ▶ Decoders
- ▶ Encoder-Decoder Models





Processing Inputs



- ▶ Computes contextual representation of a word using all other words
- ▶ **Inputs:** Query (Q), Key (K), Value (V)
- ▶ **Output:** Weighted sum of values


Attention formula:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$$

- ▶ Enables each word to attend to all others

- ▶ **Query:** What we're looking for (e.g., the meaning of “bank” in context)
- ▶ **Keys:** Tags or features associated with each word/sentence (e.g., “river”, “money”)
- ▶ **Values:** Information linked to each key (e.g., context-specific meaning)
- ▶ Uses dot-product similarity between Query and Keys to assign attention weights
- ▶ Weighted sum of Values produces the contextual representation

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



The

animal

didn't

cross

the

street


because

it


was


too


wide




coreference resolution?









SELF



Query Inputs



=



Key Inputs



=



Value Inputs



















Coreference resolution ✓



Sentence boundaries ?

Coreference resolution





Context ?

Semantic relationships ?

Part of Speech ?

Comparisons ?



- ▶ Multi-head attention allows the model to jointly attend to information from different representation subspaces at different positions.
- ▶ It consists of multiple attention heads, each with its own set of parameters.
- ▶ The outputs of all heads are concatenated and linearly transformed.


Formula:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O$$

where each head is defined as:

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

- ▶ W_i^Q , W_i^K , and W_i^V are learned projection matrices for each head.
- ▶ W^O is the output projection matrix that combines the outputs of all heads.



Sentence boundaries



Coreference resolution



Context



Semantic relationships





Part of speech




Comparisons










$$d_h = \frac{d_{model}}{h}$$



Multi-Head Attention









Multi Head Attention : Z

$$d_h = \frac{d_{model}}{h}$$

$$R^{T \times d_{model}}$$





Sentence boundaries



Coreference resolution



Context



Semantic relationships



Part of speech



Comparisons






- ▶ Enables the model to jointly attend to information from different representation subspaces at different positions.
- ▶ Captures various linguistic and syntactic patterns by using multiple attention heads.
- ▶ Improves the model's ability to represent complex relationships in the data.
- ▶ Provides richer and more diverse feature representations.

- ▶ Transformers lack sequence order awareness.
- ▶ **Positional encodings** are added to input embeddings to provide information about token positions.
- ▶ **Sinusoidal encoding:**


$$PE(pos, 2i) = \sin \left(\frac{pos}{10000^{\frac{2i}{d_{\text{model}}}}} \right)$$
$$PE(pos, 2i + 1) = \cos \left(\frac{pos}{10000^{\frac{2i}{d_{\text{model}}}}} \right)$$

I ate an apple <eos>






Requirements for Positional Encodings???



Requirements for Positional Encodings

- ▶ Some representation of time? (like seq2seq?)
- ▶ Should be unique for each position – not cyclic



Requirements for Positional Encodings


- ▶ Some representation of time? (like seq2seq?)
- ▶ Should be unique for each position – not cyclic

Possible Candidates:

$$P_{t+1} = P_t + \Delta c$$

$$P_{t+1} = e^{t\Delta c}$$

$$P_{t+1} = P_t e^{\Delta c}$$



Requirements for Positional Encodings



- Some representation of time? (like **seq2seq**?)
- Should be unique for each position – not cyclic

Possible Candidates :

$$P_{t+1} = P_t + \Delta c$$

$$P_{t+1} = e^{P_t \cdot c}$$

$$P_{t+1} = P_t \cdot t \Delta c$$



Requirements for Positional Encodings



- Some representation of time? (like **seq2seq**?)
- Should be unique for each position – not cyclic

Possible Candidates :

$$P_{t+1} = P_t + \Delta c$$

$$P_{t+1} = P_t \cdot c$$

$$P_{t+1} = P_t^{t \Delta c}$$




Requirements for Positional Encodings

- ▶ Some representation of time? (like seq2seq?)
- ▶ Should be unique for each position – not cyclic
- ▶ **Bounded**

Possible Candidates

$$P(t + t') = M^{t'} \times P(t)$$



Requirements for Positional Encodings


- ▶ Some representation of time? (like seq2seq?)
- ▶ Should be unique for each position – not cyclic
- ▶ **Bounded**

Possible Candidates

$$P(t + t') = M^{t'} \times P(t)$$

M?

1. Should be a unitary matrix
2. Magnitudes of eigen value should be 1 $\rightarrow$ norm preserving
3. The matrix can be learnt
4. Produces unique rotated embeddings each time




Requirements for Positional Encodings

- Some representation of time? (like **seq2seq**?)
- Should be unique for each position – not cyclic
- **Bounded**

Possible

Candidates

$$P(t + t') = M^{t'} \times P(t)$$




Requirements for Position Encodings

- Some representation of time? (like **seq2seq**?)
- Should be unique for each position
- Bounded

Actual Candidates


$\text{sine}(g(t))$

$\text{cosine}(g(t))$



Requirements for $g(t)$

- Must have same dimensions as input embeddings
- Must produce overall unique encodings



For each position, an embedded input is moved the same distance but at a different angle. **Inputs that are close to each other in the sequence have similar perturbations, but inputs that are far apart are perturbed in different directions.**

$$PE_{(pos, 2i)} = \sin\left(pos/10000^{2i/d_{model}}\right)$$


$$PE_{(pos, 2i+1)} = \cos\left(pos/10000^{2i/d_{model}}\right)$$

pos idx of the token in input sentence

i i^{th} dimension out of d

d_{model} embedding dimension of each token

Different calculations for odd and even embedding indices.







Positional Encoding:

	0	1	2	3	4
Dim 1	0.000	0.841	0.909	0.141	-0.757
Dim 2	1.000	0.540	-0.416	-0.990	-0.654
Dim 3	0.000	0.025	0.050	0.075	0.100
Dim 4	1.000	1.000	0.999	0.997	0.995
Dim 5	0.000	0.001	0.001	0.002	0.003

Positional Encoding






Final Input Embeddings

Position Encodings

Input Embeddings



▶ Learnable position embeddings:

- Position embeddings are parameters learned during training.
- Allow the model to adapt position representations to the task.

▶ Rotary Position Embeddings (RoPE):

- Encodes positions by rotating query and key vectors in multi-head attention.
- Enables better extrapolation to longer sequences.

▶ ALiBi (Attention with Linear Biases):


- Adds linear biases to attention scores based on distance between tokens.
- Encourages attention to nearby tokens without explicit position embeddings.

- ▶ **Multi-head Attention**
- ▶ **Add & Norm**
- ▶ **Feed-forward Layer**
- ▶ **Add & Norm**

Note



Residual connections and layer normalization are crucial for effective training.



Transformer Block







- ▶ **Encoder:** Stack of self-attention and feed-forward (FF) layers.
- ▶ **Decoder:** Adds masked self-attention and encoder-decoder attention on top of FF layers.
- ▶ **Applications:** Widely used in tasks like translation, summarization, and more.



WHERE IS THE
CONTEXT?












- ▶ **Example:** BERT
- ▶ Only the encoder stack is used.
- ▶ **Applications:** Classification, question answering, embeddings.

Encoder

ENCODER



Encoder




⋮





Input to Encoder_{*i*+1}


↑

Output from Encoder_{*i*}





Encoder Decoder Attention ?




Encoder Self Attention

1. Queries from Encoder Inputs
2. Keys from Encoder Inputs
3. Values from Encoder Inputs

Decoder Masked Self Attention

1. Queries from Decoder Inputs
2. Keys from Decoder Inputs
3. Values from Decoder Inputs




Encoder

Keys from **Encoder Outputs**
Values from **Encoder Outputs**


Decoder

Queries from **Decoder Inputs**

NOTE: Every decoder block receives the same FINAL encoder output





- Non Linearity
- Complex Relationships
- Learn from each other



- ▶ **Example:** GPT, LLaMA
- ▶ Uses only the decoder block.
- ▶ **Masked self-attention (causal):** Each token can only attend to previous tokens.
- ▶ **Applications:** Language modeling, text generation.

DECODER









$$R^{T_d \times d_{model}}$$




Decoder output











Representation



Generation

Input – input tokens
Output – hidden states

Representation




Input – output tokens and hidden states*
Output – output tokens

Generation

Input – input tokens
Output – hidden states

Model can see all timesteps

Representation



Input – output tokens and hidden states*

Output – output tokens

Model can only see previous timesteps


Generation

Input – input tokens
Output – hidden states

Model can see all timesteps

**Does not usually output tokens, so
no inherent auto-regressivity**

Representation



Input – output tokens and hidden states*

Output – output tokens

**Model can only see previous
timesteps**

**Model is auto-regressive with
previous timesteps' outputs**

Generation


Input – input tokens
Output – hidden states

Model can see all timesteps

Does not usually output tokens, so no inherent auto-regressivity

Can also be adapted to generate tokens by appending a module that maps hidden state dimensionality to vocab size

Representation



Input – output tokens and hidden states*

Output – output tokens

Model can only see previous timesteps

Model is auto-regressive with previous timesteps' outputs




Can also be adapted to generate hidden states by looking before token outputs

Generation

- ▶ **Purpose:** Prevents the model from attending to future tokens in the sequence.

- **Implementation:** Uses a triangular mask in the self-attention mechanism.

- **Effect:** Ensures that predictions for a token only depend on previous tokens, maintaining the autoregressive property.



Target: S1b EOT S2a S2b EOT Target: S1b EOT S2b EOT
Model input: S1a S1b EOT S2a S2b Model input: S1a S1b S2a S2b

Possible attention masks:

	S1	S1	EOT	S2	S2
S1	1	0	0	0	0
S1	1	1	0	0	0
EOT	1	1	1	0	0
S2	1	1	1	1	0
S2	1	1	1	1	1

	S1	S1	S2	S2
S1	1	0	0	0
S1	1	1	0	0
S2	0	0	1	0
S2	0	0	1	1

Left: Causal mask when S1 and S2 can attend to each other

Right: Preventing attention from corss-attending between S1 and S2


Feature	Pre-LN	Post-LN
Stability	✓ More stable	✗ Less stable
Gradient flow	✓ Better	✗ Can explode/vanish
Used in	GPT-3, T5	Original Transformer

Table 1: Comparison of Pre-LayerNorm and Post-LayerNorm Architectures

Pre-LN: LayerNorm applied **before** sublayers


Post-LN: LayerNorm applied **after** sublayers

Pre-LayerNorm vs. Post-LayerNorm



Pre-LayerNorm vs. Post-LayerNorm

Post-LN Transformer





Here, LayerNorms are after attention and after fully connected layers




Better gradients if LayerNorm is placed inside residual connection, before the attention and fully connected layers

Places the layer normalization between the residual blocks: the expected gradients of the parameters near the output layer are large

Pre-LN Transformer



Pre-LayerNorm vs. Post-LayerNorm




LayerNorm: Normalizes across the feature dimension by subtracting the mean and dividing by the standard deviation.


RMSNorm: Removes the mean and scales by the root mean square (RMS) of the features.

- ▶ Simpler
- ▶ Faster
- ▶ Slightly better in large-scale settings [zhou2022rmsnorm](#)

Layer Normalization vs RMSNorm



Layer Normalization vs RMSNorm



- ▶ **Encoder-only:** BERT, RoBERTa
- ▶ **Decoder-only:** GPT, LLaMA, Falcon
- ▶ **Encoder-Decoder:** T5, BART, mT5

Modifications:

- ▶ **Long-range:** Longformer, BigBird
- ▶ **Efficient:** Linformer, Performer
- ▶ **Sparse:** Sparse Transformer, Reformer

Key Ingredients of BERT:

► Pre-training Corpora:


- English Wikipedia (~2.5B words)
- BookCorpus (~800M words)

► Pre-training Tasks:

- **MLM** (Masked Language Modeling): Predict randomly masked words in context
- **NSP** (Next Sentence Prediction): Predict if two sentences naturally follow each other

► Model Sizes:

- **BERT-Base**: 110M parameters
- **BERT-Large**: 340M parameters



What is our takeaway from BERT?

► Pre-training tasks can be invented flexibly...


- Effective representations can be derived from a flexible regime of pre-training tasks.

► Different NLP tasks seem to be highly transferable with each other...

- As long as we have effective representations, that forms a general model which can serve as the backbone for many specialized models.

And scaling works!!!

340M was considered large in 2018



GPT Pre-Training Corpus:


- ▶ Similarly, BooksCorpus and English Wikipedia

GPT Pre-Training Tasks:

- ▶ Predict the next token, given the previous tokens
- ▶ More learning signals than MLM

GPT Pre-Training Results:

- ▶ GPT – 117M Params
- ▶ Similarly competitive on GLUE and SQuAD



What is our takeaway from GPT?


► The Effectiveness of Self-Supervised Learning

- Specifically, the model seems to be able to learn from generating the language *itself*, rather than from any specific task we might cook up.

► Language Model as a Knowledge Base

- Specifically, a generatively pretrained model seems to have a decent zero-shot performance on a range of NLP tasks.

And scaling works!!!




The LLM Era – Paradigm Shift in Machine Learning

BERT – 2018
DistilBERT – 2019
RoBERTa – 2019
ALBERT – 2019
ELECTRA – 2020
DeBERTa – 2020

...

Representation



GPT – 2018
GPT-2 – 2019
GPT-3 – 2020
GPT-Neo – 2021
GPT-3.5 (ChatGPT) – 2022
LLaMA – 2023
GPT-4 – 2023

...

Generation

- ▶ High compute and memory requirements
- ▶ Poor extrapolation capabilities
- ▶ Limitations due to hard-coded positional encoding
- ▶ Quadratic attention cost ($O(n^2)$)
- ▶ Lack of reasoning and grounding

- ▶ Transformers revolutionized deep learning by removing recurrence and leveraging attention.
- ▶ Self-attention allows contextual understanding.
- ▶ Multi-head and positional encodings add flexibility.
- ▶ Variants serve different tasks (BERT, GPT, T5).
- ▶ Ongoing research improves efficiency, grounding, and capabilities.

- ▶ These slides have been adapted from
- ▶ Bhiksha Raj & Rita Singh, 11-785 Introduction to Deep Learning, CMU

1. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017). Attention is All You Need. *In NeurIPS*.
2. Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of Deep Bidirectional Transformers. *In NAACL*.
3. Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). Language Models are Unsupervised Multitask Learners (GPT-2). OpenAI.
4. Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., & Liu, P. J. (2020). Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer (T5). *JMLR*.
5. Zhou, H., Zhang, S., Ma, S., & Sun, J. (2022). RMSNorm: Root Mean Square Layer Normalization. arXiv preprint arXiv:1910.07467.
6. Dao, T., Fu, T., Ermon, S., Rudra, A., & Ré, C. (2022). FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. *In NeurIPS*.
7. Beltagy, I., Peters, M. E., & Cohan, A. (2020). Longformer: The Long-Document Transformer. arXiv preprint arXiv:2004.05150.
8. Tay, Y., Dehghani, M., Bahri, D., & Metzler, D. (2020). Efficient Transformers: A Survey. arXiv preprint arXiv:2009.06732.