Large Language Models

Naeemullah Khan
naeemullah.khan@kaust.edu.sa

جامعة الملك عبدالله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

July 21, 2026

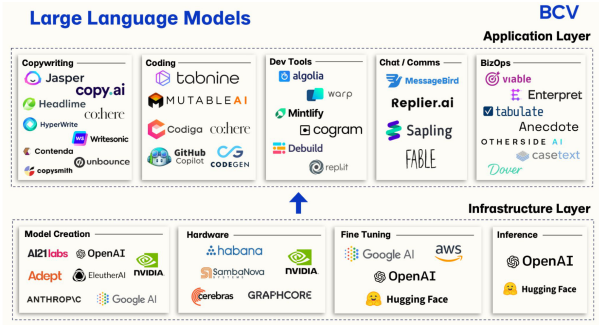
Large Language Models
BCV
Application Layer
Copywriting
Jasper
copy.ai
Headlime
cohere
HyperWrite
Writesonic
Contenda
unbounce
copysmith
Coding
tabnine
MUTABLE AI
Codiga
cohere
GitHub Copilot
CODEGEN
Dev Tools
algolia
warp
Mintlify
cogram
Debuild
replit
Chat / Comms
MessageBird
Replier.ai
Sapling
FABLE
BizOps
viable
Enterpret
tabulate
Anecdote
OTHERSIDE AI
casetext
Dover
Infrastructure Layer
Model Creation
AI21labs
OpenAI
Adept
EleutherAI
NVIDIA
ANTHROP\C
Google AI
Hardware
habana
SambaNova
NVIDIA
cerebras
GRAPHCORE
Fine Tuning
Google AI
aws
OpenAI
Hugging Face
Inference
OpenAI
Hugging Face

Table of Contents

Table of Contents (cont.)

Motivation

- **Why LLMs?**
  - Transforming NLP: ChatGPT, Claude, Gemini, etc.
  - Achieving human-like generation and comprehension.
  - Pivotal for tasks like summarization, translation, reasoning.
- **Need for deeper understanding:**
  - LLMs are expensive to train and operate.
  - Design decisions impact performance significantly (e.g., tokenization, scaling laws).

Learning Outcomes

By the end of this session, you should be able to:

- Explain core architectures like BERT and GPT.
- Understand scaling laws for LLM development.
- Describe tokenization strategies and their impact.
- Discuss limitations of current models (e.g., context window).
- Explore emerging directions like token-free LLMs.

BERT and GPT Concepts

- **BERT (Bidirectional Encoder Representations from Transformers)**
  - Uses transformer encoder only.
  - Trained with Masked Language Modeling (MLM).
  - **Bidirectional:** Considers both left and right context.
  - Fine-tuned for tasks like QA, classification.
  - **Architecture:**
    - Layers of encoder blocks.
    - Positional encodings added to embeddings.
    - Self-attention heads capture dependencies.

BERT – Bidirectional Encoder Representations

- One of the biggest challenges in LM-building used to be the lack of task-specific training data.
- What if we learn an effective representation that can be applied to a variety of downstream tasks?
  - Word2vec (2013)
  - GloVe (2014)

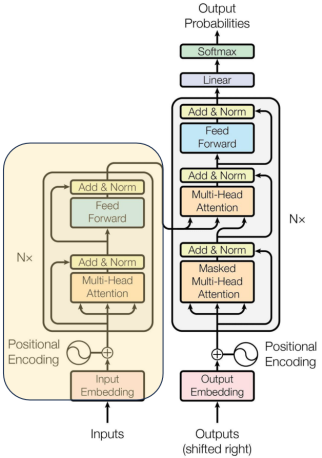

BERT – Bidirectional Encoder Representations

BERT Pre-Training Corpus:

- English Wikipedia – 2,500 million words
- Book Corpus – 800 million words

BERT Pre-Training Tasks:

- MLM (Masked Language Modeling)
- NSP (Next Sentence Prediction)

BERT Pre-Training Results:

- BERT-Base – 110M Params
- BERT-Large – 340M Params

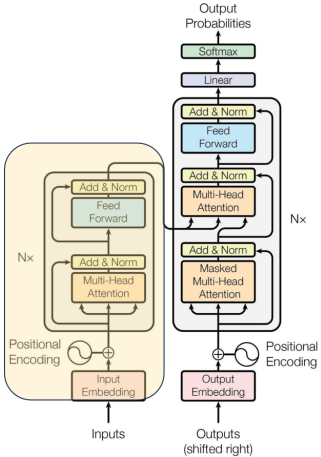

BERT – Bidirectional Encoder Representations

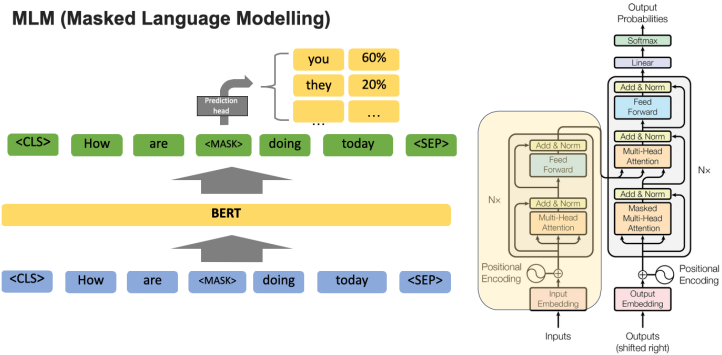

BERT – Bidirectional Encoder Representations

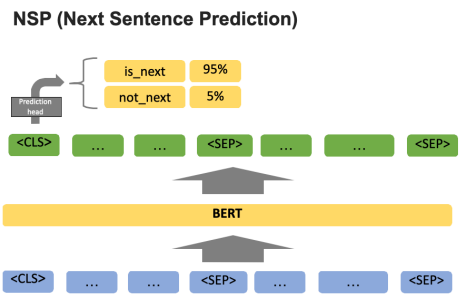

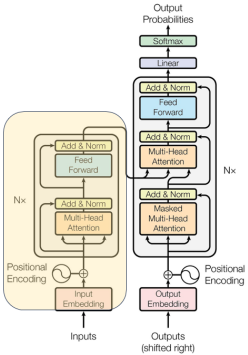

BERT – Bidirectional Encoder Representations

BERT Fine-Tuning:

- Simply add a task-specific module after the last encoder layer to map it to the desired dimension.
- **Classification Tasks:**
  - Add a feed-forward layer on top of the encoder output for the [CLS] token.
- **Question Answering Tasks:**
  - Train two extra vectors to mark the beginning and end of the answer from the paragraph.
- ...

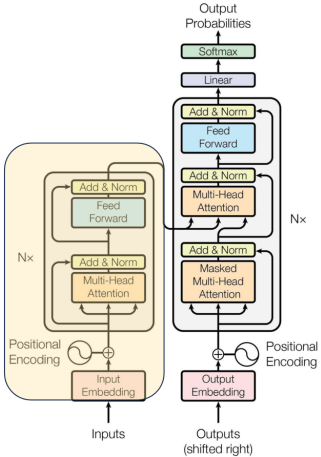

BERT – Bidirectional Encoder Representations

BERT Evaluation:

- **General Language Understanding Evaluation (GLUE):**
  - Sentence pair tasks
  - Single sentence classification
- **Stanford Question Answering Dataset (SQuAD)**

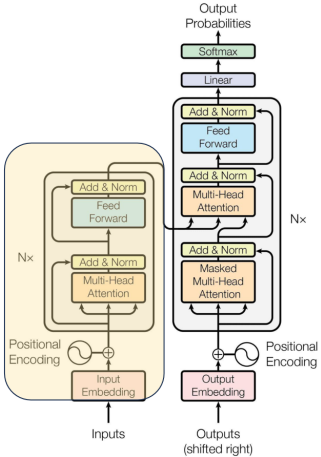

BERT – Bidirectional Encoder Representations

BERT Evaluation:

System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average -
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
$\text{BERT}_{\text{BASE}}$	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
$\text{BERT}_{\text{LARGE}}$	**86.7/85.9**	**72.1**	**92.7**	**94.9**	**60.5**	**86.5**	**89.3**	**70.1**	**82.1**

System	Dev EM	Dev F1	Test EM	Test F1
Leaderboard (Oct 8th, 2018)				
Human	-	-	82.3	91.2
#1 Ensemble - nlnet	-	-	86.0	91.7
#2 Ensemble - QANet	-	-	84.5	90.5
#1 Single - nlnet	-	-	83.5	90.1
#2 Single - QANet	-	-	82.5	89.3
Published				
BiDAF+ELMo (Single)	-	85.8	-	-
R.M. Reader (Single)	78.9	86.3	79.5	86.6
R.M. Reader (Ensemble)	81.2	87.9	82.3	88.5
Ours				
$\text{BERT}_{\text{BASE}}$ (Single)	80.8	88.5	-	-
$\text{BERT}_{\text{LARGE}}$ (Single)	84.1	90.9	-	-
$\text{BERT}_{\text{LARGE}}$ (Ensemble)	85.8	91.8	-	-
$\text{BERT}_{\text{LARGE}}$ (Sgl.+TriviaQA)	**84.2**	**91.1**	**85.1**	**91.8**
$\text{BERT}_{\text{LARGE}}$ (Ens.+TriviaQA)	**86.2**	**92.2**	**87.4**	**93.2**

Table 2: SQuAD results. The BERT ensemble is 7x systems which use different pre-training checkpoints and fine-tuning seeds.

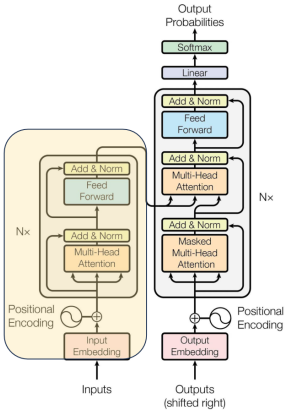

BERT and GPT Concepts

- **GPT (Generative Pre-trained Transformer)**
  - Uses transformer **decoder only**.
  - Trained with next-token prediction.
  - **Unidirectional:** Considers left context only.
  - Fine-tuned for text generation, dialogue systems.
  - **Architecture:**
    - Layers of decoder blocks.
    - Causal masking to prevent future token access.
    - Self-attention heads capture sequential dependencies.

GPT – Generative Pre-trained Transformer

- Similarly motivated as BERT, though differently designed
- Can we leverage large amounts of unlabeled data to pretrain an LM that understands general patterns?

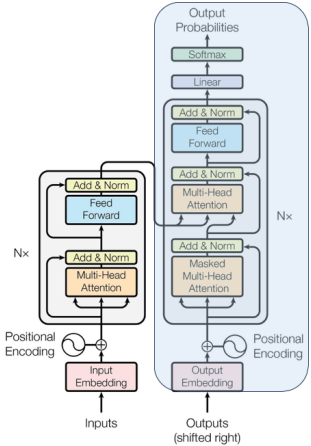

GPT Fine-Tuning:

- Prompt-format task-specific text as a continuous stream for the model to fit

Summarization

Summarize this article:

The summary is:

QA

Answer the question based on the context.

Context:

Question:

Answer:

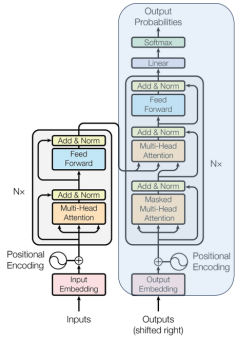

- Key Differences:

Feature	**BERT**	**GPT**
Directionality	Bidirectional	Unidirectional
Objective	Masked Language Modeling (MLM)	Causal Language Modeling (CLM)
Output	Contextual embeddings	Text generation
Usage	Downstream tasks (e.g., classification, QA)	Generation, few-shot learning

Scaling Laws for LLMs

- **Kaplan et al. (2020):** "Scaling Laws for Neural Language Models"
- Performance improves predictably with:
  - More parameters
  - More compute
  - Larger datasets
- **Optimal allocation of compute:** Train bigger models with less data, rather than small models with lots of data.
- **Implication:** LLMs like GPT-3 (175B), GPT-4 (est. >500B) are products of scaling laws.

Scaling in GPT-2

- Scaling improves the perplexity of the LM and improves performance

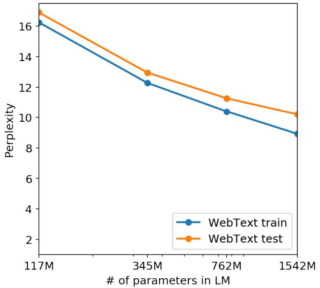

Why is this interesting? Look at data scaling

- We know that typical scaling effects look like this when we increase the amount of training data

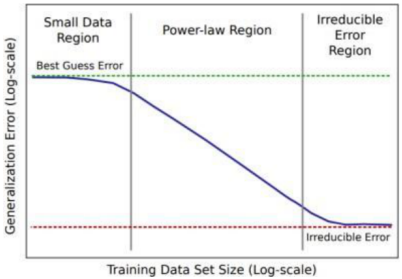

Why is this interesting? Look at data scaling

- Loss and dataset size is linear on a log-log plot
- This is “power-law scaling”

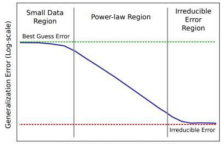

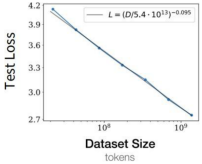

Scaling – (Kaplan, 2020)

- Can we understand scaling by positing scaling laws?
- With scaling laws, we can make decisions on architecture, data, and hyperparameters by training smaller models.
- **OpenAI Study:** Scaling Laws for Neural Language Models (Kaplan et al., 2020)

Scaling – (Kaplan, 2020)

- **OpenAI Study:** Scaling Laws for Neural Language Models (Kaplan et al., 2020)

Key Findings:

- Performance depends strongly on scale, and only weakly on the model shape.
- Larger models are more sample-efficient.
- Smooth power laws ($y = ax^b$) observed between empirical performance and:
  - N – number of parameters
  - D – dataset size
  - C – compute

Scaling Effects

- The effect of some hyperparameters on large language models (LMs) can be predicted before full training—such as optimizer choice (Adam vs. SGD), model depth, and architecture (LSTM vs. Transformer).

Idea:

- Train a few smaller models
- Establish a scaling law (e.g., Adam vs. SGD scaling law)
- Select optimal hyperparameters based on the scaling law prediction

Model Scaling: GPT-3

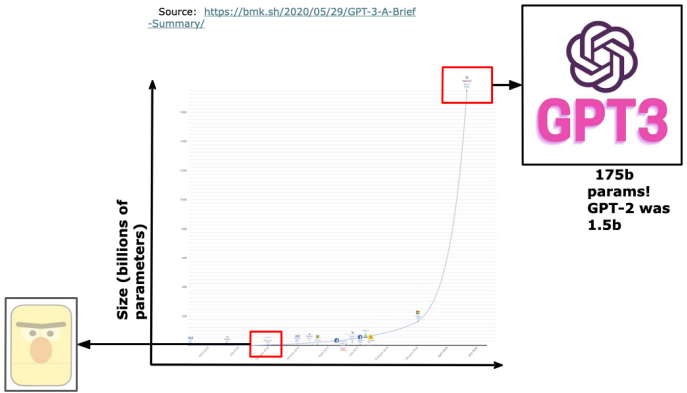

Emergent Abilities with GPT-3 – Wei et al., 2022

- **Emergent abilities:**
  - Not present in smaller models but is present in larger models
  - Do LLMs like GPT-3 have these?
- **Findings:**
  - GPT-3 trained on text can do arithmetic problems like addition and subtraction
  - Different abilities "emerge" at different scales

Emergent Abilities with GPT-3 – Wei et al., 2022

- **Emergent abilities:**
  - Not present in smaller models but is present in larger models
  - Do LLMs like GPT-3 have these?
- **Findings:**
  - GPT-3 trained on text can do arithmetic problems like addition and subtraction
  - Different abilities "emerge" at different scales
  - Model scale is **not** the only contributor to emergence – for 14 BIG-Bench tasks, LaMDA 137B and GPT-3 175B models perform at near-random, but PaLM 62B achieves above-random performance
  - Problems LLMs can't solve today may be emergent for future LLMs

Pre-training Overview

- **Pre-training Phase:**
  - Large-scale unsupervised training on corpus (e.g., Common Crawl, Books)
  - Objective: learn general-purpose language representations
- **Why Pre-train?**
  - Data-efficient fine-tuning
  - Enables zero-shot and few-shot capabilities
  - Foundation for instruction tuning, alignment
- **Challenges:**
  - Massive compute costs
  - Environmental concerns (carbon footprint)

Pre-training Overview

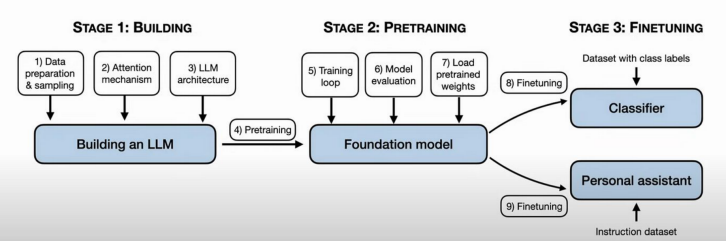

Training of Decoder-only LLMs — Llama 2

- **Auto-regressive Pre-training**
  Train to predict the next token on very large scale corpora ($\sim$2 trillion tokens)
- **Instruction Fine-tuning / Supervised Fine-tuning (SFT)**
  Fine-tune the pre-trained model with pairs of (instruction + input, output) using a large dataset, then further with a small high-quality dataset

Instruction fine-tuning provides, as a prefix, a natural language description of the task along with the input.

- **Example:** Translate into French this sentence: `my name is` $\rightarrow$ `je m'appelle`

Supervised Fine-tuning versus Pre-training

- **Objective function**
  - Loss computed only for target tokens in SFT, all tokens are targets in pre-training
- **Input and Target**
  - Instruction + input as input with the target in SFT and only input as input with shifted input as target
- **Purpose**
  - Pre-training makes good generalist auto-completes but good SFT builds models that can do many unseen tasks
  - SFT can also guide nature of outputs in terms of safety and helpfulness

Instruction Tuning (Wei et al. 2021)

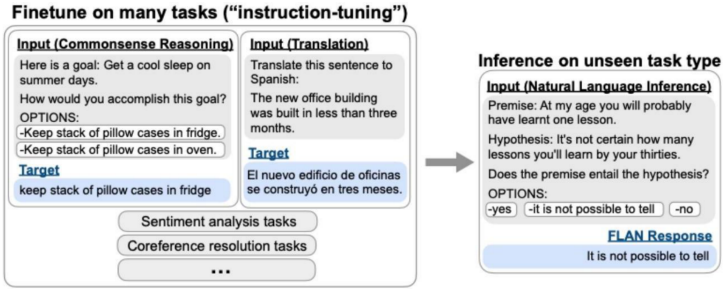

Unsafe Outputs – Alignment Problem

- LLMs may produce
  - Harmful text – unparliamentary language, bias and discrimination
  - Text that can cause direct harm – allowing easy access to dangerous information
- Therefore, LLMs should be trained to produce outputs that align with human preferences and values
- Modern LLMs do so by using SFT and by using human preference directly in model training

Tokenization: Why It Matters?

- **Language $\neq$ Input**
  - Text must be converted to numbers $\rightarrow$ tokens.
- **Tokenization:**
  - Splits text into manageable units.
  - Balances granularity and vocabulary size.
- **Good tokenizer =**
  - Efficient sequence length
  - High vocabulary coverage
  - Robust across domains (code, multilingual, etc.)

Machine Translation

Targets

Ich habe einen
Apfel gegessen

Inputs

I ate an apple

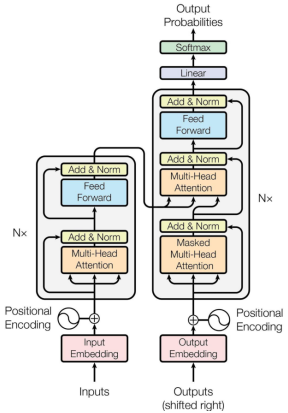

Processing Inputs

Inputs

I ate an apple

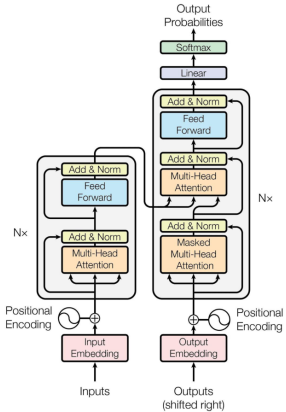

Tokenization

Tokenizer (split into individual words)

I ate an apple

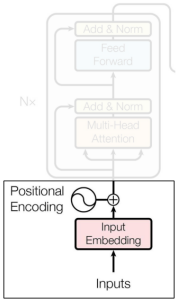

Tokenization

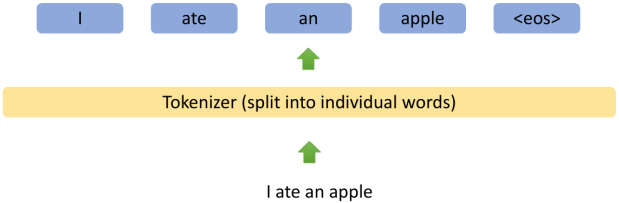

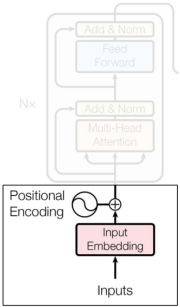

Different Levels of Tokenization

- **Character-level**
  - Fine-grained, handles unknowns
  - Long sequences, slow training
- **Word-level**
  - Intuitive, natural boundaries
  - Large vocab, OOV (Out-of-Vocab) issues
- **Subword-level**
  - Balance of generalization & compactness
  - Requires segmentation algorithm
- **SentencePiece/Unigram LM**
  - Learned from data, language-agnostic

Different Levels of Tokenization

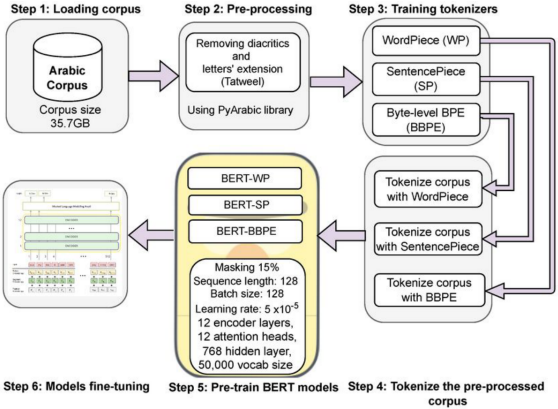

Different Levels of Tokenization

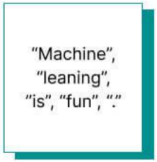

Word Tokenization

Character Tokenization

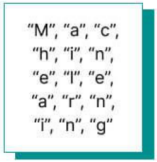

Tokenization Of Subwords

Byte-Pair Encoding (BPE)

- **How BPE works:**
  - Start with characters.
  - Merge most frequent pair (e.g., "t", "h" $\rightarrow$ "th").
  - Repeat until vocab size reached.
- **Advantages:**
  - Handles unknown words well (e.g., unbelievably $\rightarrow$ un + believ + ably)
  - Reduces sequence length
- **Limitations:**
  - Not optimal for all scripts (e.g., Chinese)
  - Word boundaries may be unclear

Byte-Pair Encoding (BPE)

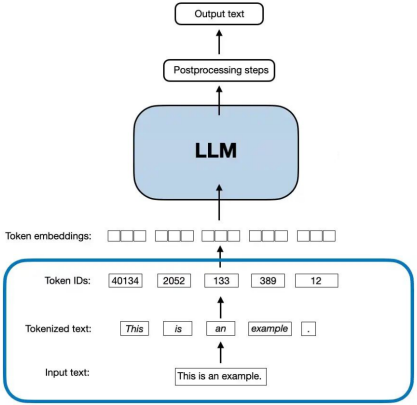

Token-Free LLMs

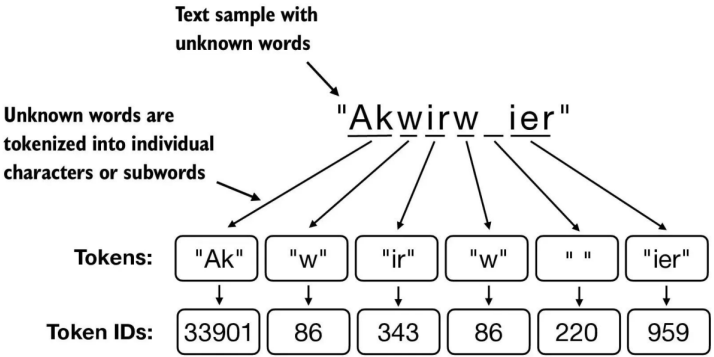

Token-Free LLMs

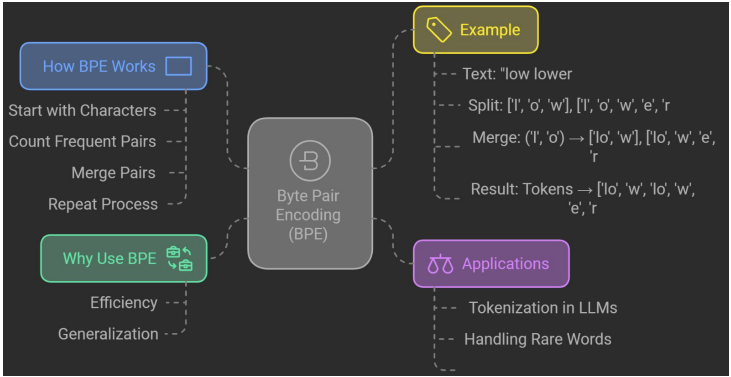

Token-Free LLMs

- **What are Token-Free LLMs?**
  - Models that operate directly on raw text without tokenization.
  - Aim to eliminate the need for pre-defined tokens.
- **Advantages:**
  - Avoids tokenization errors and biases.
  - Can handle arbitrary text lengths and formats.
  - Potentially more efficient for certain tasks.
- **Challenges:**
  - Requires novel architectures to process raw text effectively.
  - May struggle with long-range dependencies without tokenization.

ByT5: Towards a Token-Free Future with Pre-trained Byte-to-Byte Models

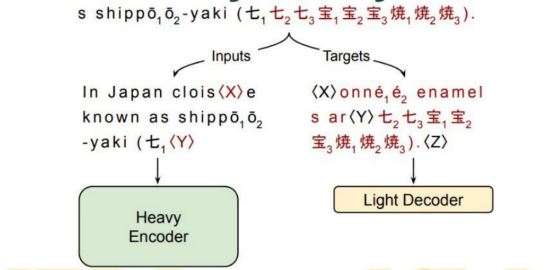

Token-Free LLMs

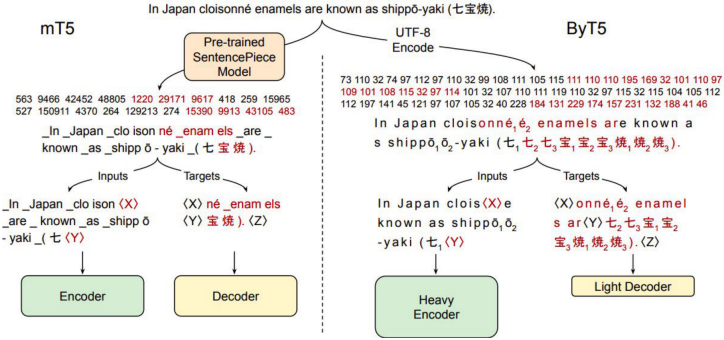

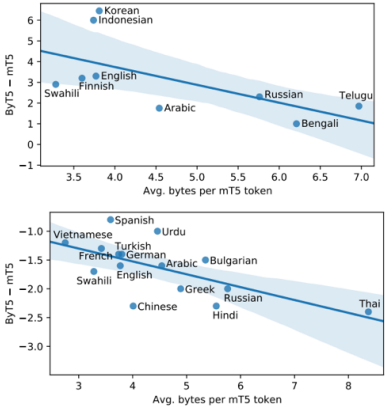

Figure 3: Per-language performance gaps between ByT5-Large and mT5-Large, as a function of each language's "compression rate". **Top**: TyDiQA-GoldP gap. **Bottom**: XNLI zero-shot gap.

Token-Free LLMs: Example — ByT5

- **Example: ByT5 (Xue et al., 2022)**
  - Character-level variant of T5 that processes raw bytes instead of tokens.
  - Eliminates the need for a tokenizer, enabling better multilingual and low-resource language support.
- **Pros:**
  - No preprocessing or tokenization pipeline required.
  - Handles any language or script without modification.
  - Performs better on low-resource and unseen languages.
- **Cons:**
  - Training is slower due to longer input sequences.
  - Increased computational requirements.
  - May require more data to achieve comparable performance.

Token-Free LLMs

- **Current Research:**
  - Exploring architectures like *Raw Transformer* and *Text2Vec*.
  - Investigating how to maintain performance without tokenization.
- **Future Directions:**
  - Integrating token-free approaches with existing LLMs.
  - Enhancing efficiency and scalability of token-free models.

Tokenization Is Not Language-Neutral

- Tokenizers are trained on English-heavy corpora, so the same sentence costs very different numbers of tokens across languages.
- On parallel multilingual text, the shortest languages are already about 50% longer than English in tokens, and the worst cases are far larger: roughly $3\times$ for Arabic, $11\times$ for Burmese, and up to $15\times$ for scripts such as Shan.
- This is an **equity and cost problem**: longer token sequences mean higher API cost, more latency, and less usable context for speakers of under-represented languages, for the exact same content.
- Byte- and character-level models (previous section) reduce, but do not fully remove, the disparity.

Source: Petrov, La Malfa, Torr, Bibi, "Language Model Tokenizers Introduce Unfairness Between Languages," NeurIPS 2023, arXiv:2305.15425.

Glitch Tokens

- The tokenizer is trained first, then the model is trained on a (sometimes different) corpus. Tokens that live in the vocabulary but almost never appear in training end up with **untrained embeddings**.
- These **glitch tokens** trigger strange behavior: the model cannot repeat them, spells them wrong, or produces evasive or erratic output when they appear.
- The original examples were odd strings scraped into the GPT-2 vocabulary (the "SolidGoldMagikarp" family); similar artifacts appear as rare Arabic and Burmese tokens in multilingual tokenizers.
- Lesson: the tokenizer and the training data must stay in sync, and rare-token embeddings are a genuine failure surface.

Source: Rumbelow & Watkins, "SolidGoldMagikarp (plus, prompt generation)," 2023.

Context Window Problems

- Current LLMs have a finite "context window":
  - **GPT-3:** 2K tokens
  - **GPT-4:** Up to 128K tokens
  - **Claude 3.5:** Up to 200K tokens
- **Problems:**
  - Long documents get truncated
  - Token limit affects reasoning
  - Difficult to do document-level QA or summarization

Example of a Context Window

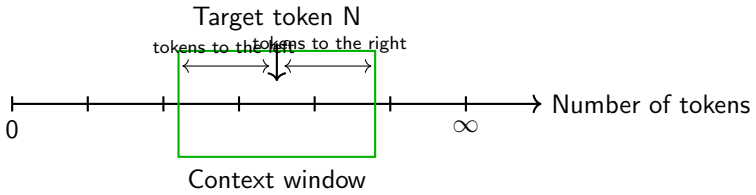

Why Bigger Context Windows Aren't Always Better?

1

Information Overload

- Excess data can overwhelm the LLM.
- Reduces focus and slows performance.

2

Getting Lost in Data

- Large windows prioritize edges,
- Missing key middle info.

3

Poor Management

- More data leads to noise,
- Redundancy, and potential bias.

4

Long-Range Issues

- Understanding relationships

becomes harder with large context windows.

Context Window Problems

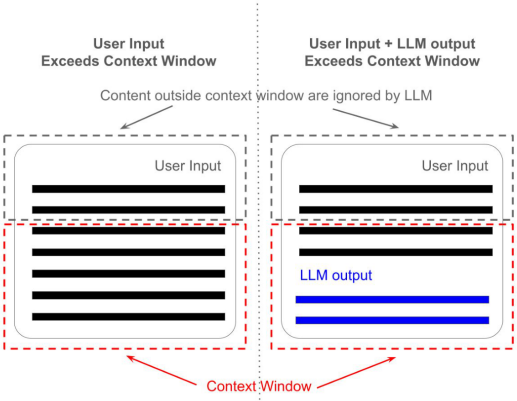

Moving Towards Larger Context Windows

- **Efficient Attention Variants:**
  - Longformer, BigBird, FlashAttention-2
- **Memory-Augmented Models:**
  - Retrieval-augmented generation (RAG)
  - Memory layers (e.g., RetNet)
- **Chunking + Re-ranking:**
  - Process in parts, summarize/join results
- **Recurrence or State Passing:**
  - Transformer-XL, RWKV (RNN-inspired)

Moving Towards Larger Context Windows

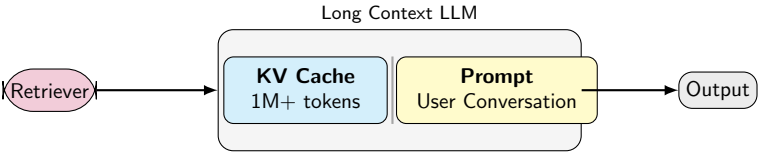

Limitations of LLMs

- High inference costs & latency
- Bias and toxic outputs
- Hallucination and factual inconsistency
- Limited interpretability
- Require vast pre-training data
- Poor domain generalization (medical, legal, etc.)

Future Directions for LLMs

- **Long Context Handling:** Efficient token reuse, sparse attention, recurrence.
- **Token-Free Models:** Improved byte/character models.
- **Multimodal LLMs:** Integrating vision, speech, and more.
- **Smaller, Efficient LLMs:** Distillation, quantization, sparse models.
- **Open-Weight & Ethical LLMs:** Responsible, accessible models.

References

- These slides have been adapted from:
  - Bhiksha Raj & Rita Singh, 11-785 Introduction to Deep Learning, CMU

References

1. Kaplan et al., “Scaling Laws for Neural Language Models”, 2020.
2. Devlin et al., “BERT: Pre-training of Deep Bidirectional Transformers”, 2018.
3. Brown et al., “Language Models are Few-Shot Learners (GPT-3)”, 2020.
4. Xue et al., “ByT5: Towards a token-free future with pre-trained byte-to-byte models”, 2022.
5. Petrov, La Malfa, Torr, Bibi, “Language Model Tokenizers Introduce Unfairness Between Languages”, NeurIPS 2023.
6. Beltagy et al., “Longformer: The Long-Document Transformer”, 2020.
7. Tay et al., “Efficient Transformers: A Survey”, 2020.
8. OpenAI, “Technical Report on GPT-4”, 2023.

References

8. Google Research Blog: Scaling Transformer Models.
9. Anthropic, “Claude 3.5 Release Notes”, 2024.
10. FlashAttention: https://arxiv.org/abs/2205.14135