

Recent Advancements in Natural Language Processing

Naeemullah Khan

naeemullah.khan@kaust.edu.sa




جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

July 16, 2026

The History of NLP



1. Motivation
2. Learning Outcomes
3. The LLaMA 3 Herd of Models
4. Byte Latent Transformer (BLT)
5. Large Language Diffusion Models (LLDMs)
6. Summary
7. References

Three departures from the standard recipe of an autoregressive transformer over subword tokens:

- ▶ **Open-weight LLMs:** Releases like the LLaMA 3 herd bring frontier-level models into the open, widening access to research and deployment.
- ▶ **Tokenizer-free modeling:** the Byte Latent Transformer (BLT) replaces subword tokenization with learned byte patches.
- ▶ **Non-autoregressive generation:** Large Language Diffusion Models (LLDMs) generate text by iterative denoising instead of left-to-right prediction.

By the end of this session, you should be able to:

- ▶ Describe the LLaMA 3 herd of models and what open-weight frontier models enable.
- ▶ Explain the Byte Latent Transformer (BLT): why tokenization is a bottleneck, and how entropy-based byte patching replaces it.
- ▶ Summarize Large Language Diffusion Models (LLDMs) and how diffusion-based text generation differs from autoregressive decoding.
- ▶ Situate these emerging paradigms relative to the standard transformer recipe and identify the trade-offs each one makes.

The LLaMA 3 Herd of Models



The Llama 3 Herd of Models

Llama Team, AI @ Meta¹

¹A detailed contributor list can be found in the appendix of this paper.

Modern artificial intelligence (AI) systems are powered by foundation models. This paper presents a new set of foundation models, called Llama 3. It is a herd of language models that natively support multilinguality, coding, reasoning, and tool usage. Our largest model is a dense Transformer with 405B parameters and a context window of up to 128K tokens. This paper presents an extensive empirical evaluation of Llama 3. We find that Llama 3 delivers comparable quality to leading language models such as GPT-4 on a plethora of tasks. We publicly release Llama 3, including pre-trained and post-trained versions of the 405B parameter language model and our Llama Guard 3 model for input and output safety. The paper also presents the results of experiments in which we integrate image, video, and speech capabilities into Llama 3 via a compositional approach. We observe this approach performs competitively with the state-of-the-art on image, video, and speech recognition tasks. The resulting models are not yet being broadly released as they are still under development.

Date: July 23, 2024

Website: <https://llama.meta.com/>

Meta AI (2024)

- ▶ **Model Sizes:** 8B, 70B, and 405B
- ▶ **Open Weights:** Available for commercial use
- ▶ **Training:** Over 15 trillion tokens, multilingual data, improved alignment techniques

	Finetuned	Multilingual	Long context	Tool use	Release
Llama 3 8B	✗	✗ ¹	✗	✗	April 2024
Llama 3 8B Instruct	✓	✗	✗	✗	April 2024
Llama 3 70B	✗	✗ ¹	✗	✗	April 2024
Llama 3 70B Instruct	✓	✗	✗	✗	April 2024
Llama 3.1 8B	✗	✓	✓	✗	July 2024
Llama 3.1 8B Instruct	✓	✓	✓	✓	July 2024
Llama 3.1 70B	✗	✓	✓	✗	July 2024
Llama 3.1 70B Instruct	✓	✓	✓	✓	July 2024
Llama 3.1 405B	✗	✓	✓	✗	July 2024
Llama 3.1 405B Instruct	✓	✓	✓	✓	July 2024

Overview of the Llama 3 Herd of models. All results in this paper are for the Llama 3.1 models.

Key Features:

- ▶ 128k context window for extended input sequences
- ▶ FlashAttention-2 and grouped-query attention for efficient inference
- ▶ Fine-tuned variants: Instruct and Chat models

The LLaMA 3 Herd of Models (cont.)

Category	Benchmark	Llama 3 8B	Gemma 2 9B	Mistral 7B	Llama 3 70B	Mixtral 8x22B	GPT 3.5 Turbo	Llama 3 405B	Nemotron 4 340B	GPT-4 (o1us)	GPT-4o	Claude 3.5 Sonnet
General	MMLU (5-shot)	69.4	72.3	61.1	83.6	76.9	70.7	87.3	82.6	85.1	89.1	89.9
	MMLU (0-shot, CoT)	73.0	72.3 [△]	60.5	86.0	79.9	69.8	88.6	78.7 [◊]	85.4	88.7	88.3
	MMLU-Pro (5-shot, CoT)	48.3	–	36.9	66.4	56.3	49.2	73.3	62.7	64.8	74.0	77.0
	IFEval	80.4	73.6	57.6	87.5	72.7	69.9	88.6	85.1	84.3	85.6	88.0
Code	HumanEval (0-shot)	72.6	54.3	40.2	80.5	75.6	68.0	89.0	73.2	86.6	90.2	92.0
	MBPP EvalPlus (0-shot)	72.8	71.7	49.5	86.0	78.6	82.0	88.6	72.8	83.6	87.8	90.5
Math	GSM8K (8-shot, CoT)	84.5	76.7	53.2	95.1	88.2	81.6	96.8	92.3 [◊]	94.2	96.1	96.4 [◊]
	MATH (0-shot, CoT)	51.9	44.3	13.0	68.0	54.1	43.1	73.8	41.1	64.5	76.6	71.1
Reasoning	ARC Challenge (0-shot)	83.4	87.6	74.2	94.8	88.7	83.7	96.9	94.6	96.4	96.7	96.7
	GPQA (0-shot, CoT)	32.8	–	28.8	46.7	33.3	30.8	51.1	–	41.4	53.6	59.4
Tool use	BFCL	76.1	–	60.4	84.8	–	85.9	88.5	86.5	88.3	80.5	90.2
	Nexus	38.5	30.0	24.7	56.7	48.5	37.2	58.7	–	50.3	56.1	45.7
Long context	ZeroSCROLLS/QuALITY	81.0	–	–	90.5	–	–	95.2	–	95.2	90.5	90.5
	InfiniteBench/En.MC	65.1	–	–	78.2	–	–	83.4	–	72.1	82.5	–
	NIH/Multi-needle	98.8	–	–	97.5	–	–	98.1	–	100.0	100.0	90.8
Multilingual	MGSM (0-shot, CoT)	68.9	53.2	29.9	86.9	71.1	51.4	91.6	–	85.9	90.5	91.6

Table 2 Performance of finetuned Llama 3 models on key benchmark evaluations. The table compares the performance of the 8B, 70B, and 405B versions of Llama 3 with that of competing models. We **boldface** the best-performing model in each of three model-size equivalence classes. [△]Results obtained using 5-shot prompting (no CoT). [◊]Results obtained without CoT. [◊]Results obtained using zero-shot prompting.

The LLaMA 3 Herd of Models (cont.)




Figure 1 Illustration of the overall architecture and training of Llama 3. Llama 3 is a Transformer language model trained to predict the next token of a textual sequence. See text for details.

The LLaMA 3 Herd of Models (cont.)



Figure 5 Illustration of 4D parallelism. GPUs are divided into parallelism groups in the order of [TP, CP, PP, DP], where DP stands for FSDP. In this example, 16 GPUs are configured with a group size of $|TP|=2$, $|CP|=2$, $|PP|=2$, and $|DP|=2$. A GPU's position in 4D parallelism is represented as a vector, $[D_1, D_2, D_3, D_4]$, where D_i is the index on the i -th parallelism dimension. In this example, GPU0[TP0, CP0, PP0, DP0] and GPU1[TP1, CP0, PP0, DP0] are in the same TP group, GPU0 and GPU2 are in the same CP group, GPU0 and GPU4 are in the same PP group, and GPU0 and GPU8 are in the same DP group.

The LLaMA 3 Herd of Models (cont.)

User prompt

Based on the inflation rate for the US in February 2024, and the cost of this list of tuples containing my daily groceries, [(("eggs",3),("chicken",4),("rice",2),("beans",2.50),("apples",6.25))], can you give me the estimated total cost of my groceries over a 30 day month when you add 1 cookie per day at Costco's food court giant cookie debut price and adjust the total for the inflation percentage in \$5.00?

System Prompt

Environment: `ipython`,
Tools: `brave_search`, `wolfram_alpha`
Cutting Knowledge Date: `December 2023`
Today's Date: `Jul 11, 2024`

Assistant response

To answer your question, I will need to:

- Look up the inflation rate for the US in February 2024
- Look up the Costco food court giant cookie debut price
- Calculate the estimated total cost of your groceries over a 30 day month when you add 1 cookie per day at Costco's food court giant cookie debut price and adjust the total for the inflation percentage

First, I will look up the inflation rate for the US in February 2024.

Assistant tool call

`brave_search.call(query="US Inflation rate February 2024")`

Tool output

```
{
  "title": "US Inflation Remained Higher Than Economists Had Expected in February",
  "url": "https://www.investopedia.com/february-2024-cpi-8697820",
  "description": "Consumer prices rose 3.2% (strong) over the year in February, higher than the 3.1% annual rate in January and above what economists had predicted. An uptick in gas prices and housing prices was behind the unexpectedly high inflation rate."
}
```

The inflation rate for the US in February 2024 was 3.2%. Next, I will look up the Costco food court giant cookie debut price.

Assistant tool call

`brave_search.call(query="Costco food court giant cookie debut price")`

Tool output

```
{
  "title": "Costco's food court replaced its churro with a giant cookie. Here's how it tastes",
  "url": "https://www.today.com/food/groceries/costco-double-chocolate-chunk-cookie-food-court-review-rcn33537",
  "description": "Costco's Kirkland Signature Double Chocolate Chunk (strong) Cookie (strong), the latest addition to its (strong) Food (strong) Court (strong), costs $2.49 and has 750 calories. Here's my review of the buttery treat."
}
```

The Costco food court giant cookie debut price is \$2.49. Now, I will calculate the estimated total cost of your groceries over a 30 day month when you add 1 cookie per day at Costco's food court giant cookie debut price and adjust the total for the inflation percentage.

Assistant tool call

```
# Define the list of tuples containing the daily groceries
groceries = [(("eggs",3),("chicken",4),("rice",2),("beans",2.50),("apples",6.25))]
```

```
# Calculate the total cost of the groceries per day
total_cost_per_day = sum([item[1] for item in groceries])
```

```
# Calculate the total cost of the groceries for 30 days including the cookie
total_cost_for_30_days_with_cookie = total_cost_per_day_with_cookie * 30
```

```
# Define the inflation rate
inflation_rate = 0.832
```

```
# Calculate the total cost of the groceries for 30 days including the cookie and adjusted for inflation
total_cost_for_30_days_with_cookie_and_inflation = total_cost_for_30_days_with_cookie * (1 + inflation_rate)
```

```
# Print the result
print(f"The estimated total cost of your groceries over a 30 day month when you add 1 cookie per day at Costco's food court giant cookie debut price and adjust the total for the inflation percentage is ${total_cost_for_30_days_with_cookie_and_inflation:.2f}")
```

Tool output

The estimated total cost of your groceries over a 30 day month when you add 1 cookie per day at Costco's food court giant cookie debut price and adjust the total for the inflation percentage is \$626.63.

Figure 10 Multi-step tool usage. Example of Llama 3 performing multi-step planning, reasoning, and tool calling to solve a task.

KAUST Academy

Recent Advancements in NLP

14

	Reading Comprehension		
	SQuAD	QuAC	RACE
Llama 3 8B	77.0 $\pm$ 0.8	44.9 $\pm$ 1.1	54.3 $\pm$ 1.4
Mistral 7B	73.2 $\pm$ 0.8	44.7 $\pm$ 1.1	53.0 $\pm$ 1.4
Gemma 7B	81.8 $\pm$ 0.7	42.4 $\pm$ 1.1	48.8 $\pm$ 1.4
Llama 3 70B	81.8 $\pm$ 0.7	51.1 $\pm$ 1.1	59.0 $\pm$ 1.4
Mixtral 8 $\times$ 22B	84.1 $\pm$ 0.7	44.9 $\pm$ 1.1	59.2 $\pm$ 1.4
Llama 3 405B	81.8 $\pm$ 0.7	53.6 $\pm$ 1.1	58.1 $\pm$ 1.4
GPT-4	—	—	—
Nemotron 4 340B	—	—	—
Gemini Ultra	—	—	—

Table 9 Pre-trained model performance on reading comprehension tasks. Results include 95% confidence intervals.

	Code	
	HumanEval	MBPP
Llama 3 8B	37.2 $\pm$ 7.4	47.6 $\pm$ 4.4
Mistral 7B	30.5 $\pm$ 7.0	47.5 $\pm$ 4.4
Gemma 7B	32.3 $\pm$ 7.2	44.4 $\pm$ 4.4
Llama 3 70B	58.5 $\pm$ 7.5	66.2 $\pm$ 4.1
Mixtral 8 $\times$ 22B	45.1 $\pm$ 7.6	71.2 $\pm$ 4.0
Llama 3 405B	61.0 $\pm$ 7.5	73.4 $\pm$ 3.9
GPT-4	67.0 $\pm$ 7.2	—
Nemotron 4 340B	57.3 $\pm$ 7.6	—
Gemini Ultra	74.4 $\pm$ 6.7	—

Table 10 Pre-trained model performance on coding tasks. Results include 95% confidence intervals.

Highlights:

- ▶ Competitive performance with GPT-4 and Gemini
- ▶ Foundation for Code LLaMA, Audio LLaMA, and multimodal LLaMA models

Byte Latent Transformer (BLT)

Byte Latent Transformer: Patches Scale Better Than Tokens

Artidoro Pagnoni[‡], Ram Pasunuru[‡], Pedro Rodriguez[‡], John Nguyen[‡], Benjamin Muller, Margaret Li^{1,◊},
Chunting Zhou[◊], Lili Yu, Jason Weston, Luke Zettlemoyer, Gargi Ghosh, Mike Lewis, Ari Holtzman^{‡,2,◊},
Srinivasan Iyer[‡]

FAIR at Meta, ¹Paul G. Allen School of Computer Science & Engineering, University of Washington,

²University of Chicago

[‡]Joint second author, [‡]Joint last author, [◊]Work done at Meta

We introduce the Byte Latent Transformer (BLT), a new byte-level LLM architecture that, for the first time, matches tokenization-based LLM performance at scale with significant improvements in inference efficiency and robustness. BLT encodes bytes into dynamically sized patches, which serve as the primary units of computation. Patches are segmented based on the entropy of the next byte, allocating more compute and model capacity where increased data complexity demands it. We present the first FLOP controlled scaling study of byte-level models up to 8B parameters and 4T training bytes. Our results demonstrate the feasibility of scaling models trained on raw bytes without a fixed vocabulary. Both training and inference efficiency improve due to dynamically selecting long patches when data is predictable, along with qualitative improvements on reasoning and long tail generalization. Overall, for fixed inference costs, BLT shows significantly better scaling than tokenization-based models, by simultaneously growing both patch and model size.

Date: December 16, 2024

Correspondence: artidoro at cs.washington.edu, sviyer at meta.com

Code: <https://github.com/facebookresearch/blt>



BLT is a compression-based approach for efficient LLMs:

- ▶ Compress input text into a smaller latent space
 - ▶ Run transformer on latent tokens
 - ▶ Decode with a decoder-only model
-
- ✓ Reduces memory & compute
 - ✓ Maintains quality (BLEU, perplexity)
 - ✓ Works well with long context (128k+)

Related: VQ-VAE, tokenizer-free models, compression transformers **Meta AI (2024)** <https://ai.meta.com/research/publications/byte-latent-transformer-patches-scale-better-than-tokens/>

Byte Latent Transformer (BLT) (cont.)




Figure 1 Scaling trends for fixed inference FLOP models (fully) trained with varying training budgets. In token-based models, a fixed inference budget determines the model size. In contrast, the BLT architecture provides a new scaling axis allowing simultaneous increases in model and patch size while keeping the same training and inference budget. BLT patch-size (ps) 6 and 8 models quickly overtake scaling trends of BPE Llama 2 and 3. Moving to the larger inference budget makes the larger patch size 8 model more desirable sooner. Both BPE compute-optimal point and crossover point are indicated with vertical lines.

Byte Latent Transformer (BLT) (cont.)




Figure 2 BLT comprises three modules, a lightweight *Local Encoder* that encodes input bytes into patch representations, a computationally expensive *Latent Transformer* over patch representations, and a lightweight *Local Decoder* to decode the next patch of bytes. BLT incorporates byte n -gram embeddings and a cross-attention mechanism to maximize information flow between the *Latent Transformer* and the byte-level modules (Figure 5). Unlike fixed-vocabulary tokenization, BLT dynamically groups bytes into patches preserving access to the byte-level information.




Figure 3 Patching schemes group bytes in different ways, each leading to a different number of resulting patches. Since each patch is processed using a large transformer step, the number of patches directly determines the bulk of the compute expended in terms of FLOPs. These schemes group bytes into patches by (a) striding every four bytes (§2.1) as in MegaByte (Yu et al., 2023), (b) tokenizing with Byte-Pair Encoding (BPE), in this case the Llama-3 (Dubey et al., 2024) tokenizer, (c & d) entropy-based patching as in this work (§2.3), (e) patching on space-bytes (Slagle, 2024), (f) and patching on entropy using a small CNN byte-level model with 2-byte context.




Figure 4 This figure plots the entropy $H(x_i)$ of each byte in “Daenerys Targaryen is in Game of Thrones, a fantasy epic by George R.R. Martin.” with spaces shown as underscores. Patches end when $H(x_i)$ exceeds the global threshold θ_g , shown as a red horizontal line. The start of new patches are shown with vertical gray lines. For example, the entropies of “G” and “e” in “George R.R. Martin” exceed θ_g , so “G” is the start of a single byte patch and “e” of a larger patch extending to the end of the named entity as the entropy $H(x_i)$ stays low, resulting in no additional patches.

Space Patching (Slagle, 2024) improves over strided patching by creating new patches after any space-like byte, which often aligns with word boundaries in many languages.

- ▶ Allocates a latent transformer step for every word.
- ▶ Ensures consistent patching of words across sequences.
- ▶ Focuses compute (flops) on hard predictions, typically following spaces.

Example: Predicting the first byte of the answer to “Who composed the Magic Flute?” is harder than predicting subsequent bytes in “Mozart”.

- ▶ The first character narrows down possible completions, making the rest easier to predict.

Limitations:

- ▶ Not suitable for all languages or domains.
- ▶ Cannot vary patch size.

Next: A new patching method that leverages the difficulty of predicting initial bytes in words and enables control over patch size.

Entropy Patching: Using Next-Byte Entropies from a Small Byte LM

- ▶ Instead of rule-based heuristics (e.g., whitespace), entropy patching uses a data-driven approach to identify high-uncertainty next-byte predictions.
- ▶ Train a small byte-level auto-regressive LM on BLT training data.
- ▶ Compute next-byte entropies $H(x_i)$ under the LM distribution p_e over the byte vocabulary V :

$$H(x_i) = - \sum_{v \in V} p_e(x_i = v \mid x_{<i}) \log p_e(x_i = v \mid x_{<i}) \quad (1)$$

- ▶ Two methods for identifying patch boundaries:
 1. **Global Constraint:** $H(x_t) > \theta_g$ (entropy above a global threshold)
 2. **Approx. Monotonic Constraint:** $H(x_t) - H(x_{t-1}) > \theta_r$ (entropy rises relative to previous byte)
- ▶ The second method finds points that break approximate monotonically decreasing entropy within a patch.
- ▶ Patch boundaries are identified during lightweight preprocessing at dataloading.
- ▶ Differs from Nawrot et al. (2023), which trains a classifier to predict entropy-based patch boundaries.

In experiments: Both methods are compared for distinguishing between low and high entropy bytes.

Encoder Multi-Headed Cross-Attention:

- ▶ Follows Perceiver input cross-attention (Jaegle et al., 2021), but latent representations correspond to variable patch representations (not fixed set).
- ▶ Each patch p_j attends only to its own bytes.
- ▶ Query vector for each patch is initialized by pooling byte representations for p_j , then projecting via $EC \in \mathbb{R}^{h_E \times (h_E \times U_E)}$, where U_E is the number of encoder cross-attention heads.

Formulation:

$$P_{0,j} = EC(f_{\text{bytes}}(p_j)) \quad (\text{Pooling function } f \text{ over bytes in patch } p_j) \quad (2)$$

$$P_l = P_{l-1} + W_o \left[\text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V \right] \quad (3)$$

$$Q_j = W_q(P_{l-1,j}) \quad (4)$$

$$K_i = W_k(h_{l-1,i}) \quad (5)$$

$$V_i = W_v(h_{l-1,i}) \quad (6)$$

$$h_l = \text{Encoder-Transformer-Layer}_l(h_{l-1}) \quad (7)$$

- ▶ $P \in \mathbb{R}^{n_p \times h_G}$: n_p patch representations for the global model.
- ▶ W_q, W_k, W_v, W_o : learned projection matrices.

- ▶ Patch representations are initialized by pooling byte embeddings e_i for each patch p_j .

Key Difference: Latent representations are variable and correspond to patch boundaries, not a fixed latent set.

Encoder Multi-Headed Cross-Attention (cont.)




Figure 5 The local encoder uses a cross-attention block with patch representations as queries, and byte representations as keys/values to encode byte representations into patch representations. The local decoder uses a similar block but with the roles reversed i.e. byte representations are now the queries and patch representations are the keys/values. Here we use Cross-Attn $k = 2$.

Decoder Multi-Headed Cross-Attention:

- ▶ In the decoder, the roles of queries and key/values are swapped compared to the encoder cross-attention.
- ▶ Byte representations act as queries, while patch representations serve as keys and values.
- ▶ Initial byte representations for cross-attention are set as the byte embeddings from the last encoder layer, h^E .

Formulation:

$$D_0 = h^E \quad (8)$$

$$B_l = D_{l-1} + W_o \left[\text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V \right] \quad (9)$$

$$Q_i = W_q(d_{l-1,i}) \quad (10)$$

$$K_j = W_k(DC(o_j)) \quad (11)$$

$$V_j = W_v(DC(o_j)) \quad (12)$$

$$D_l = \text{Decoder-Transformer-Layer}_l(B_l) \quad (13)$$

- ▶ DC is a linear transformation and split operation applied to the final patch representations o_j from the global model.
- ▶ W_q, W_k, W_v, W_o are learned projection matrices.

- ▶ $B \in \mathbb{R}^{h_D \times n_b}$, where n_b is the number of output bytes.
- ▶ Multiple attention heads, pre-LayerNorm, no positional embeddings, and residual connections are used.

Summary: Decoder cross-attention enables each output byte to attend to patch-level latent representations, facilitating efficient and flexible decoding.

BLT Performance:

- ▶ BLT achieves state-of-the-art performance on long-context benchmarks.
- ▶ Outperforms existing LLMs in terms of perplexity and BLEU scores.
- ▶ Efficiently handles long sequences (128k+ tokens) with reduced memory and compute requirements.

Language	Language → English		English → Language	
	Llama 3	BLT	Llama 3	BLT
Arabic	22.3	24.6	10.4	8.8
German	41.3	42.0	29.8	31.2
Hindi	20.7	20.9	7.8	7.2
Italian	34.0	33.9	24.4	26.2
Vietnamese	31.2	31.0	28.4	23.7
Thai	17.9	18.1	10.5	7.7
Armenian	1.7	6.3	0.6	0.9
Amharic	1.3	3.1	0.4	0.5
Assamese	2.7	5.4	0.8	1.6
Bengali	4.7	12.7	1.7	4.1
Bosnian	36.0	37.3	16.9	19.6
Cebuano	18.2	20.6	5.8	9.1
Georgian	1.7	7.4	1.0	2.5
Gujarati	2.0	5.8	1.0	2.2
Hausa	5.75	5.9	1.2	1.3
Icelandic	16.1	17.9	4.8	5.3
Kannada	1.6	3.9	0.7	1.7
Kazakh	5.6	7.0	1.0	2.6
Kabuverdianu	20.3	20.9	5.1	6.8
Khmer	4.4	9.5	0.8	0.8
Kyrgyz	4.6	5.1	0.9	2.0
Malayalam	1.8	3.5	0.7	1.4
Odia	1.6	2.7	0.8	1.1
Somali	5.0	5.0	1.1	1.4
Swahili	10.1	12.0	1.4	2.3
Urdu	9.3	9.5	2.0	1.4
Zulu	4.7	5.0	0.6	0.5
Overall Average	12.1	14.0	5.9	6.4

Table 4 Performance of 8B BLT and 8B Llama 3 trained for 1T tokens on translating into and from six widely-used languages and twenty one lower resource languages with various scripts from the FLORES-101 benchmark (Goyal et al., 2022).

Task	Prompt	Llama 3	BLT
Substitute Word	Question: Substitute " and " with " internet " in " She went to the kitchen and saw two cereals. ". Answer:	She went to the kitchen and saw two cereals.	She went to the kitchen internet saw two cereals.
Swap Char	Question: Swap " h " and " a " in " that ". Answer:	that	taht
Substitute Char	Question: Substitute " a " with " m " in " page ". Answer:	-	pmge
Semantic Similarity	Question: More semantically related to " are ": " seem ", " acre ". Answer:	acre	seem
Orthographic Similarity	Question: Closer in Levenshtein distance to " time ": " timber ", " period ". Answer:	period	timber
Insert Char	Question: Add an " z " after every " n " in " not ". Answer:	znatz	nzot

Figure 7 Output responses from Llama 3 and BLT models for various tasks from CUTE benchmark. BLT model performs better on sequence manipulation tasks compared to the tokenizer-based Llama 3 model. Note that few-shot examples are not shown in the above prompts to maintain clarity.

Large Language Diffusion Models (LLDMs)

Large Language Diffusion Models

Shen Nie^{1,†} Fengqi Zhu^{1,†} Zebin You^{1,†} Xiaolu Zhang^{2,‡} Jingyang Ou¹ Jun Hu^{2,‡} Jun Zhou²
Yankai Lin^{1,‡} Ji-Rong Wen¹ Chongxuan Li^{1,‡,¶}

Abstract

Autoregressive models (ARMs) are widely regarded as the cornerstone of large language models (LLMs). We challenge this notion by introducing **LLaDA**, a diffusion model trained from scratch under the pre-training and supervised fine-tuning (SFT) paradigm. LLaDA models distributions through a forward data masking process and a reverse process, parameterized by a vanilla Transformer to predict masked tokens. By optimizing a likelihood bound, it provides a principled generative approach for probabilistic inference. Across extensive benchmarks, LLaDA demonstrates strong *scalability*, outperforming our self-constructed ARM baselines. Remarkably, LLaDA 8B is competitive with strong LLMs like LLaMA3 8B in *in-context learning* and, after SFT, exhibits impressive *instruction-following* abilities in case studies such as multi-turn dialogue. Moreover, LLaDA addresses the reversal curse, surpassing GPT-4o in a reversal poem completion task. Our findings establish diffusion models as a viable and promising alternative to ARMs, challenging the assumption that key LLM capabilities discussed above are inherently tied to ARMs. Project page and codes: <https://ml-gsai.github.io/LLaDA-demo/>.




Figure 1. **Zero/Few-Shot Benchmarks.** We scale LLaDA to an unprecedented size of 8B parameters from scratch, achieving competitive performance with strong LLMs (Dubey et al., 2024).

distribution $p_{\text{data}}(\cdot)$ by optimizing a model distribution $p_{\theta}(\cdot)$ through maximum likelihood estimation, or equivalently KL divergence minimization between the two distributions:

$$\underbrace{\max_{\theta} \mathbb{E}_{p_{\text{data}}(x)} \log p_{\theta}(x) \Leftrightarrow \min_{\theta} \text{KL}(p_{\text{data}}(x) || p_{\theta}(x))}_{\text{Generative modeling principles}}. \quad (1)$$

- ▶ **Emerging trend:** Use diffusion instead of autoregression for text generation

Key Idea:

- ▶ Generate latent text representation over diffusion steps
- ▶ Use decoder to map latent $\rightarrow$ text

Benefits:

- ▶ Better global coherence
- ▶ Parallel generation
- ▶ Less repetition, hallucination

Related work:

- ▶ Diffusion-LM (Li et al.)
- ▶ Masked-Diffusion Transformers
- ▶ UniDiffuser (Vision + Text)

Note: Still experimental, but promising for the future of text generation

Large Language Diffusion Models




Figure 2. A Conceptual Overview of LLaDA. (a) Pre-training. LLaDA is trained on text with random masks applied independently to all tokens at the same ratio $t \sim U[0, 1]$. (b) SFT. Only response tokens are possibly masked. (c) Sampling. LLaDA simulates a diffusion process from $t = 1$ (fully masked) to $t = 0$ (unmasked), predicting all masks simultaneously at each step with flexible remask strategies.

Large Language Diffusion Models




Figure 3. Scalability of LLaDA. We evaluate the performance of LLaDA and our ARM baselines trained on the same data across increasing computational FLOPs. LLaDA exhibits strong scalability, matching the overall performance of ARMs on six tasks.

Large Language Diffusion Models

Table 1. Benchmark Results of Pre-trained LLMs. * indicates that LLaDA 8B Base, LLaMA2 7B Base, and LLaMA3 8B Base are evaluated under the same protocol, detailed in Appendix B.5. Results indicated by [†] and [¶] are sourced from Chu et al. (2024); Yang et al. (2024) and Bi et al. (2024) respectively. The numbers in parentheses represent the number of shots used for evaluation. “-” indicates unknown data.

	LLaDA 8B*	LLaMA3 8B*	LLaMA2 7B*	Qwen2 7B [†]	Qwen2.5 7B [†]	Mistral 7B [†]	Deepseek 7B [¶]
Model	Diffusion	AR	AR	AR	AR	AR	AR
Training tokens	2.3T	15T	2T	7T	18T	-	2T
General Tasks							
MMLU	65.9 (5)	65.4 (5)	45.9 (5)	70.3 (5)	74.2 (5)	64.2 (5)	48.2 (5)
BBH	49.8 (3)	57.6 (3)	37.3 (3)	62.3 (3)	70.4 (3)	56.1 (3)	39.5 (3)
ARC-C	47.9 (0)	53.1 (0)	46.3 (0)	60.6 (25)	63.7 (25)	60.0 (25)	48.1 (0)
Hellaswag	72.5 (0)	79.1 (0)	76.0 (0)	80.7 (10)	80.2 (10)	83.3 (10)	75.4 (0)
TruthfulQA	46.4 (0)	44.0 (0)	39.0 (0)	54.2 (0)	56.4 (0)	42.2 (0)	-
WinoGrande	74.8 (5)	77.3 (5)	72.5 (5)	77.0 (5)	75.9 (5)	78.4 (5)	70.5 (0)
PIQA	74.4 (0)	80.6 (0)	79.1 (0)	-	-	-	79.2 (0)
Mathematics & Science							
GSM8K	70.7 (4)	53.1 (4)	14.3 (4)	80.2 (4)	85.4 (4)	36.2 (4)	17.4 (8)
Math	27.3 (4)	15.1 (4)	3.2 (4)	43.5 (4)	49.8 (4)	10.2 (4)	6.0 (4)
GPQA	26.1 (5)	25.9 (5)	25.7 (5)	30.8 (5)	36.4 (5)	24.7 (5)	-
Code							
HumanEval	33.5 (0)	34.2 (0)	12.8 (0)	51.2 (0)	57.9 (0)	29.3 (0)	26.2 (0)
HumanEval-FIM	73.8 (2)	73.3 (2)	26.9 (2)	-	-	-	-
MBPP	38.2 (4)	47.4 (4)	18.4 (4)	64.2 (0)	74.9 (0)	51.1 (0)	39.0 (3)
Chinese							
CMMLU	69.9 (5)	50.7 (5)	32.5 (5)	83.9 (5)	-	-	47.2 (5)
C-Eval	70.5 (5)	51.7 (5)	34.0 (5)	83.2 (5)	-	-	45.0 (5)

Table 4. **Visualization of the Sampling Process and a Generated Multi-round Dialogue.** In the response of LLaDA, darker colors indicate tokens predicted in the later stages of sampling, while lighter colors correspond to earlier predictions.

<i>Sampling Process</i>	
User	Lily can run 12 kilometers per hour for 4 hours. After that, she runs 6 kilometers per hour. How many kilometers can she run in 8 hours?
LLaDA	In 4 hours , Lily runs 12 * 4 = 48 kilometers . After 4 hours she runs 6 * 4 = 24 kilometers . In total she runs 48 + 24 = 72 kilometers in 8 hours . Result 72
<i>Multi-round Dialogue Case</i>	
User	Could you please share with me what the initial two lines of the renowned poem "The Road Not Taken"?
LLaDA	Certainly! The initial two lines of the renowned poem "The Road Not Taken" by Robert Frost are: "Two roads diverged in a yellow wood, And sorry I could not travel both." These lines set the stage for the poem, introducing the idea of a speaker facing a choice between two paths.
User	Please help me translate into Chinese.
LLaDA	Sure! The first two lines of "The Road Not Taken" by Robert Frost can be translated into Chinese as: "两条路分岔在黄色的树林中，遗憾我不能同时走"
User	Please translate into German.
LLaDA	Sure! The first two lines of "The Road Not Taken" by Robert Frost can be translated into German as: "Zwei Wege trennten sich im gelben Wald, und leider konnte ich nicht beide tun."
User	Please also write a poem about life choices. I require 5 sentences, and the first word of each sentence starts with C as the first letter.
LLaDA	Certainly! Here's a poem about life choices, with each sentence starting with the letter "C": Choices cascade through life's journey, Creating connections with each decision, Careful consideration leads to better paths, Courageous choices shape our destiny, Choosing wisely guides us through the years.

7: Return p_θ

Algorithm 2 Supervised Fine-Tuning of LLaDA

Require: mask predictor p_θ , pair data distribution p_{data}

1: **repeat**

2: $p_0, r_0 \sim p_{\text{data}}, t \sim \text{U}(0, 1]$ # please refer to Appendix B.1 for details on the SFT data processing.

3: $r_t \sim q_{t|0}(r_t|r_0)$ # $q_{t|0}$ is defined in Eq. (7)

4: Calculate $\mathcal{L} = -\frac{1}{t \cdot L'} \sum_{i=1}^{L'} \mathbf{1}[r_t^i = \text{M}] \log p_\theta(r_0^i | p_0, r_t)$ # L' is the sequence length of r_0

5: Calculate $\nabla_\theta \mathcal{L}$ and run optimizer.

6: **until** Converged

7: **Return** p_θ

Algorithm 3 Conditional Log-likelihood Evaluation of LLaDA

Require: mask predictor p_θ , prompt p_0 , response r_0 , the number of Monte Carlo estimations n_{mc}

```
1: log_likelihood = 0
2: for  $i \leftarrow 1$  to  $n_{mc}$  do
3:    $l \sim \{1, 2, \dots, L\}$  #  $L$  is the sequence length of  $r_0$ 
4:   Obtain  $r_l$  by uniformly sampling  $l$  tokens from  $r_0$  without replacement for masking
5:    $\text{log\_likelihood} = \text{log\_likelihood} + \frac{L}{l} \sum_{i=1}^L \mathbf{1}[r_l^i = \text{M}] \log p_\theta(r_0^i | p_0, r_l)$ 
6: end for
7:  $\text{log\_likelihood} = \text{log\_likelihood} / n_{mc}$ 
8: Return log_likelihood
```

Summary

Advancement	Domain	Key Benefit
LLaMA 3	NLP	Open frontier-level models
BLT	NLP/LLM	Tokenizer-free byte-level modeling
LLDMs	NLP	Diffusion-based text generation

- ▶ Multimodal Foundation Models
- ▶ Efficient Training (QLoRA, MoE, FlashAttention)
- ▶ LLMs with Memory, Tools, Agents
- ▶ Transformer Alternatives (MLPs, State Space Models)
- ▶ Unified Models: Language, Vision, Audio, Action

Integration of these advances will define AI agents, robotics, and autonomous systems.

- ▶ **LLaMA 3**: Open weights at frontier quality widen access to research and deployment.
- ▶ **BLT**: Learned byte patches remove the tokenizer, trading fixed vocabulary for dynamic, entropy-driven compute allocation.
- ▶ **LLDMs**: Diffusion-based generation is an alternative to left-to-right decoding, enabling parallel refinement of whole sequences.
- ▶ Each paradigm removes a different fixed assumption of the standard transformer recipe.

References

- ▶ **SEEM**: <https://segment-anything.com/seem>
- ▶ **FlashAttention-2**: Dao et al., 2023 ([arXiv:2307.08691](https://arxiv.org/abs/2307.08691))
- ▶ **LoRA**: Hu et al., 2021
- ▶ **Mixture of Experts**: Fedus et al., 2021 (Switch Transformer), Lepikhin et al., 2020 (GShard)
- ▶ **LLaMA 3**: Meta AI 2024 release notes
- ▶ **DeepSeek + GRPO**: <https://github.com/deepseek-ai>
- ▶ **BLT**: Meta AI, 2024 ([arXiv:2412.09871](https://arxiv.org/abs/2412.09871))
- ▶ **LLDM**: Li et al., 2022 (Diffusion-LM), MaskGIT, UniDiffuser
- ▶ **GRPO**: Shao et al., 2024 (DeepSeekMath, [arXiv:2402.03300](https://arxiv.org/abs/2402.03300))

Resources:

- ▶ HuggingFace ModelHub & Papers
- ▶ Meta AI blog
- ▶ DeepSeek official GitHub
- ▶ Open LLM Leaderboard:
<https://huggingface.co/spaces/HuggingFaceH4/open-llm-leaderboard>