

Reinforcement Learning

Naeemullah Khan

naeemullah.khan@kaust.edu.sa




جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

May 12, 2026

Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning



Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, Sergey Levine


- ▶ Number of times the agent must interact with the environment in order to learn a task
- ▶ Good sample complexity is the first prerequisite for successful skill acquisition.
- ▶ Learning skills in the real world can take a substantial amount of time
 - can get damaged through trial and error

- ▶ “Learning Hand-Eye Coordination for Robotic Grasping with Deep Learning and Large-Scale Data Collection”, Levine et al., 2016
 - 14 robot arms learning to grasp in parallel
 - objects started being picked up at around 20,000 grasps



<https://spectrum.ieee.org/automaton/robotics/artificial-intelligence/google-large-scale-robotic-grasping-project>

Observing the behavior of the robot after over 800,000 grasp attempts, which is equivalent to about 3000 robot-hours of practice, we can see the beginnings of intelligent reactive behaviors.



https://www.youtube.com/watch?v=cXaic_k80uM

- ▶ Solution?
- ▶ Off-Policy Learning!

- ▶ **On-policy learning:** use the deterministic outcomes or samples from the target policy to train the algorithm
 - has low sample efficiency (TRPO, PPO, A3C)
 - require new samples to be collected for nearly every update to the policy
 - becomes extremely expensive when the task is complex
- ▶ **Off-policy methods:** training on a distribution of transitions or episodes produced by a different behavior policy rather than that produced by the target policy
 - Does not require full trajectories and can reuse any past episodes (experience replay) for much better sample efficiency
 - relatively straightforward for Q-learning based methods

- **Value Function:** How good is a state?

$$\begin{aligned} V(s) &= \mathbb{E}[G_t \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) \mid S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\ &= \mathbb{E}[\underbrace{R_{t+1} + \gamma V(S_{t+1})}_{\text{temporal difference target}} \mid S_t = s] \end{aligned}$$

- Similarly, for **Q-Function:** How good is a state-action pair?

$$\begin{aligned} Q(s, a) &= \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \mathbb{E}[R_{t+1} + \gamma \mathbb{E}_{a \sim \pi} Q(S_{t+1}, a) \mid S_t = s, A_t = a] \end{aligned}$$

- ▶ $\dots, S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}, \dots$ (**on-policy**):

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

- ▶ **Q-Learning (off-policy)**

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(S_{t+1}, a) - Q(S_t, A_t))$$

- ▶ **DQN**, Minh et al., 2015

$$\mathcal{L}(\theta) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2 \right]$$

- ▶ Function Approximation
- ▶ Experience Replay: samples randomly drawn from replay memory
- ▶ Doesn't scale to continuous action space

- **Critic:** updates value function parameters w and depending on the algorithm it could be action-value $Q(a | s; w)$ or state-value $V(s; w)$.
- **Actor:** updates policy parameters θ , in the direction suggested by the critic, $\pi(a | s; \theta)$.

$$\nabla \mathcal{J}(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla \ln \pi(a | s, \theta) Q_{\pi}(s, a)]$$

policy gradient

$$\theta \leftarrow \theta + \alpha_{\theta} Q(s, a; w) \nabla_{\theta} \ln \pi(a | s; \theta)$$

update actor


$$G_{t:t+1} = r_t + \gamma Q(s', a'; w) - Q(s, a; w)$$

correction for action-value

$$w \leftarrow w + \alpha_w G_{t:t+1} \nabla_w Q(s, a; w)$$



update critic

- ▶ DDPG = DQN + DPG (Lillicrap et al., 2015)
 - off-policy actor-critic method that learns a deterministic policy in continuous domain
 - exploration noise added to the deterministic policy when select action
 - difficult to stabilize and brittle to hyperparameters (Duan et al., 2016, Henderson et al., 2017)
 - unscalable to complex tasks with high dimensions (Gu et al., 2017)





<https://www.youtube.com/watch?v=zR11FLZ-09M>

- ▶ Training is sensitive to randomness in the environment, initialization of the policy and the algorithm implementation




<https://gym.openai.com/envs/Walker2d-v2/>


- ▶ Knowing only one way to act makes agents vulnerable to environmental changes that are common in the real-world



<https://bair.berkeley.edu/blog/2017/10/06/soft-q-learning/>



Traditional Graph of MDP



Graphical Model with Optimality Variables

Normal trajectory distribution

$$p(\tau) = p(\mathbf{s}_1, \mathbf{a}_1, \dots, \mathbf{s}_T, \mathbf{a}_T \mid \theta) = p(\mathbf{s}_1) \prod_{t=1}^T p(\mathbf{a}_t \mid \mathbf{s}_t, \theta) p(\mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t).$$

Posterior trajectory distribution

$$p(\mathcal{O}_t = 1 \mid \mathbf{s}_t, \mathbf{a}_t) = \exp(r(\mathbf{s}_t, \mathbf{a}_t)).$$

$$\begin{aligned} p(\tau \mid \mathbf{o}_{1:T}) &\propto p(\tau, \mathbf{o}_{1:T}) = p(\mathbf{s}_1) \prod_{t=1}^T p(\mathcal{O}_t = 1 \mid \mathbf{s}_t, \mathbf{a}_t) p(\mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t) \\ &= p(\mathbf{s}_1) \prod_{t=1}^T \exp(r(\mathbf{s}_t, \mathbf{a}_t)) p(\mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t) \\ &= \left[p(\mathbf{s}_1) \prod_{t=1}^T p(\mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t) \right] \exp\left(\sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t)\right). \end{aligned}$$

Variational Inference

$$\begin{aligned} D_{\text{KL}}(\hat{p}(\tau) \| p(\tau)) &= -\mathbb{E}_{\tau \sim \hat{p}(\tau)} [\log p(\tau) - \log \hat{p}(\tau)]. \\ -D_{\text{KL}}(\hat{p}(\tau) \| p(\tau)) &= \mathbb{E}_{\tau \sim \hat{p}(\tau)} \left[\log p(\mathbf{s}_1) + \sum_{t=1}^T \left(\log p(\mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t) + r(\mathbf{s}_t, \mathbf{a}_t) \right) \right. \\ &\quad \left. - \log p(\mathbf{s}_1) - \sum_{t=1}^T \left(\log p(\mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t) + \log \pi(\mathbf{a}_t \mid \mathbf{s}_t) \right) \right] \\ &= \mathbb{E}_{\tau \sim \hat{p}(\tau)} \left[\sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t) - \log \pi(\mathbf{a}_t \mid \mathbf{s}_t) \right] \\ &= \sum_{t=1}^T \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \hat{p}(\mathbf{s}_t, \mathbf{a}_t)} [r(\mathbf{s}_t, \mathbf{a}_t) - \log \pi(\mathbf{a}_t \mid \mathbf{s}_t)] \\ &= \sum_{t=1}^T \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \hat{p}(\mathbf{s}_t, \mathbf{a}_t)} [r(\mathbf{s}_t, \mathbf{a}_t)] + \mathbb{E}_{\mathbf{s}_t \sim \hat{p}(\mathbf{s}_t)} [\mathcal{H}(\pi(\mathbf{a}_t \mid \mathbf{s}_t))]. \end{aligned}$$

Conventional RL Objective — Expected Reward

$$\sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [r(\mathbf{s}_t, \mathbf{a}_t)]$$





Maximum Entropy RL Objective — Expected Reward + Entropy of Policy

$$\sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [r(\mathbf{s}_t, \mathbf{a}_t) + \alpha \mathcal{H}(\pi(\cdot | \mathbf{s}_t))]$$


Entropy of a RV x

$$H(P) = \mathbb{E}_{x \sim P} [-\log P(x)]$$

- MaxEnt RL agent can capture different modes of optimality to improve robustness against environmental changes



<https://bair.berkeley.edu/blog/2017/10/06/soft-q-learning/>



$$\begin{aligned}
 & \min_{\pi} D_{\text{KL}}(\pi(\cdot | s_0) \| \exp(Q(s_0, \cdot))) \\
 &= \min_{\pi} \mathbb{E}_{\pi} \left[\log \frac{\pi(a_0 | s_0)}{\exp(Q(s_0, a_0))} \right] \\
 &= \max_{\pi} \mathbb{E}_{\pi} [Q(s_0, a_0) - \log \pi(a_0 | s_0)] \\
 &= \max_{\pi} \mathbb{E}_{\pi} \left[\sum_t r(s_t, a_t) + \mathcal{H}(\pi(\cdot | s_0)) \mid s_0 \right] \\
 &= J_{\text{MaxEnt}}(\pi(\cdot | s_0))
 \end{aligned}$$

► Soft Q-Learning (Haarnoja et al., 2017)

- off-policy algorithms under MaxEnt RL objective
- Learns Q^* directly
- Instead of taking max action, sample the action
- use approximate inference methods to sample
 - Stein variational gradient descent
- not true actor-critic

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(S_{t+1}, a) - Q(S_t, A_t))$$

$$Q_{\text{soft}}(\mathbf{s}_t, \mathbf{a}_t) \leftarrow r_t + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p_s} [V_{\text{soft}}(\mathbf{s}_{t+1})], \quad \forall \mathbf{s}_t, \mathbf{a}_t$$

$$V_{\text{soft}}(\mathbf{s}_t) \leftarrow \alpha \log \int_{\mathcal{A}} \exp\left(\frac{1}{\alpha} Q_{\text{soft}}(\mathbf{s}_t, \mathbf{a}')\right) d\mathbf{a}', \quad \forall \mathbf{s}_t$$

- ▶ One of the most efficient model-free algorithms
 - SOTA off-policy
 - well suited for real world robotics learning
- ▶ Can learn stochastic policy on continuous action domain
- ▶ Robust to noise
- ▶ **Ingredients:**
 - Actor-critic architecture with separate policy and value function networks
 - Off-policy formulation to reuse of previously collected data for efficiency
 - Entropy-constrained objective to encourage stability and exploration

- **policy evaluation:** compute value of π according to Max Entropy RL Objective
- modified Bellman backup operator $\mathcal{T}$:

$$\mathcal{T}^\pi Q(\mathbf{s}_t, \mathbf{a}_t) \triangleq r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p} \left[\boxed{V(\mathbf{s}_{t+1})} \right]$$
$$\boxed{V(\mathbf{s}_t)} = \mathbb{E}_{\mathbf{a}_t \sim \pi} \left[Q(\mathbf{s}_t, \mathbf{a}_t) - \boxed{\alpha \log \pi(\mathbf{a}_t \mid \mathbf{s}_t)} \right]$$

- **Lemma 1:** Contraction Mapping for Soft Bellman Updates

$$Q^{k+1} = \mathcal{T}^\pi Q^k \quad \text{converges to the soft Q-function of } \pi$$

- ▶ **policy improvement:** update policy towards the exponential of the new soft Q-function
- ▶ modified Bellman backup operator $\mathcal{T}$:
 - choose tractable family of distributions big Π
 - choose KL divergence to project the improved policy into big Π

$$\pi_{\text{new}} = \arg \min_{\pi' \in \Pi} D_{\text{KL}} \left(\pi'(\cdot | \mathbf{s}_t) \parallel \frac{\exp(\frac{1}{\alpha} Q^{\pi_{\text{old}}}(\mathbf{s}_t, \cdot))}{Z^{\pi_{\text{old}}}(\mathbf{s}_t)} \right)$$

▶ **Lemma 2**

$$Q^{\pi_{\text{new}}}(\mathbf{s}_t, \mathbf{a}_t) \geq Q^{\pi_{\text{old}}}(\mathbf{s}_t, \mathbf{a}_t) \quad \text{for any state action pair}$$

- ▶ soft policy iteration: soft policy evaluation $\longleftrightarrow$ soft policy improvement
- ▶ **Theorem 1:** Repeated application of soft policy evaluation and soft policy improvement from any policy $\pi \in \Pi$ converges to the optimal MaxEnt policy among all policies in Π
 - exact form applicable only in discrete case
 - need function approximation to represent Q-values in continuous domains
 - $\rightarrow$ Soft Actor-Critic (SAC)!

$$Q_{\theta}(\mathbf{s}_t, \mathbf{a}_t)$$

- e.g. neural network

parameterized tractable policy

$$\pi_{\phi}(\mathbf{a}_t \mid \mathbf{s}_t)$$

- e.g. Gaussian with mean and covariances given by neural networks

soft Q-function objective and its stochastic gradient wrt its parameters

$$J_Q(\theta) \quad \hat{\nabla}_{\theta} J_Q(\theta)$$

policy objective and stochastic gradient wrt its parameters

$$J_{\pi}(\phi) \quad \hat{\nabla}_{\phi} J_{\pi}(\phi)$$

► Critic — Soft Q-function

- minimize square error
- $\bar{\theta}$ exponential moving average of soft Q-function weights to stabilize training (DQN)

$$J_Q(\theta) = \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \mathcal{D}} \left[\frac{1}{2} \left(Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - (r(\mathbf{s}_t, \mathbf{a}_t) + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p} [V_{\bar{\theta}}(\mathbf{s}_{t+1})]) \right)^2 \right]$$

$$V(\mathbf{s}_t) = \mathbb{E}_{\mathbf{a}_t \sim \pi} [Q(\mathbf{s}_t, \mathbf{a}_t) - \alpha \log \pi(\mathbf{a}_t | \mathbf{s}_t)]$$

$$\hat{\nabla}_\theta J_Q(\theta) = \nabla_\theta Q_\theta(\mathbf{a}_t, \mathbf{s}_t) (Q_\theta(\mathbf{s}_t, \mathbf{a}_t) - (r(\mathbf{s}_t, \mathbf{a}_t) + \gamma (Q_{\bar{\theta}}(\mathbf{s}_{t+1}, \mathbf{a}_{t+1}) - \alpha \log(\pi_\phi(\mathbf{a}_{t+1} | \mathbf{s}_{t+1}))))))$$

► **Actor — Policy**

$$\pi_{\text{new}} = \arg \min_{\pi' \in \Pi} D_{\text{KL}} \left(\pi'(\cdot | \mathbf{s}_t) \parallel \frac{\exp(\frac{1}{\alpha} Q^{\pi_{\text{old}}}(\mathbf{s}_t, \cdot))}{Z^{\pi_{\text{old}}}(\mathbf{s}_t)} \right)$$

► multiply by alpha and ignoring the normalization Z

$$J_{\pi}(\phi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}} [\mathbb{E}_{\mathbf{a}_t \sim \pi_{\phi}} [\alpha \log(\pi_{\phi}(\mathbf{a}_t | \mathbf{s}_t)) - Q_{\theta}(\mathbf{s}_t, \mathbf{a}_t)]]$$

► reparameterize with neural network f $\mathbf{a}_t = f_{\phi}(\epsilon_t; \mathbf{s}_t)$

- epsilon: input noise vector, sampled from a fixed distribution (spherical Gaussian)

$$J_{\pi}(\phi) = \mathbb{E}_{\mathbf{s}_t \sim \mathcal{D}, \epsilon_t \sim \mathcal{N}} [\alpha \log \pi_{\phi}(f_{\phi}(\epsilon_t; \mathbf{s}_t) | \mathbf{s}_t) - Q_{\theta}(\mathbf{s}_t, f_{\phi}(\epsilon_t; \mathbf{s}_t))]$$

$$\hat{\nabla}_{\phi} J_{\pi}(\phi) = \nabla_{\phi} \alpha \log(\pi_{\phi}(\mathbf{a}_t | \mathbf{s}_t)) + (\nabla_{\mathbf{a}_t} \alpha \log(\pi_{\phi}(\mathbf{a}_t | \mathbf{s}_t)) - \nabla_{\mathbf{a}_t} Q(\mathbf{s}_t, \mathbf{a}_t)) \nabla_{\phi} f_{\phi}(\epsilon_t; \mathbf{s}_t)$$

► **Unbiased gradient estimator that extends DDPG style policy gradients to any tractable stochastic policy**

Algorithm 1 Soft Actor-Critic

```

1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi_1, \phi_2$ , empty replay buffer  $\mathcal{D}$ 
2: Set target parameters equal to main parameters  $\phi_{\text{target},1} \leftarrow \phi_1, \phi_{\text{target},2} \leftarrow \phi_2$ 
3: repeat
4:   Observe state  $s$  and select action  $a \sim \pi_\theta(\cdot|s)$ 
5:   Execute  $a$  in the environment
6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$ 
8:   If  $s'$  is terminal, reset environment state.
9:   if it's time to update then
10:    for  $j$  in range(however many updates) do
11:      Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$ 
12:      Compute targets for the Q functions:

$$y(r, s', d) = r + \gamma(1 - d) \left( \min_{i=1,2} Q_{\phi_{\text{target},i}}(s', \tilde{a}') - \alpha \log \pi_\theta(\tilde{a}'|s') \right), \quad \tilde{a}' \sim \pi_\theta(\cdot|s')$$

13:      Update Q-functions by one step of gradient descent using

$$\nabla_{\phi_i} \frac{1}{|B|} \sum_{(s,a,r,s',d) \in B} (Q_{\phi_i}(s, a) - y(r, s', d))^2 \quad \text{for } i = 1, 2$$

14:      Update policy by one step of gradient ascent using

$$\nabla_\theta \frac{1}{|B|} \sum_{s \in B} \left( \min_{i=1,2} Q_{\phi_i}(s, \tilde{a}_\theta(s)) - \alpha \log \pi_\theta(\tilde{a}_\theta(s)|s) \right),$$

      where  $\tilde{a}_\theta(s)$  is a sample from  $\pi_\theta(\cdot|s)$  which is differentiable wrt  $\theta$  via the reparametrization trick.
15:      Update target networks with

$$\phi_{\text{target},i} \leftarrow \rho \phi_{\text{target},i} + (1 - \rho) \phi_i \quad \text{for } i = 1, 2$$

16:    end for
17:  end if
18: until convergence

```


<https://spinningup.openai.com/en/latest/algorithms/sac.html>

► Tasks


- A range of continuous control tasks from the OpenAI gym benchmark suite
- RL-Lab implementation of the Humanoid task
- The easier tasks can be solved by a wide range of different algorithms, the more complex benchmarks, such as the 21-dimensional Humanoid (rllab) are exceptionally difficult to solve with off-policy algorithms.

► Baselines:


- DDPG, SQL, PPO, TD3 (concurrent)
- TD3 is an extension to DDPG that first applied the double Q-learning trick to continuous control along with other improvements.




(a) Hopper-v1




(b) Walker2d-v1




(c) HalfCheetah-v1



(d) Ant-v1




(e) Humanoid-v1




(f) Humanoid (rllab)

- ▶ How does the stochasticity of the policy and entropy maximization affect the performance?
- ▶ Comparison with a deterministic variant of SAC that does not maximize the entropy and that closely resembles DDPG




<https://arxiv.org/abs/1801.01290>


Experimental Results: Hyperparameter Sensitivity



(a) Evaluation



(b) Reward Scale



(c) Target Smoothing Coefficient (τ)

<https://arxiv.org/abs/1801.01290>

- ▶ Unfortunately, SAC also suffers from brittleness to the alpha temperature hyperparameter that controls exploration
 - → automatic temperature tuning!

- Adaptive temperature coefficient


$$\sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} \left[r(\mathbf{s}_t, \mathbf{a}_t) + \underbrace{\alpha}_{\text{temperature}} \mathcal{H}(\pi(\cdot | \mathbf{s}_t)) \right]$$

- Extend to real-world tasks such as locomotion for a quadrupedal robot and robotic manipulation with a dexterous hand



<https://arxiv.org/abs/1801.01290>

- ▶ Dexterous Hand Manipulations
- ▶ 20 hour end-to-end learning
- ▶ valve position as input: SAC 3 hours vs. PPO 7.4 hours



<https://sites.google.com/view/sac-and-applications>

- ▶ Choosing the optimal temperature is non-trivial (tuned for each task)
- ▶ Constrained optimization problem:

$$\max \sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [r(\mathbf{s}_t, \mathbf{a}_t) + \alpha \mathcal{H}(\pi(\cdot | \mathbf{s}_t))]$$
$$\max_{\pi_{0:T}} \mathbb{E}_{\rho_\pi} \left[\sum_{t=0}^T r(\mathbf{s}_t, \mathbf{a}_t) \right] \quad \text{s.t.} \quad \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_\pi} [-\log(\pi_t(\mathbf{a}_t | \mathbf{s}_t))] \geq \mathcal{H} \quad \forall t$$

Unroll the expectation

$$\max_{\pi_0} \left(\mathbb{E}[r(\mathbf{s}_0, \mathbf{a}_0)] + \max_{\pi_1} \left(\mathbb{E}[\dots] + \boxed{\max_{\pi_T} \mathbb{E}[r(\mathbf{s}_T, \mathbf{a}_T)]} \right) \right)$$

For the last time step in the trajectory

$$\begin{aligned} \max_{\pi_T} \mathbb{E}_{(\mathbf{s}_T, \mathbf{a}_T) \sim \rho_\pi} [r(\mathbf{s}_T, \mathbf{a}_T)] &= \min_{\alpha_T \geq 0} \max_{\pi_T} \mathbb{E}[r(\mathbf{s}_T, \mathbf{a}_T) - \alpha_T \log \pi(\mathbf{a}_T | \mathbf{s}_T)] - \alpha_T \\ \arg \min_{\alpha_T} \mathbb{E}_{\mathbf{s}_T, \mathbf{a}_T \sim \pi_T^*} [-\alpha_T \log \pi_T^*(\mathbf{a}_T | \mathbf{s}_T; \alpha_T) - \alpha_T \mathcal{H}]. \end{aligned}$$

Similarly, for the previous time step

$$\begin{aligned}
 & \max_{\pi_{T-1}} \left(\mathbb{E}[r(\mathbf{s}_{T-1}, \mathbf{a}_{T-1})] + \max_{\pi_T} \mathbb{E}[r(\mathbf{s}_T, \mathbf{a}_T)] \right) \\
 &= \max_{\pi_{T-1}} \left(Q_{T-1}^*(\mathbf{s}_{T-1}, \mathbf{a}_{T-1}) - \alpha_T \mathcal{H} \right) \\
 &= \min_{\alpha_{T-1} \geq 0} \max_{\pi_{T-1}} \left(\mathbb{E}[Q_{T-1}^*(\mathbf{s}_{T-1}, \mathbf{a}_{T-1})] - \mathbb{E}[\alpha_{T-1} \log \pi(\mathbf{a}_{T-1} | \mathbf{s}_{T-1})] - \alpha_{T-1} \mathcal{H} \right) \\
 & \alpha_t^* = \arg \min_{\alpha_t} \mathbb{E}_{\mathbf{a}_t \sim \pi_t^*} \left[-\alpha_t \log \pi_t^*(\mathbf{a}_t | \mathbf{s}_t; \alpha_t) - \alpha_t \bar{\mathcal{H}} \right]
 \end{aligned} \tag{16}$$







Figure 1: Training curves on continuous control benchmarks. Soft actor-critic (blue and yellow) performs consistently across all tasks and outperforming both on-policy and off-policy methods in the most challenging tasks.

Figure 1: Training curves on continuous control benchmarks. Soft actor-critic (blue and yellow) performs consistently across all tasks and outperforming both on-policy and off-policy methods in the most challenging tasks.




- ▶ Lack of experiments on hard-exploration problem

- ▶ Lack of experiments on hard-exploration problem
- ▶ Approximating a multi-modal Boltzmann distribution with a unimodal Gaussian



- ▶ Lack of experiments on hard-exploration problem
- ▶ Approximating a multi-modal Boltzmann distribution with a unimodal Gaussian
- ▶ High-variance using automatic temperature tuning

- ▶ Lack of experiments on hard-exploration problem
- ▶ Approximating a multi-modal Boltzmann distribution with a unimodal Gaussian
- ▶ High-variance using automatic temperature tuning



- ▶ An off-policy maximum entropy deep reinforcement learning algorithm
 - Sample-efficient
 - Scale to high-dimensional observation/action space
 - Robustness to random seed, noise and etc.
- ▶ Theoretical Results
 - Convergence of soft policy iteration
 - Derivation of soft-actor critic algorithm
- ▶ Empirical Results
 - SAC outperforms SOTA model-free deep RL methods, including DDPG, PPO and Soft Q-learning, in terms of the policy's optimality, sample complexity and robustness.

These slides have been adapted from

- ▶ Animesh Garg, [CSC2621: Reinforcement Learning in Robotics](#), University of Toronto