

Reinforcement Learning

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية

King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

June 14, 2026


Article | Published: 09 June 2021

A graph placement methodology for fast chip design


Azalia Mirhoseini, Anna Goldie, Mustafa Yazgan, Joe Wenjie Jiang, Ebrahim Songhori, Shen Wang, Young-Joon Lee, Eric Johnson, Omkar Pathak, Azade Nazi, Jiwoo Pak, Andy Tong, Kavya Srinivasa, William Hang, Emre Tuncer, Quoc V. Le, James Laudon, Richard Ho, Roger Carpenter & Jeff Dean

Nature **594**, 207–212 (2021)

In the past decade, systems and hardware have transformed ML.




In the past decade, systems and hardware have transformed ML. Now, it's time for ML to transform systems and hardware.



Demand for Compute Outpacing Supply (Moore's Law)

Benchmark	Error rate	Polynomial		
		Computation Required (Gflops)	Environmental Cost (CO_2)	Economic Cost (\$)
ImageNet	Today: 11.5%	10^{14}	10^6	10^6
	Target 1: 5%	10^{19}	10^{10}	10^{11}
	Target 2: 1%	10^{28}	10^{20}	10^{20}

Implications of achieving performance on the computation, carbon emissions, and economic costs from deep learning on projections from polynomial models. *The Computational Limits of Deep Learning*, Thompson et al., 2020



Since 2012, the amount of compute used in the largest AI training runs doubled every
3.4 months, OpenAI, 2019





Figure 1 Language modeling performance improves smoothly as we increase the model size, dataset size, and amount of compute² used for training. For optimal performance all three factors must be scaled up in tandem. Empirical performance has a power-law relationship with each individual factor when not bottlenecked by the other two.

0

⁰Kaplan, McCandlish, Henighan, Brown, Chess, Child, Gray, Radford, Wu, Amodei. [Scaling Laws for Neural Language Models](#). OpenAI 2020.


- ▶ Enabling cheaper, faster, and more environmentally friendly chips
- ▶ Potential to reduce the design cycle from 1.5–2 years to weeks
 - Today, we design chips for the NN architectures of 2–5 years from now
 - Shortening the chip design cycle would enable us to be far more adaptive to the rapidly advancing field of machine learning
- ▶ New possibilities emerge if we evolve NN architectures and chips together
 - Discovering the next generation of NN architectures (which would not be computationally feasible with today's chips)

- ▶ A form of graph resource optimization
- ▶ Place the chip components to minimize the latency of computation, power consumption, chip area and cost, while adhering to constraints, such as congestion, cell utilization, heat profile, etc.




Complexity of Chip Placement Problem

Chess




Number of states $\sim 10^{123}$

Go



Number of states $\sim 10^{360}$

Chip Placement



Number of states $\sim 10^{9000}$

Partitioning-Based Methods
(e.g. MinCut)

Stochastic/Hill-
Climbing Methods
(e.g. Simulated Annealing)

Analytic Solvers
(e.g. RePIAce)

Partitioning-Based Methods
(e.g. MinCut)

Stochastic/Hill-
Climbing Methods
(e.g. Simulated Annealing)


Analytic Solvers
(e.g. RePIAce)

Learning-Based Methods

State: Graph embedding of chip netlist, embedding of the current node, and the canvas.

Action: Placing the current node onto a grid cell.

Reward: A weighted average of total wirelength, density, and congestion.




$$J(\theta, G) = \frac{1}{K} \sum_{g \sim G} \mathbb{E}_{p \sim \pi_{\theta}} [R_{p,g}]$$

- ▶ G : set of training graphs, K : size of the training set
- ▶ π_{θ} : RL policy parameterized by θ
- ▶ $R_{p,g}$: reward corresponding to placement p of netlist (graph) g

$$R_{p,g} = -\text{Wirelength}(p, g) - \lambda \text{Congestion}(p, g) - \gamma \text{Density}(p, g)$$

A Hybrid Approach to Placement Optimization




White area are macros and the green area is composed of standard cell clusters. The agent finds smoother, rounder macro placements to reduce the wirelength.

Human Expert



Time taken: ~6–8 weeks
Total wirelength: 57.07m
Route DRC* violations: 1766

ML Placer




Time taken: 24 hours
Total wirelength: 55.42m (−2.9% shorter)
Route DRC violations: 1789 (+23 – negligible difference)


*DRC: Design Rule Checking

Moving Towards Generalized Placements

Before: Training from scratch for each chip netlist



Then: Pre-training the policy and fine-tuning on new netlists




Key observation: A value network trained only on placements generated by a single policy is unable to accurately predict the quality of placements generated by another policy, limiting the ability of the policy network to generalize.

To decompose the problem, they trained models capable of accurately predicting reward from off-policy data. (*Inverse RL*)



What to do in domains where reward is hard to specify?

- ▶ One solution is to ask humans to provide feedback — however, this is prohibitively expensive in the naive formulation, as RL typically requires thousands to millions of labels to learn an effective policy (depending on the complexity of the task)
- ▶ But what if you train a model to predict human judgments and then use this predictive model as the reward signal?




Reinforcement Learning for **Human Feedback** (RLHF)

- ▶ Without access to the true reward function and labeling $< 1\%$ of the environment interactions, able to perform complex tasks, including Atari games and MuJoCo.



0

- ▶ “Secret sauce” behind powerful LLMs like ChatGPT!
- ▶ Humans rank-order pairs of behavior, train a preference model, use preference model as reward, and RL-finetune to optimize “good” behavior
- ▶ Performing RLHF on top of pretrained large language models (LLMs) greatly improves instruction-following / in-context learning / prompting.



أكاديمية كاوست
KAUST ACADEMY



Step 1: Collect Human Judgments

Playground task 1

Talk to the assistant

Task 1 of 5

Decide on a task you'd like the assistant to help with and enter it into the task box. Interact with the AI assistant. When you're finished, select a final response from the assistant and leave a comment on how the assistant did in the comment box. Click to go to the next task.

Comments

I thought the assistant was ...

Rating

☐ Bad ☐ ☐ ☐ ☐ Good

Next Task

Human

I have to come up with a conversation with an AI assistant for the purposes of showing it off in a demo. What sorts of conversation topics would you recommend?

Assistant

I think the best topics for conversation are common themes and interests that everyone could have. Things like weather, sports, hobbies, current events, popular culture, and general life topics like relationships and work would all be good conversation topics. You could also have more advanced conversations discussing philosophy and morality if you want to demonstrate some higher level AI capabilities.

Human

I think I'm more interested in an advanced conversation. What's a good starter question if I wanted to discuss something philosophical, like epistemology or the nature of language?

Assistant

I think an excellent philosophical topic would be epistemology, which is the theory of knowledge, and specifically how we can know and understand the world around us. A really good starting question would be something like "How can we truly know and understand the world?", or "What is the nature of knowledge?". If you want to go even deeper, I think philosophy of language would also be fascinating - questions like "What is meaning?", "How do language and thought relate?", and "How can language mislead and deceive?" would be great discussion topics.

Human

How would you answer a question like: How do language and thought relate?

Choose the most helpful and honest response

A

I think a good answer to that would be that language helps us to structure our thoughts and express what we are thinking. Language gives us a symbolic system of signs and meanings that we use to communicate our ideas and make sense of the world. It's critical for how we record and relate knowledge, and I think it's impossible to completely understand and express our thoughts without the use of language. Without language, our thoughts would just remain as instinctual impulses and basic desires, and we would have a much harder time comprehending and reflecting on the world around us.

B



I think the key relationship between language and thought is that language allows us to communicate and express our thoughts and ideas to others. Language is a system of symbols and signs that we use to organize and express our thoughts and mental concepts. So language is how we translate our inner thoughts and ideas into something shared and communicable with others. Without language, we wouldn't be able to express or understand the thoughts and ideas of others, and so language is essentially bridge between our inner minds and the shared external world.

A A A A B B B

A is better B is better

Step 2: Train Preference Models (PMs)

- ▶ Train PM to assign a higher score to the response preferred by a human rater
- ▶ Base models from 13M through 52B parameters (in increments of 4x)



Step 3: Perform RL-Finetuning with PM as Reward Signal


- ▶ Extract all prompts from the previous steps, prompt the base LM to respond, and then use the PM score as the reward signal
- ▶ Train with Proximal Policy Optimization (PPO) with an auxiliary KL penalty

$$r_{\text{total}} = r_{\text{PM}} - \lambda_{\text{KL}} D_{\text{KL}}(\text{policy} \parallel \text{policy}_0)$$


- ▶ Alignment tax for small models but alignment bonus for 13B+ models
- ▶ Tradeoff between helpfulness and harmlessness, but performance improves on both distributions as model scale up
- ▶ RLHF improves programming ability for models pretrained on code
- ▶ RLHF boosts performance on MMLU, Lambada, Hellaswag, OpenBookQA, and ARC, but hurt performance on TriviaQA compared to a base models

Next Step: RL from AI Feedback (RLAIF)!


- **Motivation:** Scaling supervision — as models approach or exceed human-level performance, it becomes difficult for humans to supervise them.
- **RLAIF:** Perform RL-finetuning using AI feedback derived from a “constitution” describing desired behavior. Humans don’t need to be in the loop, except to write the constitution!




- **Supervised Learning:** Improves initial model, which helps with exploration and sample efficiency
- **Reinforcement Learning:** Significantly boosts performance and reliability of the final policy




1. Sample from an initial policy
2. Generate “self-critiques” and revisions
3. Finetune the original model with the revised responses



1. Sample from a finetuned model
2. Use a model to evaluate which of two responses is “better”
3. Train a preference model on the AI-labeled data
4. Perform RL-finetuning with the PM as the reward signal (just like RLAIIF)




- Finetuning with AI-generated feedback can generate results that match or exceed models that are finetuned with human feedback




RL for **Device Placement**

What is device placement and why is it important?


Trend towards many-device training, bigger models, larger batch sizes



Google neural machine translation '16
300 million parameters, trained on 128
GPUs



Sparsely gated mixture of experts '17
130 billion parameters, trained on 128
GPUs




BigGAN '18
355 million parameters,
trained on 512 TPU cores

- ▶ Often based on greedy heuristics
- ▶ Requires deep understanding of devices: nonlinear FLOPs, bandwidth, latency behavior
- ▶ Requires modeling parallelism and pipelining
- ▶ Does not generalize well

- ▶ ML is repeatedly replacing rule based heuristics
- ▶ RL can be applied to device placement
 - Effective search across large state and action spaces to find optimal solutions
 - Automated learning from underlying environment only based on reward function (e.g. runtime of a program)


Posing device placement as an RL problem




0


⁰Mirhoseini et al. **Hierarchical Planning for Device Placement**. ICLR 2018.

Posing device placement as an RL problem




An end-to-end hierarchical placement model





White represents CPU (Ixion Haswell 2300). Each other color represents a separate GPU (Nvidia Tesla K80). Searching over a space of 5^{280} possible assignments.⁰

⁰Mirhoseini et al. **Hierarchical Planning for Device Placement**. ICLR 2018.



Results (runtime in seconds)

Tasks	CPU Only	GPU Only	#GPUs	Human Expert	Scotch	MinCut	Hierarchical Planner	Runtime Reduction
Inception-V3	0.61	0.15	2	0.15	0.93	0.82	0.13	16.3%
ResNet	-	1.18	2	1.18	6.27	2.92	1.18	0%
RNNLM	6.89	1.57	2	1.57	5.62	5.21	1.57	0%
NMT (2-layer)	6.46	OOM	2	2.13	3.21	5.34	0.84	60.6%
NMT (4-layer)	10.68	OOM	4	3.64	11.18	11.63	1.69	53.7%
NMT (8-layer)	11.52	OOM	8	3.88	17.85	19.01	4.07	-4.9%

Resource Management with Deep Reinforcement Learning

Hongzi Mao*, Mohammad Alizadeh*, Ishai Menache†, Srikanth Kandula†
Massachusetts Institute of Technology*, Microsoft Research†
{hongzi, alizadeh}@mit.edu, {ishai, srikanth}@microsoft.com


0

⁰Mao et al. [Resource Management with Deep Reinforcement Learning](#).

- ▶ Resource Management is a difficult online decision making task
- ▶ Requires to understand the workload and environment
- ▶ RL solution for resource management!

```
for each iteration:
     $\Delta\theta \leftarrow 0$ 
    for each jobset:
        run episode  $i = 1, \dots, N$ :
             $\{s_1^i, a_1^i, r_1^i, \dots, s_{L_i}^i, a_{L_i}^i, r_{L_i}^i\} \sim \pi_\theta$ 
            compute returns:  $v_t^i = \sum_{s=t}^{L_i} \gamma^{s-t} r_s^i$ 
            for  $t = 1$  to  $L$ :
                compute baseline:  $b_t = \frac{1}{N} \sum_{i=1}^N v_t^i$ 
                for  $i = 1$  to  $N$ :
                     $\Delta\theta \leftarrow \Delta\theta + \alpha \nabla_\theta \log \pi_\theta(s_t^i, a_t^i)(v_t^i - b_t)$ 
                end
            end
        end
    end
     $\theta \leftarrow \theta + \Delta\theta$  % batch parameter update
end
```

Figure 3: Pseudo-code for training algorithm.






Figure 4: Job slowdown at different levels of load.

Figure 4: Job slowdown at different levels of load.




January 2024

SERL: A Software Suite for Sample-Efficient Robotic Reinforcement Learning

Jianlan Luo^{1*}, Zheyuan Hu^{1*}, Charles Xu¹, You Liang Tan¹, Jacob Berg², Archit Sharma³, Stefan Schaal⁴, Chelsea Finn³, Abhishek Gupta² and Sergey Levine¹


^{*}Equal Contribution, ¹Department of EECS, University of California, Berkeley, ²Department of Computer Science and Engineering, University of Washington, ³Department of Computer Science, Stanford University, ⁴Intrinsic Innovation LLC


0



⁰<https://github.com/rail-berkeley/serl>

- ▶ Provide a software framework for reinforcement learning with robotics
- ▶ Contains a sample efficient off-policy deep RL method, together with methods for computing rewards and resetting the environment, a high-quality controller for a widely-adopted robot, and a number of challenging example tasks.
- ▶ The accompanied algorithm, SERL, is derived from RLPD which itself is a variant of SAC





Real-Time Bidding with Multi-Agent Reinforcement Learning in Display Advertising

Junqi Jin[§], Chengru Song[§], Han Li[§], Kun Gai[§], Jun Wang[†], Weinan Zhang[‡]

[§]Alibaba Group, [†]University College London, [‡]Shanghai JiaoTong University

{junqi.jjq, chengru.scr, lihan.lh, jingshi.gk}@alibaba-inc.com, j.wang@cs.ucl.ac.uk, wnzhang@sjtu.edu.cn

0

⁰Jin et al. Real-Time Bidding with Multi-Agent Reinforcement Learning in Display Advertising




Figure 1: An Overview of Taobao Display Advertising System. Matching, RTP and Ranking modules sequentially process user requests, and finally return specified quantity of ads. These ads are shown in *Guess What You Like* of Taobao App, tagged by *Hot* (as shown in red dashed box) and surrounded with recommendation results.

Algorithm 1: DCMAB Algorithm

```

1 Initialize  $Q_i(s^q, a_1^q, \dots, a_N^q, d | \theta_i^Q)$ , actor  $\mu_i([g, x] | \theta_i^\mu)$ , target
   network  $Q'_i, \mu'_i$  with  $\theta_i^{Q'} \leftarrow \theta_i^Q, \theta_i^{\mu'} \leftarrow \theta_i^\mu$  for each agent  $i$ .
2 Initialize replay memory  $D$ 
3 for  $episode = 1$  to  $E$  do
4   Initialize a random process  $\mathcal{N}$  for action exploration
5   Receive initial state  $s$  for all agents
6   for  $t = 1$  to  $T$  do
7     For each agent  $i$ , compute  $a_i^q$  and add  $\mathcal{N}_i$ .
8     for auctions in parallel workers in  $T_p$  do
9       For each agent  $i$ , compute bratio and combined
         with  $a_i^q$  compute adjusting ratio  $\alpha$  and execute.
10      For each agent  $i$ , save reward, cost and maintain
         distribution  $d$ .
11    end
12    For each agent  $i$ , merge rewards, cost in last  $T_p$  to get
      reward  $r_i$  and update state to  $s^{q'}$ . Store
       $(s^q, d, a_i^q, r_i, s^{q'})$  to replay memory.
13     $s^{q'} \leftarrow s^q$ 
14    for agent  $i=1$  to  $N$  do
15      Sample a random minibatch of  $S$  samples
         $(s^q, d, a_1^q, \dots, a_N^q, r_i, s^{q'}, d')$  from  $D$ 
16      Update critic by minimizing loss with Eqs.(9),(10).
17      Update actor with Eq. (11).
18      Update target network:  $\theta' \leftarrow \tau\theta + (1 - \tau)\theta$ 
19    end
20  end
21 end

```

	AgentC1			AgentC2			AgentC3			Total		
Indices	ROI	CPA	COST	ROI	CPA	COST	ROI	CPA	COST	ROI	CPA	COST
Manual	100.00%	100.00%	99.65%	100.00%	100.00%	99.71%	100.00%	100.00%	99.42%	100.00%	100.00%	99.52%
Bandit	121.38%	82.43%	99.87%	159.41%	62.73%	99.53%	102.63%	97.39%	99.64%	112.14%	84.23%	99.63%
A2C	103.30%	96.57%	99.39%	106.85%	93.58%	99.60%	170.91%	58.55%	99.66%	158.38%	68.09%	99.62%
DDPG	577.87%	17.27%	99.51%	976.80%	10.23%	99.18%	164.29%	60.85%	99.76%	305.75%	24.26%	99.58%
DCMAB	690.18%	14.46%	99.38%	584.63%	17.10%	99.43%	275.11%	36.34%	99.57%	340.38%	24.84%	99.51%

These slides have been adapted from

- ▶ Ana Goldie, [CS224R: Deep Reinforcement Learning](#), Stanford