

Stable Diffusion

Naeemullah Khan

naeemullah.khan@kaust.edu.sa




جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology



LMH
Lady Margaret Hall
University of Oxford

June 12, 2026



1. Stable Diffusion overview: text-to-image foundations and key components
2. Motivation and introduction to diffusion models
3. Denoising Diffusion Probabilistic Models (DDPM): forward and reverse processes
4. Training, sampling, and network architectures (U-Net and attention)
5. Latent Diffusion Models and Stable Diffusion (VAE, CLIP, Stable Video Diffusion)
6. Challenges: slow sampling, generative learning trilemma, and research directions
7. Applications: GLIDE, DALL·E 2, Imagen, Nano Banana 2, and extensions
8. Summary, limitations, and references

What is Stable Diffusion?

- ▶ A state-of-the-art text-to-image generative model.
- ▶ Developed by Stability AI and collaborators (2022).
- ▶ Open-source, enabling wide adoption and customization.



Advantages

- ▶ High-quality, diverse image generation.
- ▶ Efficient training and inference due to latent space operations.
- ▶ Open-source and extensible for research and applications.

What do we need to make a good generative model?

1. Method of learning to generate new images
2. Way to link text and images
3. Way to compress images
4. Way to add in good inductive biases

Forward/reverse diffusion

Text-image representation model

Autoencoder

U-net architecture + 'attention'

Making a 'good' generative model is about making all these parts work together well!

The Journey of Generative Modeling:


- ▶ Early successes: GANs and VAEs enabled impressive image synthesis and data generation.
- ▶ **Drawbacks:**
 - GANs: Training instability, mode collapse, and sensitivity to hyperparameters.
 - VAEs: Blurry outputs and limited expressiveness.
- ▶ **Key Challenge:** Achieving both high sample quality and stable, reliable training.

Where do we go from here?

- ▶ Need for models that are robust, interpretable, and capable of generating diverse, high-fidelity samples.

Enter Diffusion Models:

- ▶ Inspired by thermodynamics, diffusion models gradually transform noise into data.
- ▶ **Highlights:**
 - Simple training objective
 - Stable optimization
 - State-of-the-art sample quality
 - Probabilistic and interpretable framework



Diffusion Models: **Introduction**

Diffusion models are a class of probabilistic generative models that learn to reverse a gradual noising process applied to data. They have gained popularity due to their ability to generate samples with high fidelity and diversity.

Diffusion models are a class of probabilistic generative models that learn to reverse a gradual noising process applied to data. They have gained popularity due to their ability to generate samples with high fidelity and diversity.

- **Inspiration:** Rooted in concepts from thermodynamics and stochastic processes.

Diffusion models are a class of probabilistic generative models that learn to reverse a gradual noising process applied to data. They have gained popularity due to their ability to generate samples with high fidelity and diversity.

- ▶ **Inspiration:** Rooted in concepts from thermodynamics and stochastic processes.
- ▶ **Key Idea:**
 - *Forward process:* Progressively corrupt data by adding Gaussian noise.
 - *Reverse process:* Learn to recover the original data by reversing the noising process.

Diffusion models are a class of probabilistic generative models that learn to reverse a gradual noising process applied to data. They have gained popularity due to their ability to generate samples with high fidelity and diversity.

- ▶ **Inspiration:** Rooted in concepts from thermodynamics and stochastic processes.
- ▶ **Key Idea:**
 - *Forward process:* Progressively corrupt data by adding Gaussian noise.
 - *Reverse process:* Learn to recover the original data by reversing the noising process.
- ▶ **Popularity:** Known for producing high-quality and diverse generated samples.

- ▶ **Forward:** Clean $\rightarrow$ Noise (additive Gaussian steps)
- ▶ **Reverse:** Learn **U-Net** to remove noise step by step
- ▶ **Training:** Model predicts the noise added
- ▶ **Sampling:** Roll back from pure noise to a data sample

Diffusion Models: **Denoising**


Denoising diffusion models, also known as score-based generative models, have recently emerged as a powerful class of generative models. They achieve remarkable results in high-fidelity image generation, often outperforming GANs.




Diffusion Models Beat GANs on Image Synthesis [Dhariwal & Nichol, OpenAI, 2021](#)

Denoising diffusion models consist of two processes:

- ▶ A fixed (predefined) **forward diffusion process** q , which gradually adds Gaussian noise to an image until only pure noise remains.
- ▶ A learned **reverse denoising diffusion process** p_θ , where a neural network is trained to gradually denoise an image starting from pure noise, eventually recovering the original image.



Denoising Diffusion Probabilistic Models (cont.)





Diffusion Models: **Forward Process**

- ▶ The forward process is a Markov chain that iteratively adds Gaussian noise to the image at each timestep.
- ▶ Given clean data $\mathbf{x}_0$, we generate $\mathbf{x}_t$ as:

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I})$$


- ▶ Over many steps (e.g., 1000), the data becomes indistinguishable from pure noise.

- ▶ This process is predefined and not learned.
- ▶ The full forward process:

$$q(\mathbf{x}_{1:T}|\mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t|\mathbf{x}_{t-1})$$

- ▶ Each new (slightly noisier) image at time step t is drawn from a conditional Gaussian distribution with $\mu_t = \sqrt{1 - \beta_t}\mathbf{x}_{t-1}$ and $\sigma_t^2 = \beta_t$, which we can do by sampling $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ and then setting $\mathbf{x}_t = \sqrt{1 - \beta_t}\mathbf{x}_{t-1} + \sqrt{\beta_t}\epsilon$.

Forward Diffusion Process (cont.)



- What if we want to go directly from q_0 to q_t ?
- Let $\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$, then

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I})$$

- This allows us, during training, to optimize random terms of the loss function L (i.e., randomly sample t and optimize L_t).

Diffusion Models: **Reverse Denoising Process**

In the reverse denoising process, we denoise Gaussian noise to generate an image.

In the reverse denoising process, we denoise Gaussian noise to generate an image.

- ▶ We start with a sample from the noise distribution $p(\mathbf{x}_T) = \mathcal{N}(\mathbf{x}_T; 0, 1)$.

In the reverse denoising process, we denoise Gaussian noise to generate an image.

- ▶ We start with a sample from the noise distribution $p(\mathbf{x}_T) = \mathcal{N}(\mathbf{x}_T; 0, 1)$.
- ▶ The goal is to iteratively denoise this sample to recover the original image $\mathbf{x}_0$.

In the reverse denoising process, we denoise Gaussian noise to generate an image.

- ▶ We start with a sample from the noise distribution $p(\mathbf{x}_T) = \mathcal{N}(\mathbf{x}_T; 0, 1)$.
- ▶ The goal is to iteratively denoise this sample to recover the original image $\mathbf{x}_0$.
- ▶ At each time step t , we want to sample $\mathbf{x}_{t-1}$ given $\mathbf{x}_t$.

In the reverse denoising process, we denoise Gaussian noise to generate an image.

- ▶ We start with a sample from the noise distribution $p(\mathbf{x}_T) = \mathcal{N}(\mathbf{x}_T; 0, 1)$.
- ▶ The goal is to iteratively denoise this sample to recover the original image $\mathbf{x}_0$.
- ▶ At each time step t , we want to sample $\mathbf{x}_{t-1}$ given $\mathbf{x}_t$.
- ▶ The reverse process is defined as $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$, which we will learn using a neural network.

- ▶ The reverse process is a Markov chain that iteratively denoises the image.
- ▶ We want to sample $\mathbf{x}_{t-1}$ from $p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)$.
- ▶ However, $p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)$ is unknown and needs to be learned.

- ▶ We approximate it with a neural network that predicts the mean and variance of the Gaussian distribution.
- ▶ If β_t is small enough at each time step, we can assume $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$ is a Gaussian distribution:

$$p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t, t), \Sigma_\theta(\mathbf{x}_t, t))$$

where μ_θ and Σ_θ are neural networks conditioned on $\mathbf{x}_t$ and t .


- For simplicity, we can fix the variance Σ_θ to a constant value (e.g., $\beta_t \mathbf{I}$) and only learn the mean μ_θ .

- The reverse process can then be expressed as:

$$p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_{\theta}(\mathbf{x}_t, t), \beta_t \mathbf{I})$$

and the joint distribution over all time steps is:

$$p_{\theta}(\mathbf{x}_{0:T}) = p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t)$$



Reverse denoising process in diffusion models.



Diffusion Models: **Training & Sampling Algorithms**

Algorithm 1 Training

- 1: **repeat**
 - 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
 - 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
 - 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
 - 5: Take gradient descent step on
$$\nabla_{\theta} \left\| \epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t) \right\|^2$$
 - 6: **until** converged
-

Training Algorithm

1. Sample x_0 from the data distribution.
2. Randomly select a timestep t .
3. Corrupt x_0 with noise to obtain x_t .
4. Model predicts noise $\epsilon_{\theta}(x_t, t)$.
5. Compute loss: mean squared error between predicted and true noise.
6. Optimize using SGD or Adam.

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```


Sampling Algorithm

1. Start from Gaussian noise $\mathbf{x}_T$.
2. For each timestep $t = T, \dots, 1$:
 - 2.1 Predict noise $\boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)$.
 - 2.2 Compute posterior mean $\mu_\theta(\mathbf{x}_t, t)$.
 - 2.3 Sample $\mathbf{x}_{t-1}$ from the Gaussian posterior.
 - 2.4 (Optional) Apply guidance (e.g., classifier-free) for improved results.




Diffusion Models: **Network Architectures**


- ▶ Diffusion models commonly use U-Net architectures with ResNet blocks and self-attention layers to represent $\epsilon_{\theta}(x_t, t)$.
- ▶ Time is encoded using sinusoidal positional embeddings or random Fourier features.




Network Architectures (cont.)



Network Architectures (cont.)




- ▶ Diffusion models can be viewed as a special form of hierarchical VAEs (one VAE after another).
- ▶ In diffusion models:
 - The encoder is fixed.
 - The latent variables have the same dimension as the data.
 - The denoising model is shared across different timesteps.
 - The model is trained with a reweighted variational bound.





Diffusion Models: **Latent Diffusion Models** (LDMs)

What is a Latent Space?



- ▶ A **latent space** is an abstract, compressed space where high-dimensional data (like images) is represented in a lower-dimensional form.
- ▶ Think of it as a hidden layer of meaning—each point in this space represents a potential image or concept.
- ▶ In classical models like GANs or VAEs, the latent space is directly sampled from (e.g., a vector $\mathbf{z}$), and the generator turns it into an image.

What is a Latent Space? (cont.)

- **Latent** space = “Hidden space of ideas or features”

Two main ways latent spaces are used in diffusion:

► 1. **Latent Diffusion Models (LDMs):**

- Diffusion is performed in a compressed latent space, not directly on high-dimensional images.
- **Pipeline:** Compress image $\rightarrow$ Denoise in latent space $\rightarrow$ Decode back to image.
- Enables faster, cheaper training and inference, and allows for semantic control (e.g., text prompts, masks).

► 2. **Latent Representations in Noise Vectors:**

- Even in standard diffusion models, the initial noise vector acts as a latent code.
- Changing the noise vector generates different images or morphs between them.

- ▶ **Traditional pixel-space diffusion models are computationally expensive:** they operate directly on high-dimensional images, making both training and inference costly.
- ▶ **Latent diffusion reduces computational cost** by performing the diffusion process in a compressed, lower-dimensional latent space—while still maintaining high-quality image generation.
- ▶ This efficiency enables more complex and practical applications, such as text-guided generation and flexible image editing.
- ▶ **How is this achieved?** Instead of working in pixel space, we learn a compressed latent space using a Variational Autoencoder (VAE) with very low KL divergence weight (e.g., 10^{-6}). This yields a compact representation suitable for diffusion.
- ▶ This approach mirrors ideas from autoregressive models (e.g., VQGAN), enabling even more efficient and controllable generation pipelines.


- ▶ Text-to-image generation (e.g., Stable Diffusion) uses text embeddings in latent space.
- ▶ Image editing by modifying specific parts of the latent.
- ▶ Inversion: map a real image to latent, edit, then decode back.

- ▶ **Latent Diffusion Models (LDMs):** Operate in a compressed latent space for efficiency.
- ▶ **Diffusion Process:** Gradually adds noise to data, then learns to reverse this process.
- ▶ **Text Conditioning:** Uses CLIP text encoders to guide image generation from prompts.

Architecture Overview

- ▶ **Encoder:** Maps images to latent space.
- ▶ **UNet:** Core denoising network operating in latent space.
- ▶ **Decoder:** Reconstructs images from latent representations.
- ▶ **Text Encoder:** CLIP-based, encodes prompts for conditioning.

Latent Diffusion Models: Visual Overview



- ▶ Initialize from Stable Diffusion.
- ▶ Add temporal layers (e.g., 1D Conv, temporal attention).
- ▶ Finetune on video data.
- ▶ Data pipeline:
 - Scrape **unlabeled** videos from the internet.
 - Use image and video captioners to generate synthetic text labels.
 - Apply aggressive data filtering using several heuristics (CLIP score, optical flow score, aesthetic scores, OCR, etc.).


Stable Video Diffusion (cont.)





Diffusion Models: **Issues with Diffusion Models**

- **Slow Sampling:** Diffusion models require hundreds to thousands of iterative steps to generate a sample, making them computationally expensive and slow compared to other generative models.



Generative Learning Trilemma: Balancing sample quality, diversity, and fast sampling.

► Faster Sampling:

- **Denoising Diffusion Implicit Models (DDIM):** Propose a non-Markovian process for faster and deterministic sampling by skipping steps and refining the sample.

► Improved Training:

- **Improved Denoising Diffusion Probabilistic Models:** Learn the noise variance σ^2 during training, rather than fixing it, for better flexibility and performance.

► Score-Based Modeling:

- **Score-Based Generative Modeling through SDEs:** Model the gradient of the log-density (score function) using stochastic differential equations for improved sample quality.

► Combining GANs and Diffusion:

- **Denoising Diffusion GANs:** Use GANs to model each denoising step, addressing the trilemma by improving speed and sample quality.

► High-Fidelity Generation:

- **Cascaded Diffusion Models:** Employ a cascade of diffusion models at different resolutions to boost sample fidelity.

Diffusion Models: **Applications**



“a man wearing a white hat”

Image Inpainting with GLIDE

⁰GLIDE: Towards Photorealistic Image Generation and Editing with Text-Guided Diffusion Models



a shiba inu wearing a beret and black turtleneck




a close up of a handpalm with leaves growing from it

Text to image generation with DALL.E 2



Fix the CLIP embedding z_i
Decode using different decoder latents x_T

Image Variations



Interpolate image CLIP embeddings z_i

Use different x_T to get different interpolation trajectories.

Image interpolation



a photo of a cat → an anime drawing of a super saiyan cat, artstation



a photo of a victorian house → a photo of a modern house



a photo of an adult lion → a photo of lion cub

Change the image CLIP embedding towards the difference of the text CLIP embeddings of two prompts.

Decoder latent is kept as a constant.

Text Difference Image interpolation



A brain riding a rocketship heading towards the moon.





A dragon fruit wearing karate belt in the snow.



A relaxed garlic with a blindfold reading a newspaper while floating in a pool of tomato soup.


Application: Nano Banana 2 – Google DeepMind



Examples: water-cycle infographic and cloud-type comparison.


- Builds on Google's Gemini Image line: Nano Banana (2025) → Nano Banana Pro → Nano Banana 2.
- Production-ready output: multiple aspect ratios and resolutions from 512px to 4K.

⁰Raisinghani, "Nano Banana 2: Combining Pro capabilities with lightning-fast speed," Google DeepMind blog (Feb 2026).




Changing the semantic latent z_{sem}

Learning semantic meaningful latent representations in diffusion models



Natural Image Super Resolution $64 \times 64 \rightarrow 256 \times 256$

3D Shape Generation



⁰3D Shape Generation and Completion through Point-Voxel Diffusion

Text to Image generation with stable diffusion.
<https://huggingface.co/spaces/stabilityai/stable-diffusion>



Diffusion Models: **Summary & Limitations**

- ▶ Diffusion models add noise to data and learn to reverse it
- ▶ Core task: Predict noise at different steps
- ▶ Training is efficient; Sampling is iterative and slow
- ▶ U-Nets with attention are the winning architecture

- ▶ **Slow sampling** – iterative process
- ▶ **Data-heavy** – needs lots of varied data
- ▶ **Distribution miss** – struggles in tails of data
- ▶ **Compute needs** – high GPU/time investment
- ▶ **Misinformation concerns** – used in deepfake creation

Tutorials and Surveys

- ▶ Kreis, K., Gao, R., & Vahdat, A.: [CVPR 2022 Tutorial: Diffusion Models](#)
- ▶ Rogge, L., & Rasul, K.: [The Annotated Diffusion Model](#)
- ▶ Yang Song: [Score-Based Generative Modeling: A Brief Overview](#)
- ▶ Lilian Weng: [What are Diffusion Models?](#)
- ▶ Binxu Wang: [Stable Diffusion: A Tutorial \(Harvard ML from Scratch\)](#)
- ▶ Cornell CS4782 (2025): [Diffusion Models II \(Course Slides\)](#)

Foundational Papers

- ▶ Ho, J., Jain, A., & Abbeel, P. (2020). [Denoising Diffusion Probabilistic Models](#)
- ▶ Sohl-Dickstein, J., Weiss, E., Maheswaranathan, N., & Ganguli, S. (2015). [Deep Unsupervised Learning using Nonequilibrium Thermodynamics](#)
- ▶ Song, Y., & Ermon, S. (2019). [Generative Modeling by Estimating Gradients of the Data Distribution](#)
- ▶ Nichol, A., & Dhariwal, P. (2021). [Improved Denoising Diffusion Probabilistic Models](#)

Applications and Notable Models

- ▶ Ramesh, A., et al. (2022). [Hierarchical Text-Conditional Image Generation with CLIP Latents \(DALL·E 2\)](#)
- ▶ Saharia, C., et al. (2022). [Photorealistic Text-to-Image Diffusion Models with Deep Language Understanding \(Imagen\)](#)
- ▶ Nichol, A., Dhariwal, P., et al. (2021). [GLIDE: Towards Photorealistic Image Generation and Editing with Text-Guided Diffusion Models](#)
- ▶ Preechakul, T., et al. (2022). [Diffusion Autoencoders: Toward a Meaningful and Decodable Representation](#)
- ▶ Saharia, C., et al. (2022). [Image Super-Resolution via Iterative Refinement](#)
- ▶ Poole, B., et al. (2022). [DreamFusion: Text-to-3D using 2D Diffusion](#)